

cnn_v4

October 13, 2022

```
[ ]: %env LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$CONDA_PREFIX/lib/
```

```
env: LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$CONDA_PREFIX/lib/
```

```
[ ]: import os
print(os.environ["LD_LIBRARY_PATH"])
```

```
$LD_LIBRARY_PATH:$CONDA_PREFIX/lib/
```

```
[ ]: import tensorflow as tf
import numpy as np
import pandas as pd
import os
import keras

from keras import Sequential, models, Input
from keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, Dropout, LeakyReLU
from keras.optimizers import SGD, Adam
```

```
[ ]: from tensorflow.python.client import device_lib

#print(device_lib.list_local_devices())
print('Default GPU Device: {}'.format(tf.test.gpu_device_name()))
```

0.1 Read the csv dataset to get the values for stage and discharge of the images

```
[ ]: df = pd.read_csv("../dataset/2012_2019_PlatteRiverWeir_features_merged_all.csv")
df.head()
```

```
[ ]: Unnamed: 0      SensorTime      CaptureTime \
0          0  2012-06-09 13:15:00  2012-06-09T13:09:07
1          1  2012-06-09 13:15:00  2012-06-09T13:10:29
2          2  2012-06-09 13:45:00  2012-06-09T13:44:01
3          3  2012-06-09 14:45:00  2012-06-09T14:44:30
4          4  2012-06-09 15:45:00  2012-06-09T15:44:59
```

	Filename	Agency	SiteNumber	TimeZone	Stage	\
0	StateLineWeir_20120609_Farrell_001.jpg	USGS	6674500	MDT	2.99	
1	StateLineWeir_20120609_Farrell_002.jpg	USGS	6674500	MDT	2.99	
2	StateLineWeir_20120609_Farrell_003.jpg	USGS	6674500	MDT	2.96	
3	StateLineWeir_20120609_Farrell_004.jpg	USGS	6674500	MDT	2.94	
4	StateLineWeir_20120609_Farrell_005.jpg	USGS	6674500	MDT	2.94	

	Discharge	CalcTimestamp	...	WeirPt2X	WeirPt2Y	WwRawLineMin	\
0	916.0	2020-03-11T16:58:28	...	-1	-1	0.0	
1	916.0	2020-03-11T16:58:33	...	-1	-1	0.0	
2	873.0	2020-03-11T16:58:40	...	-1	-1	0.0	
3	846.0	2020-03-11T16:58:47	...	-1	-1	0.0	
4	846.0	2020-03-11T16:58:55	...	-1	-1	0.0	

	WwRawLineMax	WwRawLineMean	WwRawLineSigma	WwCurveLineMin	\
0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	

	WwCurveLineMax	WwCurveLineMean	WwCurveLineSigma
0	0.0	0.0	0.0
1	0.0	0.0	0.0
2	0.0	0.0	0.0
3	0.0	0.0	0.0
4	0.0	0.0	0.0

[5 rows x 60 columns]

```
[ ]: df = df[["Filename", "Stage", "Discharge"]]
```

0.1.1 Scale the data

```
[ ]: from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
```

```
[ ]: from joblib import dump, load
dump(scaler, '../best_models_weights/std_scaler.joblib', compress=True)
```

```
[ ]: ['../best_models_weights/std_scaler.joblib']
```

```
[ ]: df[["Stage", "Discharge"]] = scaler.fit_transform(df[["Stage", "Discharge"]])
df
```

```
[ ]:
```

	Filename	Stage	Discharge
0	StateLineWeir_20120609_Farrell_001.jpg	0.138117	-0.046094
1	StateLineWeir_20120609_Farrell_002.jpg	0.138117	-0.046094
2	StateLineWeir_20120609_Farrell_003.jpg	0.100875	-0.082160
3	StateLineWeir_20120609_Farrell_004.jpg	0.076046	-0.104807
4	StateLineWeir_20120609_Farrell_005.jpg	0.076046	-0.104807
...
42054	StateLineWeir_20191011_Farrell_409.jpg	-0.420526	-0.450369
42055	StateLineWeir_20191011_Farrell_410.jpg	-0.420526	-0.450369
42056	StateLineWeir_20191011_Farrell_411.jpg	-0.420526	-0.450369
42057	StateLineWeir_20191011_Farrell_412.jpg	-0.420526	-0.450369
42058	StateLineWeir_20191011_Farrell_413.jpg	-0.420526	-0.450369

[42059 rows x 3 columns]

0.2 Create the dataset pipeline

```
[ ]: IMG_SIZE = 512
      BATCH_SIZE = 32
```

```
[ ]: from glob import glob

def make_dataset(path, batch_size, df, seed=None):
    np.random.seed(seed)

    def parse_image(filename):
        image = tf.io.read_file(filename)
        image = tf.image.decode_jpeg(image, channels=3)
        #image = tf.image.resize(image, [IMG_SIZE, IMG_SIZE])
        image = tf.cast(image, tf.float32)
        image /= 255
        return image

    def configure_for_performance(ds):
        ds = ds.shuffle(buffer_size=100)
        ds = ds.batch(batch_size)
        ds = ds.repeat()
        ds = ds.prefetch(buffer_size=tf.data.experimental.AUTOTUNE)
        return ds

    filenames = glob(path + '/*')

    # make train, val and test splits of the dataset (70%, 10%, 20% split)
    split1 = int(0.7 * len(filenames))
    split2 = int(0.8 * len(filenames))

    np.random.shuffle(filenames)
```

```

train_files = filenames[:split1] # up to split 1 (ex 70%)
val_files = filenames[split1:split2] # from ex. 70% to 80%
test_files = filenames[split2:] # from ex. 80% until the end

# create stage values
stage_train_values = [df[df.Filename == file.split('/')[0]].Stage.values for
→file in train_files]
stage_val_values = [df[df.Filename == file.split('/')[0]].Stage.values for
→file in val_files]
stage_test_values = [df[df.Filename == file.split('/')[0]].Stage.values for
→file in test_files]

# create discharge values
discharge_train_values = [df[df.Filename == file.split(
    '/')[-1]].Discharge.values for file in train_files]
discharge_val_values = [df[df.Filename == file.split(
    '/')[-1]].Discharge.values for file in val_files]
discharge_test_values = [df[df.Filename == file.split(
    '/')[-1]].Discharge.values for file in test_files]

# join stage and discharge values
stage_discharge_train_values = [[np.squeeze(s), np.squeeze(d)] for s, d in
→zip(stage_train_values, discharge_train_values)]
stage_discharge_val_values = [[np.squeeze(s), np.squeeze(d)] for s, d in
→zip(stage_val_values, discharge_val_values)]
stage_discharge_test_values = [[np.squeeze(s), np.squeeze(
    d)] for s, d in zip(stage_test_values, discharge_test_values)]

# create images dataset (train, val, test)
filenames_train_ds = tf.data.Dataset.from_tensor_slices(train_files)
filenames_val_ds = tf.data.Dataset.from_tensor_slices(val_files)
filenames_test_ds = tf.data.Dataset.from_tensor_slices(test_files)

images_train_ds = filenames_train_ds.map(parse_image, num_parallel_calls=5)
images_val_ds = filenames_val_ds.map(parse_image, num_parallel_calls=5)
images_test_ds = filenames_test_ds.map(parse_image, num_parallel_calls=5)

# create stage and discharge dataset (train, val, test)
stage_discharge_train_ds = tf.data.Dataset.
→from_tensor_slices(stage_discharge_train_values)
stage_discharge_val_ds = tf.data.Dataset.
→from_tensor_slices(stage_discharge_val_values)
stage_discharge_test_ds = tf.data.Dataset.from_tensor_slices(
    stage_discharge_test_values)

# create tensorflow dataset of images and values (train, val, test)

```

```

train_ds = tf.data.Dataset.zip((images_train_ds, stage_discharge_train_ds))
train_ds = configure_for_performance(train_ds)
val_ds = tf.data.Dataset.zip((images_val_ds, stage_discharge_val_ds))
val_ds = configure_for_performance(val_ds)
test_ds = tf.data.Dataset.zip((images_test_ds, stage_discharge_test_ds))
test_ds = configure_for_performance(test_ds)

return train_ds, len(train_files), val_ds, len(val_files), test_ds,
↪len(test_files)

```

```

[ ]: path = "../dataset/images_tmp"

train_ds, train_size, val_ds, val_size, test_ds, test_size = make_dataset(path,
↪BATCH_SIZE, df, 0)

```

```

[ ]: input_shape = 0
output_shape = 0

for image, stage_discharge in train_ds.take(1):
    print(image.numpy().shape)
    print(stage_discharge.numpy().shape)

    input_shape = image.numpy().shape[1:]
    output_shape = stage_discharge.numpy().shape[1:]

```

```

(32, 512, 512, 3)
(32, 2)

```

```

[ ]: print(input_shape)
print(output_shape)

```

```

(512, 512, 3)
(2,)

```

0.3 Create model

```

[ ]: def create_model(input_shape, output_shape):
    model = Sequential()

    model.add(Input(shape=input_shape))

    model.add(Conv2D(64, kernel_size=(4, 4), strides=(2, 2), padding='same',
↪activation=LeakyReLU()))
    model.add(MaxPooling2D(pool_size=(4, 4)))
    model.add(Conv2D(64, kernel_size=(4, 4), activation=LeakyReLU(),
↪padding='same'))
    model.add(MaxPooling2D(pool_size=(2, 2)))

```

```

model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(3, 3)))
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())
model.add(Dense(64, activation='relu'))
model.add(Dropout(0.3))
model.add(Dense(32, activation='sigmoid'))
#model.add(Dropout(0.3))
model.add(Dense(32, activation='sigmoid'))
model.add(Dense(output_shape, activation='linear')) # linear regression
↳ output layer

return model

```

```
[ ]: model = create_model(input_shape, output_shape[0])
```

```
[ ]: model.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
conv2d_12 (Conv2D)	(None, 256, 256, 64)	3136
max_pooling2d_12 (MaxPooling2D)	(None, 64, 64, 64)	0
conv2d_13 (Conv2D)	(None, 64, 64, 64)	65600
max_pooling2d_13 (MaxPooling2D)	(None, 32, 32, 64)	0
conv2d_14 (Conv2D)	(None, 32, 32, 32)	18464
max_pooling2d_14 (MaxPooling2D)	(None, 10, 10, 32)	0
conv2d_15 (Conv2D)	(None, 8, 8, 32)	9248
max_pooling2d_15 (MaxPooling2D)	(None, 4, 4, 32)	0
flatten_3 (Flatten)	(None, 512)	0

dense_12 (Dense)	(None, 64)	32832
dropout_3 (Dropout)	(None, 64)	0
dense_13 (Dense)	(None, 32)	2080
dense_14 (Dense)	(None, 32)	1056
dense_15 (Dense)	(None, 2)	66

```
=====
Total params: 132,482
Trainable params: 132,482
Non-trainable params: 0
-----
```

```
[ ]: def compile_model(loss_func, optimizer, metrics=["accuracy"]):
    model.compile(loss=loss_func, optimizer=optimizer, metrics=metrics)

[ ]: sgd = SGD(learning_rate=0.01, decay=1e-4, momentum=0.9, nesterov=True)
adam = Adam(learning_rate=1e-3, decay=1e-3 / 200)

compile_model('mse', adam, [
    'mse', tf.keras.metrics.RootMeanSquaredError(name='rmse'), 'mae',
    'mape'])

[ ]: def fit_model(training_values, validation_values=None, batch_size=32,
    epochs=10, steps=32, val_steps=32, callbacks=[]):
    return model.fit(training_values, validation_data=validation_values,
    batch_size=batch_size, epochs=epochs, steps_per_epoch=steps,
    validation_steps=val_steps, callbacks=callbacks)

[ ]: import datetime

log_dir = "logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir,
    histogram_freq=1)

checkpoint_callback = tf.keras.callbacks.
    ModelCheckpoint(filepath="model_weights/cnn_best_weights.hdf5",
        monitor='val_mse',
        verbose=1,
        save_best_only=True)

[ ]: # batch_size = 0 because we already have batch size in tf dataset
```

```
history = fit_model(train_ds, val_ds, batch_size=0, epochs=20, steps=np.
↳ceil(train_size / BATCH_SIZE), val_steps=np.ceil(val_size / BATCH_SIZE),
↳callbacks=[tensorboard_callback, checkpoint_callback])
```

0.4 Evaluate model

```
[ ]: def evaluate_model(test_values, steps):
    score = model.evaluate(test_values, steps=steps)
    return score
```

```
[ ]: test_loss, test_mse, test_rmse, test_mae, test_mape = evaluate_model(test_ds,
↳steps=np.ceil(test_size / BATCH_SIZE))
```

```
99/99 [=====] - 6s 59ms/step - loss: 0.0072 - mse:
0.0072 - rmse: 0.0848 - mae: 0.0573 - mape: 24.5317
```

```
[ ]: predictions = model.predict(test_ds, steps=np.ceil(test_size / BATCH_SIZE))
```

```
99/99 [=====] - 6s 57ms/step
```

```
[ ]: for image, stage_discharge in train_ds.take(1):
    diff = predictions[:2].flatten() - stage_discharge[:2].numpy().flatten()
    percentDiff = (diff / stage_discharge[:2].numpy().flatten()) * 100
    absPercentDiff = np.abs(percentDiff)
    # compute the mean and standard deviation of the absolute percentage
    # difference
    mean = np.mean(absPercentDiff)
    std = np.std(absPercentDiff)
    # finally, show some statistics on our model
    print(mean)
    print(std)
    print(f"pred stage: {scaler.inverse_transform(predictions[:2])[0]},
↳actual stage: {scaler.inverse_transform(stage_discharge[:2].numpy())[0]}")
    print(f"pred discharge: {scaler.inverse_transform(predictions[:2])[0]},
↳actual discharge: {scaler.inverse_transform(stage_discharge[:2].
↳numpy())[0]}")
```

```
174.55231176555762
```

```
222.93393032335572
```

```
pred stage: [ 3.097657 1080.6653 ], actual stage: [ 3.13 1140. ]
```

```
pred discharge: [ 3.097657 1080.6653 ], actual discharge: [ 3.13 1140. ]
```

```
[ ]:
```


0.5 Visualize layers

```
[ ]: layer_outputs = [layer.output for layer in model.layers[:12]]
# Extracts the outputs of the top 12 layers
activation_model = models.Model(inputs=model.input, outputs=layer_outputs) #
↳ Creates a model that will return these outputs, given the model input

[ ]: activations = activation_model.predict(test_ds, steps=np.ceil(test_size / 32))

2/2 [=====] - 0s 42ms/step

[ ]: import matplotlib.pyplot as plt

layer_names = []
for layer in model.layers[:12]:
    layer_names.append(layer.name) # Names of the layers, so you can have them
↳ as part of your plot

images_per_row = 16

for layer_name, layer_activation in zip(layer_names, activations): # Displays
↳ the feature maps
    n_features = layer_activation.shape[-1] # Number of features in the feature
↳ map
    size = layer_activation.shape[1] # The feature map has shape (1, size, size,
↳ n_features).
    n_cols = n_features // images_per_row # Tiles the activation channels in
↳ this matrix
    display_grid = np.zeros((size * n_cols, images_per_row * size))

    print(layer_name)
    if ("Dense" in layer_name): break

    for col in range(n_cols): # Tiles each filter into a big horizontal grid
        for row in range(images_per_row):
            channel_image = layer_activation[0,
                                                :, :,
                                                col * images_per_row + row]
            channel_image -= channel_image.mean() # Post-processes the feature
↳ to make it visually palatable
            channel_image /= channel_image.std()
            channel_image *= 64
            channel_image += 128
            channel_image = np.clip(channel_image, 0, 255).astype('uint8')
            display_grid[col * size : (col + 1) * size, # Displays the grid
                            row * size : (row + 1) * size] = channel_image

    scale = 1. / size
```

```
plt.figure(figsize=(scale * display_grid.shape[1],
                    scale * display_grid.shape[0]))
plt.title(layer_name)
plt.grid(False)
plt.imshow(display_grid, aspect='auto', cmap='viridis')
```

```
conv2d_12
max_pooling2d_12
conv2d_13
max_pooling2d_13
conv2d_14
max_pooling2d_14
conv2d_15
max_pooling2d_15
flatten_3
```

```
/tmp/ipykernel_19196/1622603125.py:24: RuntimeWarning: invalid value encountered
in divide
```

```
channel_image /= channel_image.std()
```

```
-----
IndexError                                Traceback (most recent call last)
Cell In [71], line 20
    18 for col in range(n_cols): # Tiles each filter into a big horizontal grid
    19     for row in range(images_per_row):
---> 20         channel_image = layer_activation[0,
    21                                         :, :,
    22                                         col * images_per_row + row]
    23         channel_image -= channel_image.mean() # Post-processes the
↳ feature to make it visually palatable
    24         channel_image /= channel_image.std()

IndexError: too many indices for array: array is 2-dimensional, but 4 were
↳ indexed
```





