

cnn_v5

October 13, 2022

```
[ ]: %env LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$CONDA_PREFIX/lib/
```

```
env: LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$CONDA_PREFIX/lib/
```

```
[ ]: import os
print(os.environ["LD_LIBRARY_PATH"])
```

```
:/home/nkspartan/miniconda3/envs/tf-gpu/lib/:/home/nkspartan/miniconda3/envs/tf-gpu/lib/
```

```
[ ]: import tensorflow as tf
import numpy as np
import pandas as pd
import os
import keras

from keras import Sequential, models, Input
from keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, Dropout, LeakyReLU
from keras.optimizers import SGD, Adam
```

```
2022-10-13 22:32:01.785635: I tensorflow/core/platform/cpu_feature_guard.cc:193]
This TensorFlow binary is optimized with oneAPI Deep Neural Network Library
(oneDNN) to use the following CPU instructions in performance-critical
operations:  AVX2 FMA
```

```
To enable them in other operations, rebuild TensorFlow with the appropriate
compiler flags.
```

```
2022-10-13 22:32:02.288671: E tensorflow/stream_executor/cuda/cuda_blas.cc:2981]
Unable to register cuBLAS factory: Attempting to register factory for plugin
cuBLAS when one has already been registered
```

```
2022-10-13 22:32:03.388483: W
```

```
tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load
dynamic library 'libnvinfer.so.7'; dLError: libnvinfer.so.7: cannot open shared
object file: No such file or directory; LD_LIBRARY_PATH:
:/home/nkspartan/miniconda3/envs/tf-gpu/lib/:/home/nkspartan/miniconda3/envs/tf-gpu/lib/
```

```
2022-10-13 22:32:03.388691: W
```

```
tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load
```

```
dynamic library 'libnvinfer_plugin.so.7'; dlopen: libnvinfer_plugin.so.7:
cannot open shared object file: No such file or directory; LD_LIBRARY_PATH:
:/home/nkspartan/miniconda3/envs/tf-gpu/lib/:/home/nkspartan/miniconda3/envs/tf-
gpu/lib/
2022-10-13 22:32:03.388696: W
tensorflow/compiler/tf2tensorrt/utils/py_utils.cc:38] TF-TRT Warning: Cannot
dlopen some TensorRT libraries. If you would like to use Nvidia GPU with
TensorRT, please make sure the missing libraries mentioned above are installed
properly.
```

```
[ ]: from tensorflow.python.client import device_lib

#print(device_lib.list_local_devices())
print('Default GPU Device: {}'.format(tf.test.gpu_device_name()))
```

Default GPU Device: /device:GPU:0

```
2022-10-13 22:32:06.425925: I tensorflow/core/platform/cpu_feature_guard.cc:193]
This TensorFlow binary is optimized with oneAPI Deep Neural Network Library
(oneDNN) to use the following CPU instructions in performance-critical
operations:  AVX2 FMA
```

To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

```
2022-10-13 22:32:06.448950: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
```

```
2022-10-13 22:32:06.496166: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
```

```
2022-10-13 22:32:06.496370: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
```

```
2022-10-13 22:32:07.290819: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
```

```
2022-10-13 22:32:07.291450: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
```

```
2022-10-13 22:32:07.291606: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
```

```
2022-10-13 22:32:07.291738: I
```

```
tensorflow/core/common_runtime/gpu/gpu_device.cc:1616] Created device
/device:GPU:0 with 4078 MB memory: -> device: 0, name: NVIDIA GeForce RTX 2060,
pci bus id: 0000:08:00.0, compute capability: 7.5
```

0.1 Read the csv dataset to get the values for stage and discharge of the images

```
[ ]: df = pd.read_csv("../dataset/2012_2019_PlatteRiverWeir_features_merged_all.
↪csv")
df.head()
```

```
[ ]: Unnamed: 0      SensorTime      CaptureTime \
0          0  2012-06-09 13:15:00  2012-06-09T13:09:07
1          1  2012-06-09 13:15:00  2012-06-09T13:10:29
2          2  2012-06-09 13:45:00  2012-06-09T13:44:01
3          3  2012-06-09 14:45:00  2012-06-09T14:44:30
4          4  2012-06-09 15:45:00  2012-06-09T15:44:59
```

```
      Filename Agency  SiteNumber TimeZone  Stage \
0  StateLineWeir_20120609_Farrell_001.jpg  USGS    6674500      MDT    2.99
1  StateLineWeir_20120609_Farrell_002.jpg  USGS    6674500      MDT    2.99
2  StateLineWeir_20120609_Farrell_003.jpg  USGS    6674500      MDT    2.96
3  StateLineWeir_20120609_Farrell_004.jpg  USGS    6674500      MDT    2.94
4  StateLineWeir_20120609_Farrell_005.jpg  USGS    6674500      MDT    2.94
```

```
      Discharge      CalcTimestamp  ...  WeirPt2X  WeirPt2Y  WwRawLineMin  \
0      916.0  2020-03-11T16:58:28  ...      -1      -1           0.0
1      916.0  2020-03-11T16:58:33  ...      -1      -1           0.0
2      873.0  2020-03-11T16:58:40  ...      -1      -1           0.0
3      846.0  2020-03-11T16:58:47  ...      -1      -1           0.0
4      846.0  2020-03-11T16:58:55  ...      -1      -1           0.0
```

```
      WwRawLineMax  WwRawLineMean  WwRawLineSigma  WwCurveLineMin  \
0           0.0           0.0           0.0           0.0
1           0.0           0.0           0.0           0.0
2           0.0           0.0           0.0           0.0
3           0.0           0.0           0.0           0.0
4           0.0           0.0           0.0           0.0
```

```
      WwCurveLineMax  WwCurveLineMean  WwCurveLineSigma
0           0.0           0.0           0.0
1           0.0           0.0           0.0
2           0.0           0.0           0.0
3           0.0           0.0           0.0
4           0.0           0.0           0.0
```

```
[5 rows x 60 columns]
```

```
[ ]: df = df[["Filename", "Stage", "Discharge"]]
```

0.1.1 Scale the data

```
[ ]: from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
```

```
[ ]: df[["Stage", "Discharge"]] = scaler.fit_transform(df[["Stage", "Discharge"]])
df
```

```
[ ]:
```

	Filename	Stage	Discharge
0	StateLineWeir_20120609_Farrell_001.jpg	0.138117	-0.046094
1	StateLineWeir_20120609_Farrell_002.jpg	0.138117	-0.046094
2	StateLineWeir_20120609_Farrell_003.jpg	0.100875	-0.082160
3	StateLineWeir_20120609_Farrell_004.jpg	0.076046	-0.104807
4	StateLineWeir_20120609_Farrell_005.jpg	0.076046	-0.104807
...
42054	StateLineWeir_20191011_Farrell_409.jpg	-0.420526	-0.450369
42055	StateLineWeir_20191011_Farrell_410.jpg	-0.420526	-0.450369
42056	StateLineWeir_20191011_Farrell_411.jpg	-0.420526	-0.450369
42057	StateLineWeir_20191011_Farrell_412.jpg	-0.420526	-0.450369
42058	StateLineWeir_20191011_Farrell_413.jpg	-0.420526	-0.450369

```
[42059 rows x 3 columns]
```

```
[ ]: from joblib import dump, load
dump(scaler, 'std_scaler.joblib', compress=True)
```

```
[ ]: ['std_scaler.joblib']
```

0.2 Create the dataset pipeline

```
[ ]: IMG_SIZE = 512
BATCH_SIZE = 32
```

```
[ ]: from glob import glob

def make_dataset(path, batch_size, df, seed=None):
    np.random.seed(seed)

    def parse_image(filename):
        image = tf.io.read_file(filename)
        image = tf.image.decode_jpeg(image, channels=3)
        #image = tf.image.resize(image, [IMG_SIZE, IMG_SIZE])
        image = tf.cast(image, tf.float32)
```

```

    image /= 255
    return image

def configure_for_performance(ds):
    ds = ds.shuffle(buffer_size=100)
    ds = ds.batch(batch_size)
    ds = ds.repeat()
    ds = ds.prefetch(buffer_size=tf.data.experimental.AUTOTUNE)
    return ds

filenames = glob(path + '/*')

# make train, val and test splits of the dataset (70%, 10%, 20% split)
split1 = int(0.7 * len(filenames))
split2 = int(0.8 * len(filenames))

np.random.shuffle(filenames)
train_files = filenames[:split1] # up to split 1 (ex 70%)
val_files = filenames[split1:split2] # from ex. 70% to 80%
test_files = filenames[split2:] # from ex. 80% until the end

# create stage values
stage_train_values = [df[df.Filename == file.split('/')[0]].Stage.values for
↪file in train_files]
stage_val_values = [df[df.Filename == file.split('/')[0]].Stage.values for
↪file in val_files]
stage_test_values = [df[df.Filename == file.split('/')[0]].Stage.values for
↪file in test_files]

# create discharge values
discharge_train_values = [df[df.Filename == file.split(
    '/')[-1]].Discharge.values for file in train_files]
discharge_val_values = [df[df.Filename == file.split(
    '/')[-1]].Discharge.values for file in val_files]
discharge_test_values = [df[df.Filename == file.split(
    '/')[-1]].Discharge.values for file in test_files]

# join stage and discharge values
stage_discharge_train_values = [[np.squeeze(s), np.squeeze(d)] for s, d in
↪zip(stage_train_values, discharge_train_values)]
stage_discharge_val_values = [[np.squeeze(s), np.squeeze(d)] for s, d in
↪zip(stage_val_values, discharge_val_values)]
stage_discharge_test_values = [[np.squeeze(s), np.squeeze(
    d)] for s, d in zip(stage_test_values, discharge_test_values)]

# create images dataset (train, val, test)
filenames_train_ds = tf.data.Dataset.from_tensor_slices(train_files)

```

```

filenames_val_ds = tf.data.Dataset.from_tensor_slices(val_files)
filenames_test_ds = tf.data.Dataset.from_tensor_slices(test_files)

images_train_ds = filenames_train_ds.map(parse_image, num_parallel_calls=5)
images_val_ds = filenames_val_ds.map(parse_image, num_parallel_calls=5)
images_test_ds = filenames_test_ds.map(parse_image, num_parallel_calls=5)

# create stage and discharge dataset (train, val, test)
stage_discharge_train_ds = tf.data.Dataset.
↳from_tensor_slices(stage_discharge_train_values)
stage_discharge_val_ds = tf.data.Dataset.
↳from_tensor_slices(stage_discharge_val_values)
stage_discharge_test_ds = tf.data.Dataset.from_tensor_slices(
    stage_discharge_test_values)

# create tensorflow dataset of images and values (train, val, test)
train_ds = tf.data.Dataset.zip((images_train_ds, stage_discharge_train_ds))
train_ds = configure_for_performance(train_ds)
val_ds = tf.data.Dataset.zip((images_val_ds, stage_discharge_val_ds))
val_ds = configure_for_performance(val_ds)
test_ds = tf.data.Dataset.zip((images_test_ds, stage_discharge_test_ds))
test_ds = configure_for_performance(test_ds)

return train_ds, len(train_files), val_ds, len(val_files), test_ds,
↳len(test_files)

```

```

[ ]: path = "../..//dataset/images_tmp"

train_ds, train_size, val_ds, val_size, test_ds, test_size = make_dataset(path,
↳BATCH_SIZE, df, 0)

```

```

2022-10-13 22:33:36.393434: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-13 22:33:36.393646: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-13 22:33:36.393786: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-13 22:33:36.394135: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero

```

```

2022-10-13 22:33:36.394284: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-13 22:33:36.394423: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-13 22:33:36.394685: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-13 22:33:36.394835: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-13 22:33:36.394952: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:1616] Created device
/job:localhost/replica:0/task:0/device:GPU:0 with 4078 MB memory: -> device: 0,
name: NVIDIA GeForce RTX 2060, pci bus id: 0000:08:00.0, compute capability: 7.5

```

```

[ ]: input_shape = 0
output_shape = 0

for image, stage_discharge in train_ds.take(1):
    print(image.numpy().shape)
    print(stage_discharge.numpy().shape)

    input_shape = image.numpy().shape[1:]
    output_shape = stage_discharge.numpy().shape[1:]

```

```

(32, 512, 512, 3)
(32, 2)

```

```

[ ]: print(input_shape)
print(output_shape)

```

```

(512, 512, 3)
(2,)

```

0.3 Create model

```

[ ]: def create_model(input_shape, output_shape):
    model = Sequential()

    model.add(Input(shape=input_shape))

```

```

        model.add(Conv2D(64, kernel_size=(4, 4), strides=(2, 2), padding='same',
↪activation=LeakyReLU()))
        model.add(MaxPooling2D(pool_size=(4, 4)))
        model.add(Conv2D(64, kernel_size=(4, 4), activation=LeakyReLU(),
↪padding='same'))
        model.add(MaxPooling2D(pool_size=(2, 2)))
        model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', padding='same'))
        model.add(MaxPooling2D(pool_size=(3, 3)))
        model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
        model.add(MaxPooling2D(pool_size=(2, 2)))

        model.add(Flatten())
        model.add(Dense(64, activation='relu'))
        model.add(Dense(64, activation='relu'))
        model.add(Dense(32, activation='relu'))
        model.add(Dense(32, activation='sigmoid'))
        model.add(Dense(output_shape, activation='linear')) # linear regression
↪output layer

    return model

```

```
[ ]: model = create_model(input_shape, output_shape[0])
```

```
[ ]: model.summary()
```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
conv2d_24 (Conv2D)	(None, 256, 256, 64)	3136
max_pooling2d_24 (MaxPooling2D)	(None, 64, 64, 64)	0
conv2d_25 (Conv2D)	(None, 64, 64, 64)	65600
max_pooling2d_25 (MaxPooling2D)	(None, 32, 32, 64)	0
conv2d_26 (Conv2D)	(None, 32, 32, 32)	18464
max_pooling2d_26 (MaxPooling2D)	(None, 10, 10, 32)	0
conv2d_27 (Conv2D)	(None, 8, 8, 32)	9248

max_pooling2d_27 (MaxPoolin	(None, 4, 4, 32)	0
g2D)		
flatten_6 (Flatten)	(None, 512)	0
dense_30 (Dense)	(None, 128)	65664
dense_31 (Dense)	(None, 32)	4128
dense_32 (Dense)	(None, 32)	1056
dense_33 (Dense)	(None, 32)	1056
dense_34 (Dense)	(None, 2)	66

```

=====
Total params: 168,418
Trainable params: 168,418
Non-trainable params: 0
-----

```

```
[ ]: def compile_model(loss_func, optimizer, metrics=["accuracy"]):
      model.compile(loss=loss_func, optimizer=optimizer, metrics=metrics)
```

```
[ ]: sgd = SGD(learning_rate=0.01, decay=1e-4, momentum=0.9, nesterov=True)
      adam = Adam(learning_rate=1e-3, decay=1e-3 / 100)

      compile_model('mse', adam, [
          'mse', tf.keras.metrics.RootMeanSquaredError(name='rmse'), 'mae', ↵
          ↪ 'mape'])
```

```
[ ]: def fit_model(training_values, validation_values=None, batch_size=32, ↵
      ↪ epochs=10, steps=32, val_steps=32, callbacks=[]):
      return model.fit(training_values, validation_data=validation_values, ↵
      ↪ batch_size=batch_size, epochs=epochs, steps_per_epoch=steps, ↵
      ↪ validation_steps=val_steps, callbacks=callbacks)
```

```
[ ]: import datetime

      date_actual = datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
      log_dir = "logs/fit/" + date_actual
      tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir, ↵
      ↪ histogram_freq=1)

      checkpoint_callback = tf.keras.callbacks.
      ↪ ModelCheckpoint(filepath=f"model_weights/{date_actual}_cnn_best_weights.
      ↪ hdf5",
```

```
monitor='val_mse',  
verbose=1,  
save_best_only=True)
```

```
[ ]: # batch_size = 0 because we already have batch size in tf dataset  
history = fit_model(train_ds, val_ds, batch_size=0, epochs=20, steps=np.  
    ↳ ceil(train_size / BATCH_SIZE), val_steps=np.ceil(val_size / BATCH_SIZE),  
    ↳ callbacks=[tensorboard_callback, checkpoint_callback])
```

Epoch 1/20

347/348 [=====>.] - ETA: 0s - loss: 0.3567 - mse: 0.3567
- rmse: 0.5972 - mae: 0.3288 - mape: 162.6058

Epoch 1: val_mse improved from inf to 0.10324, saving model to
model_weights/20221013-221702_cnn_best_weights.hdf5

348/348 [=====] - 39s 111ms/step - loss: 0.3565 - mse:
0.3565 - rmse: 0.5971 - mae: 0.3287 - mape: 162.5373 - val_loss: 0.1032 -
val_mse: 0.1032 - val_rmse: 0.3213 - val_mae: 0.1697 - val_mape: 87.4317

Epoch 2/20

347/348 [=====>.] - ETA: 0s - loss: 0.0707 - mse: 0.0707
- rmse: 0.2659 - mae: 0.1487 - mape: 55.4671

Epoch 2: val_mse improved from 0.10324 to 0.03743, saving model to
model_weights/20221013-221702_cnn_best_weights.hdf5

348/348 [=====] - 34s 98ms/step - loss: 0.0707 - mse:
0.0707 - rmse: 0.2658 - mae: 0.1487 - mape: 55.4547 - val_loss: 0.0374 -
val_mse: 0.0374 - val_rmse: 0.1935 - val_mae: 0.1163 - val_mape: 57.8829

Epoch 3/20

347/348 [=====>.] - ETA: 0s - loss: 0.0309 - mse: 0.0309
- rmse: 0.1759 - mae: 0.1046 - mape: 39.8158

Epoch 3: val_mse improved from 0.03743 to 0.02378, saving model to
model_weights/20221013-221702_cnn_best_weights.hdf5

348/348 [=====] - 33s 95ms/step - loss: 0.0309 - mse:
0.0309 - rmse: 0.1758 - mae: 0.1046 - mape: 39.8099 - val_loss: 0.0238 -
val_mse: 0.0238 - val_rmse: 0.1542 - val_mae: 0.1036 - val_mape: 49.5065

Epoch 4/20

347/348 [=====>.] - ETA: 0s - loss: 0.0204 - mse: 0.0204
- rmse: 0.1430 - mae: 0.0896 - mape: 44.9386

Epoch 4: val_mse improved from 0.02378 to 0.01643, saving model to
model_weights/20221013-221702_cnn_best_weights.hdf5

348/348 [=====] - 34s 99ms/step - loss: 0.0204 - mse:
0.0204 - rmse: 0.1430 - mae: 0.0896 - mape: 44.9378 - val_loss: 0.0164 -
val_mse: 0.0164 - val_rmse: 0.1282 - val_mae: 0.0843 - val_mape: 40.7541

Epoch 5/20

347/348 [=====>.] - ETA: 0s - loss: 0.0174 - mse: 0.0174
- rmse: 0.1321 - mae: 0.0827 - mape: 52.0188

Epoch 5: val_mse improved from 0.01643 to 0.01280, saving model to
model_weights/20221013-221702_cnn_best_weights.hdf5

348/348 [=====] - 35s 101ms/step - loss: 0.0174 - mse:

0.0174 - rmse: 0.1320 - mae: 0.0827 - mape: 51.9978 - val_loss: 0.0128 - val_mse: 0.0128 - val_rmse: 0.1131 - val_mae: 0.0729 - val_mape: 40.1540

Epoch 6/20

347/348 [=====>.] - ETA: 0s - loss: 0.0149 - mse: 0.0149 - rmse: 0.1219 - mae: 0.0759 - mape: 40.7168

Epoch 6: val_mse improved from 0.01280 to 0.01113, saving model to model_weights/20221013-221702_cnn_best_weights.hdf5

348/348 [=====] - 36s 104ms/step - loss: 0.0148 - mse: 0.0148 - rmse: 0.1219 - mae: 0.0759 - mape: 40.7016 - val_loss: 0.0111 - val_mse: 0.0111 - val_rmse: 0.1055 - val_mae: 0.0702 - val_mape: 33.6337

Epoch 7/20

347/348 [=====>.] - ETA: 0s - loss: 0.0118 - mse: 0.0118 - rmse: 0.1084 - mae: 0.0680 - mape: 32.6481

Epoch 7: val_mse improved from 0.01113 to 0.01086, saving model to model_weights/20221013-221702_cnn_best_weights.hdf5

348/348 [=====] - 36s 104ms/step - loss: 0.0118 - mse: 0.0118 - rmse: 0.1084 - mae: 0.0680 - mape: 32.6601 - val_loss: 0.0109 - val_mse: 0.0109 - val_rmse: 0.1042 - val_mae: 0.0691 - val_mape: 33.7990

Epoch 8/20

347/348 [=====>.] - ETA: 0s - loss: 0.0109 - mse: 0.0109 - rmse: 0.1044 - mae: 0.0654 - mape: 36.7081

Epoch 8: val_mse did not improve from 0.01086

348/348 [=====] - 34s 97ms/step - loss: 0.0109 - mse: 0.0109 - rmse: 0.1044 - mae: 0.0654 - mape: 36.6934 - val_loss: 0.0109 - val_mse: 0.0109 - val_rmse: 0.1043 - val_mae: 0.0693 - val_mape: 31.4723

Epoch 9/20

347/348 [=====>.] - ETA: 0s - loss: 0.0097 - mse: 0.0097 - rmse: 0.0984 - mae: 0.0605 - mape: 31.5169

Epoch 9: val_mse did not improve from 0.01086

348/348 [=====] - 37s 105ms/step - loss: 0.0097 - mse: 0.0097 - rmse: 0.0984 - mae: 0.0605 - mape: 31.5105 - val_loss: 0.0116 - val_mse: 0.0116 - val_rmse: 0.1077 - val_mae: 0.0677 - val_mape: 34.1491

Epoch 10/20

347/348 [=====>.] - ETA: 0s - loss: 0.0094 - mse: 0.0094 - rmse: 0.0970 - mae: 0.0597 - mape: 29.1988

Epoch 10: val_mse did not improve from 0.01086

348/348 [=====] - 38s 108ms/step - loss: 0.0094 - mse: 0.0094 - rmse: 0.0970 - mae: 0.0597 - mape: 29.1928 - val_loss: 0.0143 - val_mse: 0.0143 - val_rmse: 0.1195 - val_mae: 0.0790 - val_mape: 36.7539

Epoch 11/20

347/348 [=====>.] - ETA: 0s - loss: 0.0089 - mse: 0.0089 - rmse: 0.0941 - mae: 0.0577 - mape: 28.1576

Epoch 11: val_mse improved from 0.01086 to 0.00949, saving model to model_weights/20221013-221702_cnn_best_weights.hdf5

348/348 [=====] - 39s 112ms/step - loss: 0.0089 - mse: 0.0089 - rmse: 0.0941 - mae: 0.0577 - mape: 28.1647 - val_loss: 0.0095 - val_mse: 0.0095 - val_rmse: 0.0974 - val_mae: 0.0650 - val_mape: 30.9810

Epoch 12/20

347/348 [=====>.] - ETA: 0s - loss: 0.0077 - mse: 0.0077
- rmse: 0.0877 - mae: 0.0540 - mape: 27.0972
Epoch 12: val_mse did not improve from 0.00949
348/348 [=====] - 39s 113ms/step - loss: 0.0077 - mse:
0.0077 - rmse: 0.0877 - mae: 0.0540 - mape: 27.0872 - val_loss: 0.0130 -
val_mse: 0.0130 - val_rmse: 0.1141 - val_mae: 0.0751 - val_mape: 30.6593
Epoch 13/20
347/348 [=====>.] - ETA: 0s - loss: 0.0082 - mse: 0.0082
- rmse: 0.0907 - mae: 0.0557 - mape: 35.7169
Epoch 13: val_mse did not improve from 0.00949
348/348 [=====] - 39s 111ms/step - loss: 0.0082 - mse:
0.0082 - rmse: 0.0907 - mae: 0.0557 - mape: 35.7048 - val_loss: 0.0128 -
val_mse: 0.0128 - val_rmse: 0.1130 - val_mae: 0.0713 - val_mape: 37.3328
Epoch 14/20
347/348 [=====>.] - ETA: 0s - loss: 0.0073 - mse: 0.0073
- rmse: 0.0855 - mae: 0.0525 - mape: 31.8779
Epoch 14: val_mse improved from 0.00949 to 0.00895, saving model to
model_weights/20221013-221702_cnn_best_weights.hdf5
348/348 [=====] - 41s 117ms/step - loss: 0.0073 - mse:
0.0073 - rmse: 0.0855 - mae: 0.0524 - mape: 31.8660 - val_loss: 0.0090 -
val_mse: 0.0090 - val_rmse: 0.0946 - val_mae: 0.0684 - val_mape: 25.2155
Epoch 15/20
347/348 [=====>.] - ETA: 0s - loss: 0.0110 - mse: 0.0110
- rmse: 0.1049 - mae: 0.0642 - mape: 47.2699
Epoch 15: val_mse improved from 0.00895 to 0.00804, saving model to
model_weights/20221013-221702_cnn_best_weights.hdf5
348/348 [=====] - 38s 110ms/step - loss: 0.0110 - mse:
0.0110 - rmse: 0.1049 - mae: 0.0642 - mape: 47.2535 - val_loss: 0.0080 -
val_mse: 0.0080 - val_rmse: 0.0897 - val_mae: 0.0605 - val_mape: 28.6601
Epoch 16/20
347/348 [=====>.] - ETA: 0s - loss: 0.0062 - mse: 0.0062
- rmse: 0.0786 - mae: 0.0471 - mape: 22.1928
Epoch 16: val_mse improved from 0.00804 to 0.00688, saving model to
model_weights/20221013-221702_cnn_best_weights.hdf5
348/348 [=====] - 39s 112ms/step - loss: 0.0062 - mse:
0.0062 - rmse: 0.0786 - mae: 0.0471 - mape: 22.1864 - val_loss: 0.0069 -
val_mse: 0.0069 - val_rmse: 0.0830 - val_mae: 0.0567 - val_mape: 28.5051
Epoch 17/20
347/348 [=====>.] - ETA: 0s - loss: 0.0073 - mse: 0.0073
- rmse: 0.0852 - mae: 0.0521 - mape: 24.7771
Epoch 17: val_mse did not improve from 0.00688
348/348 [=====] - 39s 112ms/step - loss: 0.0073 - mse:
0.0073 - rmse: 0.0852 - mae: 0.0521 - mape: 24.7728 - val_loss: 0.0070 -
val_mse: 0.0070 - val_rmse: 0.0835 - val_mae: 0.0518 - val_mape: 32.8219
Epoch 18/20
347/348 [=====>.] - ETA: 0s - loss: 0.0059 - mse: 0.0059
- rmse: 0.0765 - mae: 0.0465 - mape: 22.7724
Epoch 18: val_mse improved from 0.00688 to 0.00521, saving model to

```

model_weights/20221013-221702_cnn_best_weights.hdf5
348/348 [=====] - 39s 112ms/step - loss: 0.0059 - mse:
0.0059 - rmse: 0.0765 - mae: 0.0465 - mape: 22.7665 - val_loss: 0.0052 -
val_mse: 0.0052 - val_rmse: 0.0722 - val_mae: 0.0467 - val_mape: 24.6541
Epoch 19/20
347/348 [=====>.] - ETA: 0s - loss: 0.0053 - mse: 0.0053
- rmse: 0.0728 - mae: 0.0447 - mape: 24.0311
Epoch 19: val_mse did not improve from 0.00521
348/348 [=====] - 38s 109ms/step - loss: 0.0053 - mse:
0.0053 - rmse: 0.0728 - mae: 0.0447 - mape: 24.0230 - val_loss: 0.0070 -
val_mse: 0.0070 - val_rmse: 0.0834 - val_mae: 0.0549 - val_mape: 28.3918
Epoch 20/20
347/348 [=====>.] - ETA: 0s - loss: 0.0053 - mse: 0.0053
- rmse: 0.0726 - mae: 0.0437 - mape: 21.7055
Epoch 20: val_mse did not improve from 0.00521
348/348 [=====] - 35s 100ms/step - loss: 0.0053 - mse:
0.0053 - rmse: 0.0726 - mae: 0.0437 - mape: 21.6978 - val_loss: 0.0063 -
val_mse: 0.0063 - val_rmse: 0.0796 - val_mae: 0.0527 - val_mape: 30.9943

```

0.4 Evaluate model

```
[ ]: best_model = models.load_model(f'model_weights/20221013-221702_cnn_best_weights.
    ↪hdf5')
```

```
[ ]: def evaluate_model(model, test_values, steps):
    score = model.evaluate(test_values, steps=steps)
    return score
```

```
[ ]: test_loss, test_mse, test_rmse, test_mae, test_mape = ↪
    ↪evaluate_model(best_model, test_ds, steps=np.ceil(test_size / BATCH_SIZE))
```

```

2022-10-13 22:34:16.889444: I tensorflow/stream_executor/cuda/cuda_dnn.cc:384]
Loaded cuDNN version 8100
2022-10-13 22:34:18.077149: I
tensorflow/core/platform/default/subprocess.cc:304] Start cannot spawn child
process: No such file or directory
2022-10-13 22:34:18.077988: I
tensorflow/core/platform/default/subprocess.cc:304] Start cannot spawn child
process: No such file or directory
2022-10-13 22:34:18.078006: W tensorflow/stream_executor/gpu/asm_compiler.cc:80]
Couldn't get ptxas version string: INTERNAL: Couldn't invoke ptxas --version
2022-10-13 22:34:18.078825: I
tensorflow/core/platform/default/subprocess.cc:304] Start cannot spawn child
process: No such file or directory
2022-10-13 22:34:18.078873: W
tensorflow/stream_executor/gpu/redzone_allocator.cc:314] INTERNAL: Failed to
launch ptxas
Relying on driver to perform ptx compilation.

```

Modify \$PATH to customize ptxas location.

This message will be only logged once.

```
100/100 [=====] - 9s 59ms/step - loss: 0.0052 - mse:
0.0052 - rmse: 0.0719 - mae: 0.0469 - mape: 15.1341
```

```
[ ]: predictions = best_model.predict(test_ds, steps=np.ceil(test_size / BATCH_SIZE))
```

```
100/100 [=====] - 6s 60ms/step
```

```
[ ]: #small_test_ds = next(iter(test_ds))
```

```
[ ]: for image, stage_discharge in test_ds.take(1):
    predictions = best_model.predict(x=image)

    stage_discharge_test_values = stage_discharge[:2].numpy()
    predictions_values = predictions[:2]

    diff = predictions_values.flatten() - stage_discharge_test_values.
    ↪flatten()
    percentDiff = (diff / stage_discharge_test_values.flatten()) * 100
    absPercentDiff = np.abs(percentDiff)
    # compute the mean and standard deviation of the absolute percentage
    # difference
    mean = np.mean(absPercentDiff)
    std = np.std(absPercentDiff)
    # finally, show some statistics on our model
    print(mean)
    print(std)

    for i in range(len(stage_discharge_test_values)):
        print(f"pred stage: {scaler.
    ↪inverse_transform(predictions_values)[i][0]}, actual stage: {scaler.
    ↪inverse_transform(stage_discharge_test_values)[i][0]}")
        print(f"pred discharge: {scaler.
    ↪inverse_transform(predictions_values)[i][1]}, actual discharge: {scaler.
    ↪inverse_transform(stage_discharge_test_values)[i][1]}")
```

```
1/1 [=====] - 0s 64ms/step
3.6887279750450386
3.6212939608575954
pred stage: 2.1561872959136963, actual stage: 2.15
pred discharge: 154.75941467285156, actual discharge: 155.0
pred stage: 4.178528308868408, actual stage: 4.12
pred discharge: 2891.14892578125, actual discharge: 2730.0
```

```
[ ]:
```

0.5 Visualize layers

```
[ ]: layer_outputs = [layer.output for layer in best_model.layers[:12]]  
# Extracts the outputs of the top 12 layers  
activation_model = models.Model(inputs=best_model.input, outputs=layer_outputs)␣  
→# Creates a model that will return these outputs, given the model input
```

```
[ ]: activations = activation_model.predict(test_ds.take(1))
```

1/1 [=====] - 0s 202ms/step

```
[ ]: import matplotlib.pyplot as plt  
  
layer_names = []  
for layer in best_model.layers[:12]:  
    layer_names.append(layer.name) # Names of the layers, so you can have them␣  
→as part of your plot  
  
images_per_row = 16  
  
for layer_name, layer_activation in zip(layer_names, activations): # Displays␣  
→the feature maps  
    n_features = layer_activation.shape[-1] # Number of features in the feature␣  
→map  
    size = layer_activation.shape[1] #The feature map has shape (1, size, size,␣  
→n_features).  
    n_cols = n_features // images_per_row # Tiles the activation channels in␣  
→this matrix  
    display_grid = np.zeros((size * n_cols, images_per_row * size))  
  
    print(layer_name)  
    if ("flatten" in layer_name): break  
  
    for col in range(n_cols): # Tiles each filter into a big horizontal grid  
        for row in range(images_per_row):  
            channel_image = layer_activation[0,  
                                                :, :,  
                                                col * images_per_row + row]  
            channel_image -= channel_image.mean() # Post-processes the feature␣  
→to make it visually palatable  
            channel_image /= channel_image.std()  
            channel_image *= 64  
            channel_image += 128  
            channel_image = np.clip(channel_image, 0, 255).astype('uint8')  
            display_grid[col * size : (col + 1) * size, # Displays the grid  
                          row * size : (row + 1) * size] = channel_image  
  
    scale = 1. / size
```

```
plt.figure(figsize=(scale * display_grid.shape[1],
                    scale * display_grid.shape[0]))
plt.title(layer_name)
plt.grid(False)
plt.imshow(display_grid, aspect='auto', cmap='viridis')
```

```
conv2d_24
max_pooling2d_24
conv2d_25
max_pooling2d_25
conv2d_26
max_pooling2d_26
conv2d_27
max_pooling2d_27
flatten_6
```

```
/tmp/ipykernel_35667/2269795348.py:24: RuntimeWarning: invalid value encountered
in divide
```

```
channel_image /= channel_image.std()
```





