

cnn_v7

October 14, 2022

```
[ ]: %env LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$CONDA_PREFIX/lib/
```

```
env: LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$CONDA_PREFIX/lib/
```

```
[ ]: import os
print(os.environ["LD_LIBRARY_PATH"])
```

```
:/home/nkspartan/miniconda3/envs/tf-gpu/lib/:/home/nkspartan/miniconda3/envs/tf-gpu/lib/
```

```
[ ]: import tensorflow as tf
import numpy as np
import pandas as pd
import os
import keras

from keras import Sequential, models, Input
from keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, Dropout,
↳LeakyReLU, AveragePooling2D
from keras.optimizers import SGD, Adam
```

```
[ ]: from tensorflow.python.client import device_lib

#print(device_lib.list_local_devices())
print('Default GPU Device: {}'.format(tf.test.gpu_device_name()))
```

```
Default GPU Device: /device:GPU:0
```

```
2022-10-14 11:39:11.186822: I tensorflow/core/platform/cpu_feature_guard.cc:193]
This TensorFlow binary is optimized with oneAPI Deep Neural Network Library
(oneDNN) to use the following CPU instructions in performance-critical
operations: AVX2 FMA
```

```
To enable them in other operations, rebuild TensorFlow with the appropriate
compiler flags.
```

```
2022-10-14 11:39:11.209616: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
```

```
2022-10-14 11:39:11.261664: I
```

```

tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:39:11.261850: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:39:12.100458: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:39:12.101067: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:39:12.101224: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:39:12.101367: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:1616] Created device
/device:GPU:0 with 4023 MB memory: -> device: 0, name: NVIDIA GeForce RTX 2060,
pci bus id: 0000:08:00.0, compute capability: 7.5

```

0.1 Read the csv dataset to get the values for stage and discharge of the images

```

[ ]: df = pd.read_csv("../..//dataset/2012_2019_PlatteRiverWeir_features_merged_all.
↳csv")
df.head()

```

```

[ ]:
   Unnamed: 0  SensorTime  CaptureTime \
0           0  2012-06-09 13:15:00  2012-06-09T13:09:07
1           1  2012-06-09 13:15:00  2012-06-09T13:10:29
2           2  2012-06-09 13:45:00  2012-06-09T13:44:01
3           3  2012-06-09 14:45:00  2012-06-09T14:44:30
4           4  2012-06-09 15:45:00  2012-06-09T15:44:59

```

	Filename	Agency	SiteNumber	TimeZone	Stage	\
0	StateLineWeir_20120609_Farrell_001.jpg	USGS	6674500	MDT	2.99	
1	StateLineWeir_20120609_Farrell_002.jpg	USGS	6674500	MDT	2.99	
2	StateLineWeir_20120609_Farrell_003.jpg	USGS	6674500	MDT	2.96	
3	StateLineWeir_20120609_Farrell_004.jpg	USGS	6674500	MDT	2.94	
4	StateLineWeir_20120609_Farrell_005.jpg	USGS	6674500	MDT	2.94	

	Discharge	CalcTimestamp	...	WeirPt2X	WeirPt2Y	WwRawLineMin	\
0	916.0	2020-03-11T16:58:28	...	-1	-1	0.0	
1	916.0	2020-03-11T16:58:33	...	-1	-1	0.0	

2	873.0	2020-03-11T16:58:40	...	-1	-1	0.0
3	846.0	2020-03-11T16:58:47	...	-1	-1	0.0
4	846.0	2020-03-11T16:58:55	...	-1	-1	0.0

	WwRawLineMax	WwRawLineMean	WwRawLineSigma	WwCurveLineMin	\
0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	

	WwCurveLineMax	WwCurveLineMean	WwCurveLineSigma
0	0.0	0.0	0.0
1	0.0	0.0	0.0
2	0.0	0.0	0.0
3	0.0	0.0	0.0
4	0.0	0.0	0.0

[5 rows x 60 columns]

```
[ ]: df = df[["Filename", "Stage", "Discharge"]]
```

0.1.1 Scale the data

```
[ ]: from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
```

```
[ ]: df[["Stage", "Discharge"]] = scaler.fit_transform(df[["Stage", "Discharge"]])
df
```

```
[ ]:
      Filename      Stage  Discharge
0  StateLineWeir_20120609_Farrell_001.jpg  0.138117 -0.046094
1  StateLineWeir_20120609_Farrell_002.jpg  0.138117 -0.046094
2  StateLineWeir_20120609_Farrell_003.jpg  0.100875 -0.082160
3  StateLineWeir_20120609_Farrell_004.jpg  0.076046 -0.104807
4  StateLineWeir_20120609_Farrell_005.jpg  0.076046 -0.104807
...
42054  StateLineWeir_20191011_Farrell_409.jpg -0.420526 -0.450369
42055  StateLineWeir_20191011_Farrell_410.jpg -0.420526 -0.450369
42056  StateLineWeir_20191011_Farrell_411.jpg -0.420526 -0.450369
42057  StateLineWeir_20191011_Farrell_412.jpg -0.420526 -0.450369
42058  StateLineWeir_20191011_Farrell_413.jpg -0.420526 -0.450369
```

[42059 rows x 3 columns]

```
[ ]: from joblib import dump, load
dump(scaler, 'std_scaler.joblib', compress=True)
```

```
[ ]: ['std_scaler.joblib']
```

0.2 Create the dataset pipeline

```
[ ]: IMG_SIZE = 512
BATCH_SIZE = 32
```

```
[ ]: from glob import glob

def make_dataset(path, batch_size, df, seed=None):
    np.random.seed(seed)

    def parse_image(filename):
        image = tf.io.read_file(filename)
        image = tf.image.decode_jpeg(image, channels=3)
        #image = tf.image.resize(image, [IMG_SIZE, IMG_SIZE])
        image = tf.cast(image, tf.float32)
        image /= 255
        return image

    def configure_for_performance(ds):
        ds = ds.shuffle(buffer_size=100)
        ds = ds.batch(batch_size)
        ds = ds.repeat()
        ds = ds.prefetch(buffer_size=tf.data.experimental.AUTOTUNE)
        return ds

    filenames = glob(path + '/*')

    # make train, val and test splits of the dataset (70%, 10%, 20% split)
    split1 = int(0.7 * len(filenames))
    split2 = int(0.8 * len(filenames))

    np.random.shuffle(filenames)
    train_files = filenames[:split1] # up to split 1 (ex 70%)
    val_files = filenames[split1:split2] # from ex. 70% to 80%
    test_files = filenames[split2:] # from ex. 80% until the end

    # create stage values
    stage_train_values = [df[df.Filename == file.split('/')[0]].Stage.values for
    ↪file in train_files]
    stage_val_values = [df[df.Filename == file.split('/')[0]].Stage.values for
    ↪file in val_files]
```

```

stage_test_values = [df[df.Filename == file.split('/')[1]].Stage.values for
↪file in test_files]

# create discharge values
discharge_train_values = [df[df.Filename == file.split(
    '/')[-1]].Discharge.values for file in train_files]
discharge_val_values = [df[df.Filename == file.split(
    '/')[-1]].Discharge.values for file in val_files]
discharge_test_values = [df[df.Filename == file.split(
    '/')[-1]].Discharge.values for file in test_files]

# join stage and discharge values
stage_discharge_train_values = [[np.squeeze(s), np.squeeze(d)] for s, d in
↪zip(stage_train_values, discharge_train_values)]
stage_discharge_val_values = [[np.squeeze(s), np.squeeze(d)] for s, d in
↪zip(stage_val_values, discharge_val_values)]
stage_discharge_test_values = [[np.squeeze(s), np.squeeze(
    d)] for s, d in zip(stage_test_values, discharge_test_values)]

# create images dataset (train, val, test)
filenames_train_ds = tf.data.Dataset.from_tensor_slices(train_files)
filenames_val_ds = tf.data.Dataset.from_tensor_slices(val_files)
filenames_test_ds = tf.data.Dataset.from_tensor_slices(test_files)

images_train_ds = filenames_train_ds.map(parse_image, num_parallel_calls=5)
images_val_ds = filenames_val_ds.map(parse_image, num_parallel_calls=5)
images_test_ds = filenames_test_ds.map(parse_image, num_parallel_calls=5)

# create stage and discharge dataset (train, val, test)
stage_discharge_train_ds = tf.data.Dataset.
↪from_tensor_slices(stage_discharge_train_values)
stage_discharge_val_ds = tf.data.Dataset.
↪from_tensor_slices(stage_discharge_val_values)
stage_discharge_test_ds = tf.data.Dataset.from_tensor_slices(
    stage_discharge_test_values)

# create tensorflow dataset of images and values (train, val, test)
train_ds = tf.data.Dataset.zip((images_train_ds, stage_discharge_train_ds))
train_ds = configure_for_performance(train_ds)
val_ds = tf.data.Dataset.zip((images_val_ds, stage_discharge_val_ds))
val_ds = configure_for_performance(val_ds)
test_ds = tf.data.Dataset.zip((images_test_ds, stage_discharge_test_ds))
test_ds = configure_for_performance(test_ds)

return train_ds, len(train_files), val_ds, len(val_files), test_ds,
↪len(test_files)

```

```
[ ]: path = "../..../dataset/images"

train_ds, train_size, val_ds, val_size, test_ds, test_size = make_dataset(path,
↳ BATCH_SIZE, df, 0)
```

```
2022-10-14 11:47:25.673029: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:47:25.673317: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:47:25.673566: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:47:25.674117: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:47:25.674268: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:47:25.674407: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:47:25.674587: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:47:25.674729: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:47:25.674844: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:1616] Created device
/job:localhost/replica:0/task:0/device:GPU:0 with 4023 MB memory: -> device: 0,
name: NVIDIA GeForce RTX 2060, pci bus id: 0000:08:00.0, compute capability: 7.5
```

```
[ ]: input_shape = 0
output_shape = 0

for image, stage_discharge in train_ds.take(1):
    print(image.numpy().shape)
```

```
print(stage_discharge.numpy().shape)

input_shape = image.numpy().shape[1:]
output_shape = stage_discharge.numpy().shape[1:]
```

```
(32, 512, 512, 3)
(32, 2)
```

```
[ ]: print(input_shape)
      print(output_shape)
```

```
(512, 512, 3)
(2,)
```

0.3 Create model

```
[ ]: def create_model(input_shape, output_shape):
      model = Sequential()

      model.add(Input(shape=input_shape))

      model.add(Conv2D(64, kernel_size=(4, 4), strides=(2, 2), padding='same',
      ↪activation=LeakyReLU()))
      model.add(MaxPooling2D(pool_size=(4, 4)))

      model.add(Conv2D(64, kernel_size=(4, 4), activation=LeakyReLU(),
      ↪padding='same'))
      model.add(MaxPooling2D(pool_size=(2, 2)))

      model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', padding='same'))
      model.add(AveragePooling2D(pool_size=(2, 2)))

      model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
      model.add(AveragePooling2D(pool_size=(2, 2)))

      model.add(Flatten())
      model.add(Dense(64, activation='relu'))
      model.add(Dense(64, activation='relu'))
      model.add(Dense(32, activation='relu'))
      model.add(Dense(32, activation='relu'))
      model.add(Dense(output_shape, activation='linear')) # linear regression
      ↪output layer

      return model
```

```
[ ]: model = create_model(input_shape, output_shape[0])
```

```
[ ]: model.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 256, 256, 64)	3136
max_pooling2d (MaxPooling2D)	(None, 64, 64, 64)	0
conv2d_1 (Conv2D)	(None, 64, 64, 64)	65600
max_pooling2d_1 (MaxPooling2D)	(None, 32, 32, 64)	0
conv2d_2 (Conv2D)	(None, 32, 32, 32)	18464
average_pooling2d (AveragePooling2D)	(None, 16, 16, 32)	0
conv2d_3 (Conv2D)	(None, 14, 14, 32)	9248
average_pooling2d_1 (AveragePooling2D)	(None, 7, 7, 32)	0
flatten (Flatten)	(None, 1568)	0
dense (Dense)	(None, 64)	100416
dense_1 (Dense)	(None, 64)	4160
dense_2 (Dense)	(None, 32)	2080
dense_3 (Dense)	(None, 32)	1056
dense_4 (Dense)	(None, 2)	66

=====
Total params: 204,226
Trainable params: 204,226
Non-trainable params: 0
=====

```
[ ]: def compile_model(loss_func, optimizer, metrics=["accuracy"]):  
      model.compile(loss=loss_func, optimizer=optimizer, metrics=metrics)
```



```
[ ]: sgd = SGD(learning_rate=0.01, decay=1e-4, momentum=0.9, nesterov=True)
adam = Adam(learning_rate=1e-3, decay=1e-3 / 100)

compile_model('mse', adam, [
    'mse', tf.keras.metrics.RootMeanSquaredError(name='rmse'), 'mae',
    ↪ 'mape'])

[ ]: def fit_model(training_values, validation_values=None, epochs=10, steps=32,
    ↪ val_steps=32, callbacks=[]):
    return model.fit(training_values, validation_data=validation_values,
    ↪ epochs=epochs, steps_per_epoch=steps, validation_steps=val_steps,
    ↪ callbacks=callbacks)

[ ]: import datetime

date_actual = datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
log_dir = "logs/fit/" + date_actual
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir,
    ↪ histogram_freq=1)

checkpoint_callback = tf.keras.callbacks.
    ↪ ModelCheckpoint(filepath=f"model_weights/{date_actual}_cnn_best_weights.
    ↪ hdf5",
                                monitor='val_mse',
                                verbose=1,
                                save_best_only=True)

[ ]: # batch_size = 0 because we already have batch size in tf dataset
history = fit_model(train_ds, val_ds, epochs=25, steps=np.ceil(train_size /
    ↪ BATCH_SIZE), val_steps=np.ceil(val_size / BATCH_SIZE),
    ↪ callbacks=[tensorboard_callback, checkpoint_callback])
```

Epoch 1/25

```
2022-10-14 11:47:51.620111: I tensorflow/stream_executor/cuda/cuda_dnn.cc:384]
Loaded cuDNN version 8100
2022-10-14 11:47:52.860907: I
tensorflow/core/platform/default/subprocess.cc:304] Start cannot spawn child
process: No such file or directory
2022-10-14 11:47:52.861741: I
tensorflow/core/platform/default/subprocess.cc:304] Start cannot spawn child
process: No such file or directory
2022-10-14 11:47:52.861762: W tensorflow/stream_executor/gpu/asm_compiler.cc:80]
Couldn't get ptxas version string: INTERNAL: Couldn't invoke ptxas --version
2022-10-14 11:47:52.862848: I
tensorflow/core/platform/default/subprocess.cc:304] Start cannot spawn child
process: No such file or directory
2022-10-14 11:47:52.862909: W
```

tensorflow/stream_executor/gpu/redzone_allocator.cc:314] INTERNAL: Failed to launch ptxas

Relying on driver to perform ptx compilation.

Modify \$PATH to customize ptxas location.

This message will be only logged once.

921/921 [=====] - ETA: 0s - loss: 0.1460 - mse: 0.1460
- rmse: 0.3822 - mae: 0.2167 - mape: 86.9484

Epoch 1: val_mse improved from inf to 0.03423, saving model to
model_weights/20221014-114748_cnn_best_weights.hdf5

921/921 [=====] - 105s 109ms/step - loss: 0.1460 - mse:
0.1460 - rmse: 0.3822 - mae: 0.2167 - mape: 86.9484 - val_loss: 0.0342 -
val_mse: 0.0342 - val_rmse: 0.1850 - val_mae: 0.1270 - val_mape: 40.8599

Epoch 2/25

920/921 [=====>.] - ETA: 0s - loss: 0.0285 - mse: 0.0285
- rmse: 0.1687 - mae: 0.1139 - mape: 52.6991

Epoch 2: val_mse improved from 0.03423 to 0.03168, saving model to
model_weights/20221014-114748_cnn_best_weights.hdf5

921/921 [=====] - 110s 119ms/step - loss: 0.0285 - mse:
0.0285 - rmse: 0.1687 - mae: 0.1139 - mape: 52.6979 - val_loss: 0.0317 -
val_mse: 0.0317 - val_rmse: 0.1780 - val_mae: 0.1230 - val_mape: 33.6909

Epoch 3/25

920/921 [=====>.] - ETA: 0s - loss: 0.0159 - mse: 0.0159
- rmse: 0.1260 - mae: 0.0856 - mape: 34.2474

Epoch 3: val_mse improved from 0.03168 to 0.02314, saving model to
model_weights/20221014-114748_cnn_best_weights.hdf5

921/921 [=====] - 111s 120ms/step - loss: 0.0159 - mse:
0.0159 - rmse: 0.1260 - mae: 0.0856 - mape: 34.2464 - val_loss: 0.0231 -
val_mse: 0.0231 - val_rmse: 0.1521 - val_mae: 0.0996 - val_mape: 26.9171

Epoch 4/25

920/921 [=====>.] - ETA: 0s - loss: 0.0125 - mse: 0.0125
- rmse: 0.1119 - mae: 0.0751 - mape: 38.1384

Epoch 4: val_mse improved from 0.02314 to 0.01050, saving model to
model_weights/20221014-114748_cnn_best_weights.hdf5

921/921 [=====] - 110s 120ms/step - loss: 0.0125 - mse:
0.0125 - rmse: 0.1119 - mae: 0.0751 - mape: 38.1372 - val_loss: 0.0105 -
val_mse: 0.0105 - val_rmse: 0.1025 - val_mae: 0.0684 - val_mape: 22.3800

Epoch 5/25

920/921 [=====>.] - ETA: 0s - loss: 0.0109 - mse: 0.0109
- rmse: 0.1043 - mae: 0.0695 - mape: 29.6158

Epoch 5: val_mse did not improve from 0.01050

921/921 [=====] - 109s 118ms/step - loss: 0.0109 - mse:
0.0109 - rmse: 0.1043 - mae: 0.0695 - mape: 29.6150 - val_loss: 0.0119 -
val_mse: 0.0119 - val_rmse: 0.1089 - val_mae: 0.0718 - val_mape: 21.1961

Epoch 6/25

920/921 [=====>.] - ETA: 0s - loss: 0.0087 - mse: 0.0087
- rmse: 0.0932 - mae: 0.0615 - mape: 27.1488

Epoch 6: val_mse did not improve from 0.01050

921/921 [=====] - 109s 118ms/step - loss: 0.0087 - mse: 0.0087 - rmse: 0.0932 - mae: 0.0615 - mape: 27.1481 - val_loss: 0.0141 - val_mse: 0.0141 - val_rmse: 0.1189 - val_mae: 0.0727 - val_mape: 17.5537
Epoch 7/25
920/921 [=====>.] - ETA: 0s - loss: 0.0076 - mse: 0.0076 - rmse: 0.0874 - mae: 0.0567 - mape: 27.3970
Epoch 7: val_mse improved from 0.01050 to 0.00580, saving model to model_weights/20221014-114748_cnn_best_weights.hdf5
921/921 [=====] - 108s 118ms/step - loss: 0.0076 - mse: 0.0076 - rmse: 0.0874 - mae: 0.0567 - mape: 27.3963 - val_loss: 0.0058 - val_mse: 0.0058 - val_rmse: 0.0762 - val_mae: 0.0499 - val_mape: 14.4002
Epoch 8/25
920/921 [=====>.] - ETA: 0s - loss: 0.0071 - mse: 0.0071 - rmse: 0.0841 - mae: 0.0540 - mape: 28.0689
Epoch 8: val_mse did not improve from 0.00580
921/921 [=====] - 109s 119ms/step - loss: 0.0071 - mse: 0.0071 - rmse: 0.0841 - mae: 0.0540 - mape: 28.0681 - val_loss: 0.0149 - val_mse: 0.0149 - val_rmse: 0.1221 - val_mae: 0.0708 - val_mape: 18.1489
Epoch 9/25
920/921 [=====>.] - ETA: 0s - loss: 0.0068 - mse: 0.0068 - rmse: 0.0824 - mae: 0.0523 - mape: 24.4000
Epoch 9: val_mse improved from 0.00580 to 0.00412, saving model to model_weights/20221014-114748_cnn_best_weights.hdf5
921/921 [=====] - 108s 117ms/step - loss: 0.0068 - mse: 0.0068 - rmse: 0.0824 - mae: 0.0523 - mape: 24.3992 - val_loss: 0.0041 - val_mse: 0.0041 - val_rmse: 0.0642 - val_mae: 0.0425 - val_mape: 15.7825
Epoch 10/25
920/921 [=====>.] - ETA: 0s - loss: 0.0058 - mse: 0.0058 - rmse: 0.0763 - mae: 0.0491 - mape: 22.7052
Epoch 10: val_mse did not improve from 0.00412
921/921 [=====] - 108s 117ms/step - loss: 0.0058 - mse: 0.0058 - rmse: 0.0763 - mae: 0.0491 - mape: 22.7046 - val_loss: 0.0059 - val_mse: 0.0059 - val_rmse: 0.0771 - val_mae: 0.0525 - val_mape: 18.6875
Epoch 11/25
920/921 [=====>.] - ETA: 0s - loss: 0.0052 - mse: 0.0052 - rmse: 0.0721 - mae: 0.0467 - mape: 22.0623
Epoch 11: val_mse did not improve from 0.00412
921/921 [=====] - 108s 117ms/step - loss: 0.0052 - mse: 0.0052 - rmse: 0.0721 - mae: 0.0467 - mape: 22.0616 - val_loss: 0.0053 - val_mse: 0.0053 - val_rmse: 0.0726 - val_mae: 0.0452 - val_mape: 13.8431
Epoch 12/25
920/921 [=====>.] - ETA: 0s - loss: 0.0046 - mse: 0.0046 - rmse: 0.0680 - mae: 0.0432 - mape: 21.0112
Epoch 12: val_mse did not improve from 0.00412
921/921 [=====] - 106s 115ms/step - loss: 0.0046 - mse: 0.0046 - rmse: 0.0680 - mae: 0.0432 - mape: 21.0106 - val_loss: 0.0049 - val_mse: 0.0049 - val_rmse: 0.0698 - val_mae: 0.0453 - val_mape: 14.0386
Epoch 13/25

920/921 [=====>.] - ETA: 0s - loss: 0.0043 - mse: 0.0043
- rmse: 0.0653 - mae: 0.0419 - mape: 20.2350
Epoch 13: val_mse did not improve from 0.00412
921/921 [=====] - 100s 109ms/step - loss: 0.0043 - mse:
0.0043 - rmse: 0.0653 - mae: 0.0419 - mape: 20.2343 - val_loss: 0.0042 -
val_mse: 0.0042 - val_rmse: 0.0646 - val_mae: 0.0393 - val_mape: 12.5412
Epoch 14/25
920/921 [=====>.] - ETA: 0s - loss: 0.0051 - mse: 0.0051
- rmse: 0.0712 - mae: 0.0437 - mape: 18.7385
Epoch 14: val_mse did not improve from 0.00412
921/921 [=====] - 107s 116ms/step - loss: 0.0051 - mse:
0.0051 - rmse: 0.0712 - mae: 0.0437 - mape: 18.7380 - val_loss: 0.0062 -
val_mse: 0.0062 - val_rmse: 0.0789 - val_mae: 0.0558 - val_mape: 18.8260
Epoch 15/25
920/921 [=====>.] - ETA: 0s - loss: 0.0038 - mse: 0.0038
- rmse: 0.0613 - mae: 0.0393 - mape: 22.1371
Epoch 15: val_mse improved from 0.00412 to 0.00397, saving model to
model_weights/20221014-114748_cnn_best_weights.hdf5
921/921 [=====] - 108s 117ms/step - loss: 0.0038 - mse:
0.0038 - rmse: 0.0613 - mae: 0.0393 - mape: 22.1364 - val_loss: 0.0040 -
val_mse: 0.0040 - val_rmse: 0.0630 - val_mae: 0.0400 - val_mape: 15.2230
Epoch 16/25
920/921 [=====>.] - ETA: 0s - loss: 0.0031 - mse: 0.0031
- rmse: 0.0556 - mae: 0.0357 - mape: 18.3191
Epoch 16: val_mse improved from 0.00397 to 0.00314, saving model to
model_weights/20221014-114748_cnn_best_weights.hdf5
921/921 [=====] - 107s 116ms/step - loss: 0.0031 - mse:
0.0031 - rmse: 0.0556 - mae: 0.0357 - mape: 18.3187 - val_loss: 0.0031 -
val_mse: 0.0031 - val_rmse: 0.0560 - val_mae: 0.0375 - val_mape: 11.7202
Epoch 17/25
920/921 [=====>.] - ETA: 0s - loss: 0.0036 - mse: 0.0036
- rmse: 0.0600 - mae: 0.0374 - mape: 16.9154
Epoch 17: val_mse improved from 0.00314 to 0.00313, saving model to
model_weights/20221014-114748_cnn_best_weights.hdf5
921/921 [=====] - 107s 117ms/step - loss: 0.0036 - mse:
0.0036 - rmse: 0.0600 - mae: 0.0374 - mape: 16.9149 - val_loss: 0.0031 -
val_mse: 0.0031 - val_rmse: 0.0559 - val_mae: 0.0357 - val_mape: 13.3308
Epoch 18/25
920/921 [=====>.] - ETA: 0s - loss: 0.0028 - mse: 0.0028
- rmse: 0.0526 - mae: 0.0338 - mape: 16.8378
Epoch 18: val_mse improved from 0.00313 to 0.00271, saving model to
model_weights/20221014-114748_cnn_best_weights.hdf5
921/921 [=====] - 107s 116ms/step - loss: 0.0028 - mse:
0.0028 - rmse: 0.0526 - mae: 0.0338 - mape: 16.8374 - val_loss: 0.0027 -
val_mse: 0.0027 - val_rmse: 0.0520 - val_mae: 0.0339 - val_mape: 12.4455
Epoch 19/25
920/921 [=====>.] - ETA: 0s - loss: 0.0028 - mse: 0.0028
- rmse: 0.0531 - mae: 0.0337 - mape: 17.3883

Epoch 19: val_mse did not improve from 0.00271
921/921 [=====] - 105s 115ms/step - loss: 0.0028 - mse: 0.0028 - rmse: 0.0531 - mae: 0.0337 - mape: 17.3879 - val_loss: 0.0044 - val_mse: 0.0044 - val_rmse: 0.0665 - val_mae: 0.0454 - val_mape: 16.2450

Epoch 20/25
920/921 [=====>.] - ETA: 0s - loss: 0.0026 - mse: 0.0026 - rmse: 0.0506 - mae: 0.0323 - mape: 19.3537

Epoch 20: val_mse did not improve from 0.00271
921/921 [=====] - 110s 119ms/step - loss: 0.0026 - mse: 0.0026 - rmse: 0.0506 - mae: 0.0323 - mape: 19.3531 - val_loss: 0.0036 - val_mse: 0.0036 - val_rmse: 0.0599 - val_mae: 0.0378 - val_mape: 13.7792

Epoch 21/25
920/921 [=====>.] - ETA: 0s - loss: 0.0028 - mse: 0.0028 - rmse: 0.0532 - mae: 0.0337 - mape: 15.7936

Epoch 21: val_mse did not improve from 0.00271
921/921 [=====] - 109s 118ms/step - loss: 0.0028 - mse: 0.0028 - rmse: 0.0532 - mae: 0.0337 - mape: 15.7930 - val_loss: 0.0028 - val_mse: 0.0028 - val_rmse: 0.0532 - val_mae: 0.0334 - val_mape: 10.8679

Epoch 22/25
920/921 [=====>.] - ETA: 0s - loss: 0.0026 - mse: 0.0026 - rmse: 0.0510 - mae: 0.0317 - mape: 16.7444

Epoch 22: val_mse improved from 0.00271 to 0.00269, saving model to model_weights/20221014-114748_cnn_best_weights.hdf5
921/921 [=====] - 110s 119ms/step - loss: 0.0026 - mse: 0.0026 - rmse: 0.0510 - mae: 0.0317 - mape: 16.7439 - val_loss: 0.0027 - val_mse: 0.0027 - val_rmse: 0.0519 - val_mae: 0.0333 - val_mape: 11.0753

Epoch 23/25
920/921 [=====>.] - ETA: 0s - loss: 0.0024 - mse: 0.0024 - rmse: 0.0490 - mae: 0.0309 - mape: 15.0125

Epoch 23: val_mse did not improve from 0.00269
921/921 [=====] - 104s 113ms/step - loss: 0.0024 - mse: 0.0024 - rmse: 0.0490 - mae: 0.0309 - mape: 15.0121 - val_loss: 0.0074 - val_mse: 0.0074 - val_rmse: 0.0862 - val_mae: 0.0431 - val_mape: 10.6888

Epoch 24/25
920/921 [=====>.] - ETA: 0s - loss: 0.0094 - mse: 0.0094 - rmse: 0.0970 - mae: 0.0517 - mape: 24.8800

Epoch 24: val_mse did not improve from 0.00269
921/921 [=====] - 103s 112ms/step - loss: 0.0094 - mse: 0.0094 - rmse: 0.0970 - mae: 0.0517 - mape: 24.8794 - val_loss: 0.0031 - val_mse: 0.0031 - val_rmse: 0.0555 - val_mae: 0.0361 - val_mape: 12.2362

Epoch 25/25
920/921 [=====>.] - ETA: 0s - loss: 0.0025 - mse: 0.0025 - rmse: 0.0496 - mae: 0.0316 - mape: 16.3576

Epoch 25: val_mse did not improve from 0.00269
921/921 [=====] - 109s 119ms/step - loss: 0.0025 - mse: 0.0025 - rmse: 0.0496 - mae: 0.0316 - mape: 16.3573 - val_loss: 0.0027 - val_mse: 0.0027 - val_rmse: 0.0519 - val_mae: 0.0347 - val_mape: 10.8955

0.4 Evaluate model

```
[ ]: print(date_actual)
```

20221014-114748

```
[ ]: best_model = models.load_model(f'model_weights/{date_actual}_cnn_best_weights.  
    ↪hdf5')
```

```
[ ]: def evaluate_model(model, test_values, steps):  
    score = model.evaluate(test_values, steps=steps)  
    return score
```

```
[ ]: test_loss, test_mse, test_rmse, test_mae, test_mape = ↪  
    ↪evaluate_model(best_model, test_ds, steps=np.ceil(test_size / BATCH_SIZE))
```

263/263 [=====] - 17s 64ms/step - loss: 0.0056 - mse:
0.0056 - rmse: 0.0746 - mae: 0.0339 - mape: 17.5864

```
[ ]: predictions = best_model.predict(test_ds, steps=np.ceil(test_size / BATCH_SIZE))
```

263/263 [=====] - 17s 64ms/step

```
[ ]: for image, stage_discharge in test_ds.take(1):  
    predictions = best_model.predict(x=image)  
  
    stage_discharge_test_values = stage_discharge[:2].numpy()  
    predictions_values = predictions[:2]  
  
    diff = predictions_values.flatten() - stage_discharge_test_values.  
    ↪flatten()  
    percentDiff = (diff / stage_discharge_test_values.flatten()) * 100  
    absPercentDiff = np.abs(percentDiff)  
    # compute the mean and standard deviation of the absolute percentage  
    # difference  
    mean = np.mean(absPercentDiff)  
    std = np.std(absPercentDiff)  
    # finally, show some statistics on our model  
    print(mean)  
    print(std)  
  
    stage_discharge_test_values = stage_discharge[:10]  
    predictions_values = predictions[:10]  
  
    for i in range(len(stage_discharge_test_values.numpy())):  
        print(f"pred stage: {scaler.  
    ↪inverse_transform(predictions_values)[i][0]}, actual stage: {scaler.  
    ↪inverse_transform(stage_discharge_test_values)[i][0]}")
```

```

        print(f"pred discharge: {scaler.
↪inverse_transform(predictions_values)[i][1]}, actual discharge: {scaler.
↪inverse_transform(stage_discharge_test_values)[i][1]}")

```

```

1/1 [=====] - 0s 124ms/step
3.2816958696629652
1.3114376019170655
pred stage: 2.3442065715789795, actual stage: 2.32
pred discharge: 300.8064880371094, actual discharge: 268.0
pred stage: 3.580303192138672, actual stage: 3.6
pred discharge: 1679.95751953125, actual discharge: 1690.0
pred stage: 2.792109966278076, actual stage: 2.83
pred discharge: 689.0064086914062, actual discharge: 717.0
pred stage: 2.2538816928863525, actual stage: 2.24
pred discharge: 218.16615295410156, actual discharge: 197.0
pred stage: 3.1894640922546387, actual stage: 3.21
pred discharge: 1358.2999267578125, actual discharge: 1310.0
pred stage: 3.4943759441375732, actual stage: 3.51
pred discharge: 1605.3167724609375, actual discharge: 1650.0
pred stage: 2.461885929107666, actual stage: 2.47
pred discharge: 382.6040344238281, actual discharge: 375.0
pred stage: 2.8839025497436523, actual stage: 2.77
pred discharge: 788.4247436523438, actual discharge: 634.0
pred stage: 3.52038836479187, actual stage: 3.47
pred discharge: 1667.7935791015625, actual discharge: 1620.0
pred stage: 2.184338092803955, actual stage: 2.16
pred discharge: 150.7027130126953, actual discharge: 159.0

```

```
[ ]:
```

0.5 Visualize layers

```

[ ]: layer_outputs = [layer.output for layer in best_model.layers[:12]]
    # Extracts the outputs of the top 12 layers
    activation_model = models.Model(inputs=best_model.input, outputs=layer_outputs)
    ↪# Creates a model that will return these outputs, given the model input

```

```

[ ]: activations = activation_model.predict(test_ds.take(1))

```

```

1/1 [=====] - 0s 222ms/step

```

```

[ ]: import matplotlib.pyplot as plt

    layer_names = []
    for layer in best_model.layers[:12]:
        layer_names.append(layer.name) # Names of the layers, so you can have them
    ↪as part of your plot

```

```

images_per_row = 16

for layer_name, layer_activation in zip(layer_names, activations): # Displays
    ↪ the feature maps
        n_features = layer_activation.shape[-1] # Number of features in the feature
    ↪ map
        size = layer_activation.shape[1] # The feature map has shape (1, size, size,
    ↪ n_features).
        n_cols = n_features // images_per_row # Tiles the activation channels in
    ↪ this matrix
        display_grid = np.zeros((size * n_cols, images_per_row * size))

        print(layer_name)
        if ("flatten" in layer_name): break

        for col in range(n_cols): # Tiles each filter into a big horizontal grid
            for row in range(images_per_row):
                channel_image = layer_activation[0,
                                                :, :,
                                                col * images_per_row + row]
                channel_image -= channel_image.mean() # Post-processes the feature
    ↪ to make it visually palatable
                channel_image /= channel_image.std()
                channel_image *= 64
                channel_image += 128
                channel_image = np.clip(channel_image, 0, 255).astype('uint8')
                display_grid[col * size : (col + 1) * size, # Displays the grid
                            row * size : (row + 1) * size] = channel_image

        scale = 1. / size
        plt.figure(figsize=(scale * display_grid.shape[1],
                            scale * display_grid.shape[0]))
        plt.title(layer_name)
        plt.grid(False)
        plt.imshow(display_grid, aspect='auto', cmap='viridis')

```

```

conv2d
max_pooling2d
conv2d_1
max_pooling2d_1
conv2d_2
average_pooling2d
conv2d_3
average_pooling2d_1

```

```

/tmp/ipykernel_11977/2269795348.py:24: RuntimeWarning: invalid value encountered
in divide

```



```
channel_image /= channel_image.std()
```

```
-----  
MemoryError                                Traceback (most recent call last)  
Cell In [35], line 13  
    11 size = layer_activation.shape[1] #The feature map has shape (1, size,   
    ↪size, n_features).  
    12 n_cols = n_features // images_per_row # Tiles the activation channels i  
    ↪this matrix  
--> 13 display_grid = np.zeros((size * n_cols, images_per_row * size))  
    15 print(layer_name)  
    16 if ("flatten" in layer_name): break  
  
MemoryError: Unable to allocate 28.7 GiB for an array with shape (153664, 25088  
    ↪and data type float64
```





