

## cnn\_v8

October 14, 2022

```
[ ]: %env LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$CONDA_PREFIX/lib/
```

```
env: LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$CONDA_PREFIX/lib/
```

```
[ ]: import os
print(os.environ["LD_LIBRARY_PATH"])
```

```
:/home/nkspartan/miniconda3/envs/tf-gpu/lib/:/home/nkspartan/miniconda3/envs/tf-gpu/lib/
```

```
[ ]: import tensorflow as tf
import numpy as np
import pandas as pd
import os
import keras

from keras import Sequential, models, Input
from keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, Dropout,
↳LeakyReLU, AveragePooling2D
from keras.optimizers import SGD, Adam
```

```
[ ]: from tensorflow.python.client import device_lib

#print(device_lib.list_local_devices())
print('Default GPU Device: {}'.format(tf.test.gpu_device_name()))
```

```
Default GPU Device: /device:GPU:0
```

```
2022-10-14 11:39:11.186822: I tensorflow/core/platform/cpu_feature_guard.cc:193]
This TensorFlow binary is optimized with oneAPI Deep Neural Network Library
(oneDNN) to use the following CPU instructions in performance-critical
operations: AVX2 FMA
```

```
To enable them in other operations, rebuild TensorFlow with the appropriate
compiler flags.
```

```
2022-10-14 11:39:11.209616: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
```

```
2022-10-14 11:39:11.261664: I
```

```

tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:39:11.261850: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:39:12.100458: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:39:12.101067: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:39:12.101224: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:39:12.101367: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:1616] Created device
/device:GPU:0 with 4023 MB memory: -> device: 0, name: NVIDIA GeForce RTX 2060,
pci bus id: 0000:08:00.0, compute capability: 7.5

```

## 0.1 Read the csv dataset to get the values for stage and discharge of the images

```

[ ]: df = pd.read_csv("../dataset/2012_2019_PlatteRiverWeir_features_merged_all.
↳csv")
df.head()

```

```

[ ]:
   Unnamed: 0  SensorTime  CaptureTime \
0           0  2012-06-09 13:15:00  2012-06-09T13:09:07
1           1  2012-06-09 13:15:00  2012-06-09T13:10:29
2           2  2012-06-09 13:45:00  2012-06-09T13:44:01
3           3  2012-06-09 14:45:00  2012-06-09T14:44:30
4           4  2012-06-09 15:45:00  2012-06-09T15:44:59

```

```

   Filename Agency SiteNumber TimeZone Stage \
0  StateLineWeir_20120609_Farrell_001.jpg  USGS      6674500      MDT      2.99
1  StateLineWeir_20120609_Farrell_002.jpg  USGS      6674500      MDT      2.99
2  StateLineWeir_20120609_Farrell_003.jpg  USGS      6674500      MDT      2.96
3  StateLineWeir_20120609_Farrell_004.jpg  USGS      6674500      MDT      2.94
4  StateLineWeir_20120609_Farrell_005.jpg  USGS      6674500      MDT      2.94

```

```

   Discharge  CalcTimestamp  ... WeirPt2X WeirPt2Y WwRawLineMin \
0      916.0  2020-03-11T16:58:28  ...      -1      -1          0.0
1      916.0  2020-03-11T16:58:33  ...      -1      -1          0.0

```

2	873.0	2020-03-11T16:58:40	...	-1	-1	0.0
3	846.0	2020-03-11T16:58:47	...	-1	-1	0.0
4	846.0	2020-03-11T16:58:55	...	-1	-1	0.0

	WwRawLineMax	WwRawLineMean	WwRawLineSigma	WwCurveLineMin	\
0	0.0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0

	WwCurveLineMax	WwCurveLineMean	WwCurveLineSigma
0	0.0	0.0	0.0
1	0.0	0.0	0.0
2	0.0	0.0	0.0
3	0.0	0.0	0.0
4	0.0	0.0	0.0

[5 rows x 60 columns]

```
[ ]: df = df[["Filename", "Stage", "Discharge"]]
```

### 0.1.1 Scale the data

```
[ ]: from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
```

```
[ ]: df[["Stage", "Discharge"]] = scaler.fit_transform(df[["Stage", "Discharge"]])
df
```

```
[ ]:
      Filename      Stage  Discharge
0  StateLineWeir_20120609_Farrell_001.jpg  0.138117 -0.046094
1  StateLineWeir_20120609_Farrell_002.jpg  0.138117 -0.046094
2  StateLineWeir_20120609_Farrell_003.jpg  0.100875 -0.082160
3  StateLineWeir_20120609_Farrell_004.jpg  0.076046 -0.104807
4  StateLineWeir_20120609_Farrell_005.jpg  0.076046 -0.104807
...
42054  StateLineWeir_20191011_Farrell_409.jpg -0.420526 -0.450369
42055  StateLineWeir_20191011_Farrell_410.jpg -0.420526 -0.450369
42056  StateLineWeir_20191011_Farrell_411.jpg -0.420526 -0.450369
42057  StateLineWeir_20191011_Farrell_412.jpg -0.420526 -0.450369
42058  StateLineWeir_20191011_Farrell_413.jpg -0.420526 -0.450369
```

[42059 rows x 3 columns]

```
[ ]: from joblib import dump, load
dump(scaler, 'std_scaler.joblib', compress=True)
```

```
[ ]: ['std_scaler.joblib']
```

## 0.2 Create the dataset pipeline

```
[ ]: IMG_SIZE = 512
      BATCH_SIZE = 32
```

```
[ ]: from glob import glob

def make_dataset(path, batch_size, df, seed=None):
    np.random.seed(seed)

    def parse_image(filename):
        image = tf.io.read_file(filename)
        image = tf.image.decode_jpeg(image, channels=3)
        #image = tf.image.resize(image, [IMG_SIZE, IMG_SIZE])
        image = tf.cast(image, tf.float32)
        image /= 255
        return image

    def configure_for_performance(ds):
        ds = ds.shuffle(buffer_size=100)
        ds = ds.batch(batch_size)
        ds = ds.repeat()
        ds = ds.prefetch(buffer_size=tf.data.experimental.AUTOTUNE)
        return ds

    filenames = glob(path + '/*')

    # make train, val and test splits of the dataset (70%, 10%, 20% split)
    split1 = int(0.7 * len(filenames))
    split2 = int(0.8 * len(filenames))

    np.random.shuffle(filenames)
    train_files = filenames[:split1] # up to split 1 (ex 70%)
    val_files = filenames[split1:split2] # from ex. 70% to 80%
    test_files = filenames[split2:] # from ex. 80% until the end

    # create stage values
    stage_train_values = [df[df.Filename == file.split('/')[0]].Stage.values for
    ↪file in train_files]
    stage_val_values = [df[df.Filename == file.split('/')[0]].Stage.values for
    ↪file in val_files]
```

```

stage_test_values = [df[df.Filename == file.split('/')[1]].Stage.values for
↪file in test_files]

# create discharge values
discharge_train_values = [df[df.Filename == file.split(
    '/')[-1]].Discharge.values for file in train_files]
discharge_val_values = [df[df.Filename == file.split(
    '/')[-1]].Discharge.values for file in val_files]
discharge_test_values = [df[df.Filename == file.split(
    '/')[-1]].Discharge.values for file in test_files]

# join stage and discharge values
stage_discharge_train_values = [[np.squeeze(s), np.squeeze(d)] for s, d in
↪zip(stage_train_values, discharge_train_values)]
stage_discharge_val_values = [[np.squeeze(s), np.squeeze(d)] for s, d in
↪zip(stage_val_values, discharge_val_values)]
stage_discharge_test_values = [[np.squeeze(s), np.squeeze(
    d)] for s, d in zip(stage_test_values, discharge_test_values)]

# create images dataset (train, val, test)
filenames_train_ds = tf.data.Dataset.from_tensor_slices(train_files)
filenames_val_ds = tf.data.Dataset.from_tensor_slices(val_files)
filenames_test_ds = tf.data.Dataset.from_tensor_slices(test_files)

images_train_ds = filenames_train_ds.map(parse_image, num_parallel_calls=5)
images_val_ds = filenames_val_ds.map(parse_image, num_parallel_calls=5)
images_test_ds = filenames_test_ds.map(parse_image, num_parallel_calls=5)

# create stage and discharge dataset (train, val, test)
stage_discharge_train_ds = tf.data.Dataset.
↪from_tensor_slices(stage_discharge_train_values)
stage_discharge_val_ds = tf.data.Dataset.
↪from_tensor_slices(stage_discharge_val_values)
stage_discharge_test_ds = tf.data.Dataset.from_tensor_slices(
    stage_discharge_test_values)

# create tensorflow dataset of images and values (train, val, test)
train_ds = tf.data.Dataset.zip((images_train_ds, stage_discharge_train_ds))
train_ds = configure_for_performance(train_ds)
val_ds = tf.data.Dataset.zip((images_val_ds, stage_discharge_val_ds))
val_ds = configure_for_performance(val_ds)
test_ds = tf.data.Dataset.zip((images_test_ds, stage_discharge_test_ds))
test_ds = configure_for_performance(test_ds)

return train_ds, len(train_files), val_ds, len(val_files), test_ds,
↪len(test_files)

```

```
[ ]: path = "../..../dataset/images"

train_ds, train_size, val_ds, val_size, test_ds, test_size = make_dataset(path,
↳ BATCH_SIZE, df, 0)
```

```
2022-10-14 11:47:25.673029: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:47:25.673317: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:47:25.673566: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:47:25.674117: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:47:25.674268: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:47:25.674407: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:47:25.674587: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:47:25.674729: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-14 11:47:25.674844: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:1616] Created device
/job:localhost/replica:0/task:0/device:GPU:0 with 4023 MB memory:  -> device: 0,
name: NVIDIA GeForce RTX 2060, pci bus id: 0000:08:00.0, compute capability: 7.5
```

```
[ ]: input_shape = 0
output_shape = 0

for image, stage_discharge in train_ds.take(1):
    print(image.numpy().shape)
```

```
print(stage_discharge.numpy().shape)

input_shape = image.numpy().shape[1:]
output_shape = stage_discharge.numpy().shape[1:]
```

```
(32, 512, 512, 3)
(32, 2)
```

```
[ ]: print(input_shape)
      print(output_shape)
```

```
(512, 512, 3)
(2,)
```

### 0.3 Create model

```
[ ]: def create_model(input_shape, output_shape):
      model = Sequential()

      model.add(Input(shape=input_shape))

      model.add(Conv2D(64, kernel_size=(4, 4), strides=(2, 2), padding='same',
      ↪activation=LeakyReLU()))
      model.add(MaxPooling2D(pool_size=(4, 4)))

      model.add(Conv2D(64, kernel_size=(4, 4), activation=LeakyReLU(),
      ↪padding='same'))
      model.add(MaxPooling2D(pool_size=(2, 2)))

      model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', padding='same'))
      #model.add(AveragePooling2D(pool_size=(2, 2)))

      model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
      model.add(AveragePooling2D(pool_size=(2, 2)))

      model.add(Flatten())
      model.add(Dense(128, activation='tanh'))
      model.add(Dense(64, activation='tanh'))
      model.add(Dense(32, activation='tanh'))
      model.add(Dense(32, activation='tanh'))
      model.add(Dense(output_shape, activation='linear')) # linear regression
      ↪output layer

      return model
```

```
[ ]: model = create_model(input_shape, output_shape[0])
```

```
[ ]: model.summary()
```

```
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 256, 256, 64)	3136
max_pooling2d_2 (MaxPooling2D)	(None, 64, 64, 64)	0
conv2d_5 (Conv2D)	(None, 64, 64, 64)	65600
max_pooling2d_3 (MaxPooling2D)	(None, 32, 32, 64)	0
conv2d_6 (Conv2D)	(None, 32, 32, 32)	18464
conv2d_7 (Conv2D)	(None, 30, 30, 32)	9248
average_pooling2d_2 (AveragePooling2D)	(None, 15, 15, 32)	0
flatten_1 (Flatten)	(None, 7200)	0
dense_5 (Dense)	(None, 128)	921728
dense_6 (Dense)	(None, 64)	8256
dense_7 (Dense)	(None, 32)	2080
dense_8 (Dense)	(None, 32)	1056
dense_9 (Dense)	(None, 2)	66
Total params: 1,029,634		
Trainable params: 1,029,634		
Non-trainable params: 0		

```
[ ]: def compile_model(loss_func, optimizer, metrics=["accuracy"]):  
      model.compile(loss=loss_func, optimizer=optimizer, metrics=metrics)
```

```
[ ]: sgd = SGD(learning_rate=0.01, decay=1e-4, momentum=0.9, nesterov=True)  
      adam = Adam(learning_rate=1e-3, decay=1e-3 / 100)
```



```
compile_model('mse', adam, [
    'mse', tf.keras.metrics.RootMeanSquaredError(name='rmse'), 'mae',
    'mape'])
```

```
[ ]: def fit_model(training_values, validation_values=None, epochs=10, steps=32,
    val_steps=32, callbacks=[]):
    return model.fit(training_values, validation_data=validation_values,
    epochs=epochs, steps_per_epoch=steps, validation_steps=val_steps,
    callbacks=callbacks)
```

```
[ ]: import datetime

date_actual = datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
log_dir = "logs/fit/" + date_actual
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir,
    histogram_freq=1)

checkpoint_callback = tf.keras.callbacks.
    ModelCheckpoint(filepath=f"model_weights/{date_actual}_cnn_best_weights.
    hdf5",
                    monitor='val_mse',
                    verbose=1,
                    save_best_only=True)
```

```
[ ]: # batch_size = 0 because we already have batch size in tf dataset
history = fit_model(train_ds, val_ds, epochs=25, steps=np.ceil(train_size /
    BATCH_SIZE), val_steps=np.ceil(val_size / BATCH_SIZE),
    callbacks=[tensorboard_callback, checkpoint_callback])
```

Epoch 1/25

921/921 [=====] - ETA: 0s - loss: 0.0987 - mse: 0.0987  
- rmse: 0.3142 - mae: 0.1597 - mape: 66.8773

Epoch 1: val\_mse improved from inf to 0.01878, saving model to  
model\_weights/20221014-123925\_cnn\_best\_weights.hdf5

921/921 [=====] - 111s 120ms/step - loss: 0.0987 - mse:  
0.0987 - rmse: 0.3142 - mae: 0.1597 - mape: 66.8773 - val\_loss: 0.0188 -  
val\_mse: 0.0188 - val\_rmse: 0.1370 - val\_mae: 0.0888 - val\_mape: 32.6099

Epoch 2/25

920/921 [=====>.] - ETA: 0s - loss: 0.0129 - mse: 0.0129  
- rmse: 0.1134 - mae: 0.0691 - mape: 35.6959

Epoch 2: val\_mse improved from 0.01878 to 0.00667, saving model to  
model\_weights/20221014-123925\_cnn\_best\_weights.hdf5

921/921 [=====] - 105s 114ms/step - loss: 0.0129 - mse:  
0.0129 - rmse: 0.1134 - mae: 0.0691 - mape: 35.6947 - val\_loss: 0.0067 -  
val\_mse: 0.0067 - val\_rmse: 0.0817 - val\_mae: 0.0544 - val\_mape: 21.2948

Epoch 3/25

920/921 [=====>.] - ETA: 0s - loss: 0.0056 - mse: 0.0056

```

- rmse: 0.0747 - mae: 0.0505 - mape: 23.9313
Epoch 3: val_mse improved from 0.00667 to 0.00385, saving model to
model_weights/20221014-123925_cnn_best_weights.hdf5
921/921 [=====] - 107s 117ms/step - loss: 0.0056 - mse:
0.0056 - rmse: 0.0747 - mae: 0.0505 - mape: 23.9305 - val_loss: 0.0039 -
val_mse: 0.0039 - val_rmse: 0.0621 - val_mae: 0.0414 - val_mape: 14.9443
Epoch 4/25
920/921 [=====>.] - ETA: 0s - loss: 0.0044 - mse: 0.0044
- rmse: 0.0665 - mae: 0.0452 - mape: 21.1921
Epoch 4: val_mse did not improve from 0.00385
921/921 [=====] - 108s 118ms/step - loss: 0.0044 - mse:
0.0044 - rmse: 0.0665 - mae: 0.0452 - mape: 21.1915 - val_loss: 0.0044 -
val_mse: 0.0044 - val_rmse: 0.0664 - val_mae: 0.0528 - val_mape: 15.3347
Epoch 5/25
920/921 [=====>.] - ETA: 0s - loss: 0.0035 - mse: 0.0035
- rmse: 0.0591 - mae: 0.0401 - mape: 18.6906
Epoch 5: val_mse did not improve from 0.00385
921/921 [=====] - 110s 119ms/step - loss: 0.0035 - mse:
0.0035 - rmse: 0.0591 - mae: 0.0401 - mape: 18.6901 - val_loss: 0.0065 -
val_mse: 0.0065 - val_rmse: 0.0804 - val_mae: 0.0568 - val_mape: 16.2642
Epoch 6/25
920/921 [=====>.] - ETA: 0s - loss: 0.0032 - mse: 0.0032
- rmse: 0.0568 - mae: 0.0377 - mape: 16.6648
Epoch 6: val_mse did not improve from 0.00385
921/921 [=====] - 107s 116ms/step - loss: 0.0032 - mse:
0.0032 - rmse: 0.0568 - mae: 0.0377 - mape: 16.6654 - val_loss: 0.0048 -
val_mse: 0.0048 - val_rmse: 0.0696 - val_mae: 0.0557 - val_mape: 22.7504
Epoch 7/25
920/921 [=====>.] - ETA: 0s - loss: 0.0029 - mse: 0.0029
- rmse: 0.0536 - mae: 0.0355 - mape: 17.7932
Epoch 7: val_mse did not improve from 0.00385
921/921 [=====] - 98s 106ms/step - loss: 0.0029 - mse:
0.0029 - rmse: 0.0536 - mae: 0.0355 - mape: 17.7930 - val_loss: 0.0071 -
val_mse: 0.0071 - val_rmse: 0.0843 - val_mae: 0.0613 - val_mape: 21.6124
Epoch 8/25
920/921 [=====>.] - ETA: 0s - loss: 0.0030 - mse: 0.0030
- rmse: 0.0547 - mae: 0.0360 - mape: 19.0689
Epoch 8: val_mse improved from 0.00385 to 0.00231, saving model to
model_weights/20221014-123925_cnn_best_weights.hdf5
921/921 [=====] - 102s 111ms/step - loss: 0.0030 - mse:
0.0030 - rmse: 0.0547 - mae: 0.0360 - mape: 19.0684 - val_loss: 0.0023 -
val_mse: 0.0023 - val_rmse: 0.0480 - val_mae: 0.0321 - val_mape: 11.6860
Epoch 9/25
920/921 [=====>.] - ETA: 0s - loss: 0.0023 - mse: 0.0023
- rmse: 0.0479 - mae: 0.0313 - mape: 15.5731
Epoch 9: val_mse improved from 0.00231 to 0.00201, saving model to
model_weights/20221014-123925_cnn_best_weights.hdf5
921/921 [=====] - 110s 119ms/step - loss: 0.0023 - mse:

```

0.0023 - rmse: 0.0479 - mae: 0.0313 - mape: 15.5728 - val\_loss: 0.0020 -  
val\_mse: 0.0020 - val\_rmse: 0.0449 - val\_mae: 0.0316 - val\_mape: 14.6471  
Epoch 10/25  
920/921 [=====>.] - ETA: 0s - loss: 0.0021 - mse: 0.0021  
- rmse: 0.0462 - mae: 0.0303 - mape: 16.4243  
Epoch 10: val\_mse improved from 0.00201 to 0.00165, saving model to  
model\_weights/20221014-123925\_cnn\_best\_weights.hdf5  
921/921 [=====] - 105s 114ms/step - loss: 0.0021 - mse:  
0.0021 - rmse: 0.0462 - mae: 0.0303 - mape: 16.4241 - val\_loss: 0.0017 -  
val\_mse: 0.0017 - val\_rmse: 0.0407 - val\_mae: 0.0263 - val\_mape: 10.4725  
Epoch 11/25  
920/921 [=====>.] - ETA: 0s - loss: 0.0021 - mse: 0.0021  
- rmse: 0.0457 - mae: 0.0297 - mape: 16.0292  
Epoch 11: val\_mse did not improve from 0.00165  
921/921 [=====] - 108s 118ms/step - loss: 0.0021 - mse:  
0.0021 - rmse: 0.0457 - mae: 0.0297 - mape: 16.0287 - val\_loss: 0.0021 -  
val\_mse: 0.0021 - val\_rmse: 0.0457 - val\_mae: 0.0323 - val\_mape: 16.1571  
Epoch 12/25  
920/921 [=====>.] - ETA: 0s - loss: 0.0020 - mse: 0.0020  
- rmse: 0.0446 - mae: 0.0284 - mape: 13.2536  
Epoch 12: val\_mse did not improve from 0.00165  
921/921 [=====] - 110s 119ms/step - loss: 0.0020 - mse:  
0.0020 - rmse: 0.0446 - mae: 0.0284 - mape: 13.2532 - val\_loss: 0.0017 -  
val\_mse: 0.0017 - val\_rmse: 0.0407 - val\_mae: 0.0265 - val\_mape: 10.0987  
Epoch 13/25  
920/921 [=====>.] - ETA: 0s - loss: 0.0019 - mse: 0.0019  
- rmse: 0.0433 - mae: 0.0276 - mape: 13.7047  
Epoch 13: val\_mse improved from 0.00165 to 0.00164, saving model to  
model\_weights/20221014-123925\_cnn\_best\_weights.hdf5  
921/921 [=====] - 111s 121ms/step - loss: 0.0019 - mse:  
0.0019 - rmse: 0.0433 - mae: 0.0276 - mape: 13.7043 - val\_loss: 0.0016 -  
val\_mse: 0.0016 - val\_rmse: 0.0405 - val\_mae: 0.0261 - val\_mape: 10.0526  
Epoch 14/25  
920/921 [=====>.] - ETA: 0s - loss: 0.0016 - mse: 0.0016  
- rmse: 0.0394 - mae: 0.0252 - mape: 14.1183  
Epoch 14: val\_mse improved from 0.00164 to 0.00158, saving model to  
model\_weights/20221014-123925\_cnn\_best\_weights.hdf5  
921/921 [=====] - 111s 121ms/step - loss: 0.0016 - mse:  
0.0016 - rmse: 0.0394 - mae: 0.0252 - mape: 14.1179 - val\_loss: 0.0016 -  
val\_mse: 0.0016 - val\_rmse: 0.0398 - val\_mae: 0.0264 - val\_mape: 9.7027  
Epoch 15/25  
920/921 [=====>.] - ETA: 0s - loss: 0.0054 - mse: 0.0054  
- rmse: 0.0737 - mae: 0.0429 - mape: 20.7148  
Epoch 15: val\_mse did not improve from 0.00158  
921/921 [=====] - 111s 121ms/step - loss: 0.0054 - mse:  
0.0054 - rmse: 0.0737 - mae: 0.0429 - mape: 20.7143 - val\_loss: 0.0037 -  
val\_mse: 0.0037 - val\_rmse: 0.0612 - val\_mae: 0.0451 - val\_mape: 15.5806  
Epoch 16/25

920/921 [=====>.] - ETA: 0s - loss: 0.0022 - mse: 0.0022  
- rmse: 0.0466 - mae: 0.0301 - mape: 17.8235  
Epoch 16: val\_mse did not improve from 0.00158  
921/921 [=====] - 112s 122ms/step - loss: 0.0022 - mse:  
0.0022 - rmse: 0.0466 - mae: 0.0301 - mape: 17.8233 - val\_loss: 0.0017 -  
val\_mse: 0.0017 - val\_rmse: 0.0411 - val\_mae: 0.0280 - val\_mape: 9.7793  
Epoch 17/25  
920/921 [=====>.] - ETA: 0s - loss: 0.0014 - mse: 0.0014  
- rmse: 0.0373 - mae: 0.0233 - mape: 12.1485  
Epoch 17: val\_mse improved from 0.00158 to 0.00146, saving model to  
model\_weights/20221014-123925\_cnn\_best\_weights.hdf5  
921/921 [=====] - 109s 118ms/step - loss: 0.0014 - mse:  
0.0014 - rmse: 0.0373 - mae: 0.0233 - mape: 12.1733 - val\_loss: 0.0015 -  
val\_mse: 0.0015 - val\_rmse: 0.0382 - val\_mae: 0.0271 - val\_mape: 10.0689  
Epoch 18/25  
920/921 [=====>.] - ETA: 0s - loss: 0.0012 - mse: 0.0012  
- rmse: 0.0344 - mae: 0.0215 - mape: 12.4680  
Epoch 18: val\_mse improved from 0.00146 to 0.00134, saving model to  
model\_weights/20221014-123925\_cnn\_best\_weights.hdf5  
921/921 [=====] - 105s 114ms/step - loss: 0.0012 - mse:  
0.0012 - rmse: 0.0344 - mae: 0.0215 - mape: 12.4682 - val\_loss: 0.0013 -  
val\_mse: 0.0013 - val\_rmse: 0.0365 - val\_mae: 0.0239 - val\_mape: 8.8601  
Epoch 19/25  
920/921 [=====>.] - ETA: 0s - loss: 0.0012 - mse: 0.0012  
- rmse: 0.0342 - mae: 0.0211 - mape: 10.3338  
Epoch 19: val\_mse did not improve from 0.00134  
921/921 [=====] - 107s 116ms/step - loss: 0.0012 - mse:  
0.0012 - rmse: 0.0342 - mae: 0.0211 - mape: 10.3336 - val\_loss: 0.0018 -  
val\_mse: 0.0018 - val\_rmse: 0.0419 - val\_mae: 0.0292 - val\_mape: 13.8838  
Epoch 20/25  
920/921 [=====>.] - ETA: 0s - loss: 0.0012 - mse: 0.0012  
- rmse: 0.0347 - mae: 0.0215 - mape: 12.0982  
Epoch 20: val\_mse did not improve from 0.00134  
921/921 [=====] - 107s 117ms/step - loss: 0.0012 - mse:  
0.0012 - rmse: 0.0347 - mae: 0.0215 - mape: 12.0980 - val\_loss: 0.0024 -  
val\_mse: 0.0024 - val\_rmse: 0.0489 - val\_mae: 0.0341 - val\_mape: 13.9538  
Epoch 21/25  
920/921 [=====>.] - ETA: 0s - loss: 0.0013 - mse: 0.0013  
- rmse: 0.0366 - mae: 0.0225 - mape: 12.6725  
Epoch 21: val\_mse did not improve from 0.00134  
921/921 [=====] - 107s 116ms/step - loss: 0.0013 - mse:  
0.0013 - rmse: 0.0366 - mae: 0.0225 - mape: 12.6723 - val\_loss: 0.0049 -  
val\_mse: 0.0049 - val\_rmse: 0.0700 - val\_mae: 0.0554 - val\_mape: 20.3059  
Epoch 22/25  
920/921 [=====>.] - ETA: 0s - loss: 0.0032 - mse: 0.0032  
- rmse: 0.0565 - mae: 0.0338 - mape: 17.4778  
Epoch 22: val\_mse did not improve from 0.00134  
921/921 [=====] - 108s 117ms/step - loss: 0.0032 - mse:

```

0.0032 - rmse: 0.0565 - mae: 0.0338 - mape: 17.4773 - val_loss: 0.0014 -
val_mse: 0.0014 - val_rmse: 0.0376 - val_mae: 0.0244 - val_mape: 9.6925
Epoch 23/25
920/921 [=====>.] - ETA: 0s - loss: 0.0011 - mse: 0.0011
- rmse: 0.0331 - mae: 0.0201 - mape: 10.4921
Epoch 23: val_mse improved from 0.00134 to 0.00129, saving model to
model_weights/20221014-123925_cnn_best_weights.hdf5
921/921 [=====] - 108s 117ms/step - loss: 0.0011 - mse:
0.0011 - rmse: 0.0331 - mae: 0.0201 - mape: 10.4920 - val_loss: 0.0013 -
val_mse: 0.0013 - val_rmse: 0.0359 - val_mae: 0.0228 - val_mape: 8.9852
Epoch 24/25
920/921 [=====>.] - ETA: 0s - loss: 9.6015e-04 - mse:
9.6015e-04 - rmse: 0.0310 - mae: 0.0186 - mape: 9.2455
Epoch 24: val_mse did not improve from 0.00129
921/921 [=====] - 105s 114ms/step - loss: 9.6012e-04 -
mse: 9.6012e-04 - rmse: 0.0310 - mae: 0.0186 - mape: 9.2453 - val_loss: 0.0014 -
val_mse: 0.0014 - val_rmse: 0.0370 - val_mae: 0.0239 - val_mape: 10.2687
Epoch 25/25
920/921 [=====>.] - ETA: 0s - loss: 9.8383e-04 - mse:
9.8383e-04 - rmse: 0.0314 - mae: 0.0191 - mape: 9.8359
Epoch 25: val_mse did not improve from 0.00129
921/921 [=====] - 107s 116ms/step - loss: 9.8380e-04 -
mse: 9.8380e-04 - rmse: 0.0314 - mae: 0.0191 - mape: 9.8356 - val_loss: 0.0014 -
val_mse: 0.0014 - val_rmse: 0.0372 - val_mae: 0.0243 - val_mape: 10.0188

```

## 0.4 Evaluate model

```

[ ]: print(date_actual)

20221014-123925

[ ]: best_model = models.load_model(f'model_weights/{date_actual}_cnn_best_weights.
    ↪hdf5')

[ ]: def evaluate_model(model, test_values, steps):
    score = model.evaluate(test_values, steps=steps)
    return score

[ ]: test_loss, test_mse, test_rmse, test_mae, test_mape = ↪
    ↪evaluate_model(best_model, test_ds, steps=np.ceil(test_size / BATCH_SIZE))

263/263 [=====] - 17s 64ms/step - loss: 0.0042 - mse:
0.0042 - rmse: 0.0649 - mae: 0.0239 - mape: 18.2376

[ ]: predictions = best_model.predict(test_ds, steps=np.ceil(test_size / BATCH_SIZE))

263/263 [=====] - 17s 64ms/step

```

```
[ ]: for image, stage_discharge in test_ds.take(1):
    predictions = best_model.predict(x=image)

    stage_discharge_test_values = stage_discharge[:2].numpy()
    predictions_values = predictions[:2]

    diff = predictions_values.flatten() - stage_discharge_test_values.
    →flatten()
    percentDiff = (diff / stage_discharge_test_values.flatten()) * 100
    absPercentDiff = np.abs(percentDiff)
    # compute the mean and standard deviation of the absolute percentage
    # difference
    mean = np.mean(absPercentDiff)
    std = np.std(absPercentDiff)
    # finally, show some statistics on our model
    print(mean)
    print(std)

    stage_discharge_test_values = stage_discharge[:10]
    predictions_values = predictions[:10]

    for i in range(len(stage_discharge_test_values.numpy())):
        print(f"pred stage: {scaler.
    →inverse_transform(predictions_values)[i][0]}, actual stage: {scaler.
    →inverse_transform(stage_discharge_test_values)[i][0]}")
        print(f"pred discharge: {scaler.
    →inverse_transform(predictions_values)[i][1]}, actual discharge: {scaler.
    →inverse_transform(stage_discharge_test_values)[i][1]}")
```

```
1/1 [=====] - 0s 144ms/step
76.13851851998916
70.14508968849664
pred stage: 3.2092437744140625, actual stage: 3.24
pred discharge: 1432.246337890625, actual discharge: 1450.0
pred stage: 3.1332528591156006, actual stage: 2.98
pred discharge: 1035.6925048828125, actual discharge: 813.0
pred stage: 2.30507493019104, actual stage: 2.31
pred discharge: 250.9044952392578, actual discharge: 261.0
pred stage: 3.960446357727051, actual stage: 4.0
pred discharge: 2270.589111328125, actual discharge: 2320.0
pred stage: 1.5294480323791504, actual stage: 1.54
pred discharge: 12.187579154968262, actual discharge: 14.899999999999977
pred stage: 5.178676128387451, actual stage: 5.22
pred discharge: 4734.4560546875, actual discharge: 4760.0
pred stage: 2.1457924842834473, actual stage: 2.15
pred discharge: 114.64216613769531, actual discharge: 151.0
pred stage: 2.6019279956817627, actual stage: 2.58
```

```

pred discharge: 484.6307678222656, actual discharge: 460.0
pred stage: 2.4353623390197754, actual stage: 2.46
pred discharge: 334.6253967285156, actual discharge: 343.0
pred stage: 4.1649675369262695, actual stage: 4.18
pred discharge: 2854.3115234375, actual discharge: 2850.0

```

```
[ ]:
```

## 0.5 Visualize layers

```
[ ]: layer_outputs = [layer.output for layer in best_model.layers[:12]]
# Extracts the outputs of the top 12 layers
activation_model = models.Model(inputs=best_model.input, outputs=layer_outputs)
↪ # Creates a model that will return these outputs, given the model input
```

```
[ ]: activations = activation_model.predict(test_ds.take(1))
```

```
1/1 [=====] - 0s 140ms/step
```

```
[ ]: import matplotlib.pyplot as plt

layer_names = []
for layer in best_model.layers[:12]:
    layer_names.append(layer.name) # Names of the layers, so you can have them
    ↪ as part of your plot

images_per_row = 16

for layer_name, layer_activation in zip(layer_names, activations): # Displays
    ↪ the feature maps
    n_features = layer_activation.shape[-1] # Number of features in the feature
    ↪ map
    size = layer_activation.shape[1] # The feature map has shape (1, size, size,
    ↪ n_features).
    n_cols = n_features // images_per_row # Tiles the activation channels in
    ↪ this matrix
    display_grid = np.zeros((size * n_cols, images_per_row * size))

    print(layer_name)
    if ("flatten" in layer_name): break

    for col in range(n_cols): # Tiles each filter into a big horizontal grid
        for row in range(images_per_row):
            channel_image = layer_activation[0,
                                                :, :,
                                                col * images_per_row + row]
```

```

        channel_image -= channel_image.mean() # Post-processes the feature
↳to make it visually palatable
        channel_image /= channel_image.std()
        channel_image *= 64
        channel_image += 128
        channel_image = np.clip(channel_image, 0, 255).astype('uint8')
        display_grid[col * size : (col + 1) * size, # Displays the grid
            row * size : (row + 1) * size] = channel_image

scale = 1. / size
plt.figure(figsize=(scale * display_grid.shape[1],
                    scale * display_grid.shape[0]))
plt.title(layer_name)
plt.grid(False)
plt.imshow(display_grid, aspect='auto', cmap='viridis')

```

```

conv2d_4
max_pooling2d_2
conv2d_5
max_pooling2d_3
conv2d_6
conv2d_7
average_pooling2d_2

```

```

-----
MemoryError                                Traceback (most recent call last)
Cell In [60], line 13
     11 size = layer_activation.shape[1] #The feature map has shape (1, size,
↳size, n_features).
     12 n_cols = n_features // images_per_row # Tiles the activation channels i
↳this matrix
--> 13 display_grid = np.zeros((size * n_cols, images_per_row * size))
     15 print(layer_name)
     16 if ("flatten" in layer_name): break

MemoryError: Unable to allocate 2.72 TiB for an array with shape (3240000,
↳115200) and data type float64

```







