

# cnn\_v11

October 30, 2022

```
[ ]: %env LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$CONDA_PREFIX/lib/  
      #%env TF_GPU_ALLOCATOR=cuda_malloc_async
```

```
env: LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$CONDA_PREFIX/lib/
```

```
[ ]: import os  
      print(os.environ["LD_LIBRARY_PATH"])
```

```
$LD_LIBRARY_PATH:$CONDA_PREFIX/lib/
```

```
[ ]: import tensorflow as tf  
      import numpy as np  
      import pandas as pd  
      import os  
      import keras  
      import matplotlib.pyplot as plt  
  
      from tensorflow.keras import Sequential, models, Input  
      from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D,  
      ↪Dropout, LeakyReLU, AveragePooling2D, GlobalAveragePooling2D,  
      ↪BatchNormalization, TimeDistributed, LSTM, SpatialDropout2D  
      from tensorflow.keras.optimizers import SGD, Adam
```

```
2022-10-30 19:17:16.242430: I tensorflow/core/platform/cpu_feature_guard.cc:193]  
This TensorFlow binary is optimized with oneAPI Deep Neural Network Library  
(oneDNN) to use the following CPU instructions in performance-critical  
operations:  AVX2 FMA
```

```
To enable them in other operations, rebuild TensorFlow with the appropriate  
compiler flags.
```

```
2022-10-30 19:17:16.678262: E tensorflow/stream_executor/cuda/cuda_blas.cc:2981]  
Unable to register cuBLAS factory: Attempting to register factory for plugin  
cuBLAS when one has already been registered
```

```
2022-10-30 19:17:17.727708: W  
tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load  
dynamic library 'libnvinfer.so.7'; dlopen: libnvinfer.so.7: cannot open shared  
object file: No such file or directory; LD_LIBRARY_PATH:  
:/home/nkspartan/miniconda3/envs/tf-gpu/lib/:/home/nkspartan/miniconda3/envs/tf-  
gpu/lib/
```

```
2022-10-30 19:17:17.727788: W
tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load
dynamic library 'libnvinfer_plugin.so.7'; dLError: libnvinfer_plugin.so.7:
cannot open shared object file: No such file or directory; LD_LIBRARY_PATH:
:/home/nkspartan/miniconda3/envs/tf-gpu/lib/;/home/nkspartan/miniconda3/envs/tf-
gpu/lib/
2022-10-30 19:17:17.727793: W
tensorflow/compiler/tf2tensorrt/utils/py_utils.cc:38] TF-TRT Warning: Cannot
dlopen some TensorRT libraries. If you would like to use Nvidia GPU with
TensorRT, please make sure the missing libraries mentioned above are installed
properly.
```

```
[ ]: from tensorflow.python.client import device_lib

print('Default GPU Device: {}'.format(tf.test.gpu_device_name()))
```

```
Default GPU Device: /device:GPU:0
```

```
2022-10-30 19:17:19.248427: I tensorflow/core/platform/cpu_feature_guard.cc:193]
This TensorFlow binary is optimized with oneAPI Deep Neural Network Library
(oneDNN) to use the following CPU instructions in performance-critical
operations: AVX2 FMA
```

```
To enable them in other operations, rebuild TensorFlow with the appropriate
compiler flags.
```

```
2022-10-30 19:17:19.281426: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
```

```
2022-10-30 19:17:19.322587: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
```

```
2022-10-30 19:17:19.322783: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
```

```
2022-10-30 19:17:20.170499: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
```

```
2022-10-30 19:17:20.170698: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
```

```
2022-10-30 19:17:20.170848: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
```

```
2022-10-30 19:17:20.170982: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:1616] Created device
/device:GPU:0 with 4016 MB memory: -> device: 0, name: NVIDIA GeForce RTX 2060,
pci bus id: 0000:08:00.0, compute capability: 7.5
```

```
[ ]: physical_devices = tf.config.list_physical_devices('GPU')
tf.config.experimental.set_memory_growth(physical_devices[0], True)
```

```
2022-10-30 19:17:20.209508: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-30 19:17:20.209776: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-10-30 19:17:20.209945: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
```

## 0.1 Read the csv dataset to get the values for stage and discharge of the images

```
[ ]: df = pd.read_csv("../dataset/2012_2019_PlatteRiverWeir_features_merged_all.
↪csv")
df.head()
```

```
[ ]: Unnamed: 0      SensorTime      CaptureTime \
0          0  2012-06-09 13:15:00  2012-06-09T13:09:07
1          1  2012-06-09 13:15:00  2012-06-09T13:10:29
2          2  2012-06-09 13:45:00  2012-06-09T13:44:01
3          3  2012-06-09 14:45:00  2012-06-09T14:44:30
4          4  2012-06-09 15:45:00  2012-06-09T15:44:59

      Filename Agency  SiteNumber TimeZone  Stage \
0  StateLineWeir_20120609_Farrell_001.jpg  USGS    6674500      MDT    2.99
1  StateLineWeir_20120609_Farrell_002.jpg  USGS    6674500      MDT    2.99
2  StateLineWeir_20120609_Farrell_003.jpg  USGS    6674500      MDT    2.96
3  StateLineWeir_20120609_Farrell_004.jpg  USGS    6674500      MDT    2.94
4  StateLineWeir_20120609_Farrell_005.jpg  USGS    6674500      MDT    2.94

      Discharge      CalcTimestamp  ...  WeirPt2X  WeirPt2Y  WwRawLineMin  \
0      916.0  2020-03-11T16:58:28  ...      -1      -1            0.0
1      916.0  2020-03-11T16:58:33  ...      -1      -1            0.0
2      873.0  2020-03-11T16:58:40  ...      -1      -1            0.0
3      846.0  2020-03-11T16:58:47  ...      -1      -1            0.0
4      846.0  2020-03-11T16:58:55  ...      -1      -1            0.0
```

	WwRawLineMax	WwRawLineMean	WwRawLineSigma	WwCurveLineMin	\
0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	

	WwCurveLineMax	WwCurveLineMean	WwCurveLineSigma
0	0.0	0.0	0.0
1	0.0	0.0	0.0
2	0.0	0.0	0.0
3	0.0	0.0	0.0
4	0.0	0.0	0.0

[5 rows x 60 columns]

```
[ ]: df = df[["Filename", "Stage", "Discharge", 'SensorTime']]
```

```
[ ]: df['SensorTime'] = pd.to_datetime(df['SensorTime'])
df['Year'] = df['SensorTime'].dt.year
df.head()
```

```
[ ]:
      Filename  Stage  Discharge  \
0  StateLineWeir_20120609_Farrell_001.jpg    2.99    916.0
1  StateLineWeir_20120609_Farrell_002.jpg    2.99    916.0
2  StateLineWeir_20120609_Farrell_003.jpg    2.96    873.0
3  StateLineWeir_20120609_Farrell_004.jpg    2.94    846.0
4  StateLineWeir_20120609_Farrell_005.jpg    2.94    846.0
```

	SensorTime	Year
0	2012-06-09 13:15:00	2012
1	2012-06-09 13:15:00	2012
2	2012-06-09 13:45:00	2012
3	2012-06-09 14:45:00	2012
4	2012-06-09 15:45:00	2012

### 0.1.1 Remove outliers

```
[ ]: df = df[df.Stage > 0]
df = df[df.Discharge > 0]
```

We consider values equal to 0 as outliers because from the photos it doesn't seem that it would be possible that at this time we would have a value of 0 for stage or discharge

```
[ ]: df.shape
```

```
[ ]: (40148, 5)
```

### 0.1.2 Scale the data

```
[ ]: from sklearn.preprocessing import StandardScaler
      from joblib import load

      scaler = StandardScaler()
      #scaler = load('std_scaler.joblib') # scaler with all the 42059 observations
```

```
[ ]: df[["Stage", "Discharge"]] = scaler.fit_transform(df[["Stage", "Discharge"]])
      df
```

```
[ ]:
      Filename      Stage  Discharge  \
0      StateLineWeir_20120609_Farrell_001.jpg  0.106063  -0.084154
1      StateLineWeir_20120609_Farrell_002.jpg  0.106063  -0.084154
2      StateLineWeir_20120609_Farrell_003.jpg  0.069235  -0.119960
3      StateLineWeir_20120609_Farrell_004.jpg  0.044683  -0.142442
4      StateLineWeir_20120609_Farrell_005.jpg  0.044683  -0.142442
...
42054  StateLineWeir_20191011_Farrell_409.jpg -0.446354  -0.485510
42055  StateLineWeir_20191011_Farrell_410.jpg -0.446354  -0.485510
42056  StateLineWeir_20191011_Farrell_411.jpg -0.446354  -0.485510
42057  StateLineWeir_20191011_Farrell_412.jpg -0.446354  -0.485510
42058  StateLineWeir_20191011_Farrell_413.jpg -0.446354  -0.485510
```

```
      SensorTime  Year
0      2012-06-09 13:15:00  2012
1      2012-06-09 13:15:00  2012
2      2012-06-09 13:45:00  2012
3      2012-06-09 14:45:00  2012
4      2012-06-09 15:45:00  2012
...
42054  2019-10-11 09:00:00  2019
42055  2019-10-11 10:00:00  2019
42056  2019-10-11 11:00:00  2019
42057  2019-10-11 12:00:00  2019
42058  2019-10-11 12:45:00  2019
```

```
[40148 rows x 5 columns]
```

```
[ ]: from joblib import dump
      #dump(scaler, 'std_scaler.joblib')
```

## 0.2 Create the dataset pipeline

```
[ ]: IMG_SIZE = 224
      #IMG_SIZE = 512
      BATCH_SIZE = 32
      FRAMES = 5

[ ]: from dataset_transformer import make_dataset

[ ]: path = "../../dataset/images_tmp_draw"

      train_ds, train_size, val_ds, val_size, test_ds, test_size = make_dataset(path,
      ↪BATCH_SIZE, IMG_SIZE, FRAMES, df, 10, True, "cnn")
```

20304

7117

12727

2022-10-30 19:20:08.734097: I

tensorflow/stream\_executor/cuda/cuda\_gpu\_executor.cc:980] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero

2022-10-30 19:20:08.734291: I

tensorflow/stream\_executor/cuda/cuda\_gpu\_executor.cc:980] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero

2022-10-30 19:20:08.734431: I

tensorflow/stream\_executor/cuda/cuda\_gpu\_executor.cc:980] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero

2022-10-30 19:20:08.734614: I

tensorflow/stream\_executor/cuda/cuda\_gpu\_executor.cc:980] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero

2022-10-30 19:20:08.734757: I

tensorflow/stream\_executor/cuda/cuda\_gpu\_executor.cc:980] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero

2022-10-30 19:20:08.734871: I

tensorflow/core/common\_runtime/gpu/gpu\_device.cc:1616] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 4016 MB memory: -> device: 0, name: NVIDIA GeForce RTX 2060, pci bus id: 0000:08:00.0, compute capability: 7.5

```
[ ]: input_shape = 0
      output_shape = 0

      for image, stage_discharge in train_ds.take(1):
          print(image.numpy().shape)
```

```
print(stage_discharge.numpy().shape)

input_shape = image.numpy().shape[1:]
output_shape = stage_discharge.numpy().shape[1:]
```

```
(32, 224, 224, 3)
(32, 2)
```

```
[ ]: print(input_shape)
      print(output_shape)
```

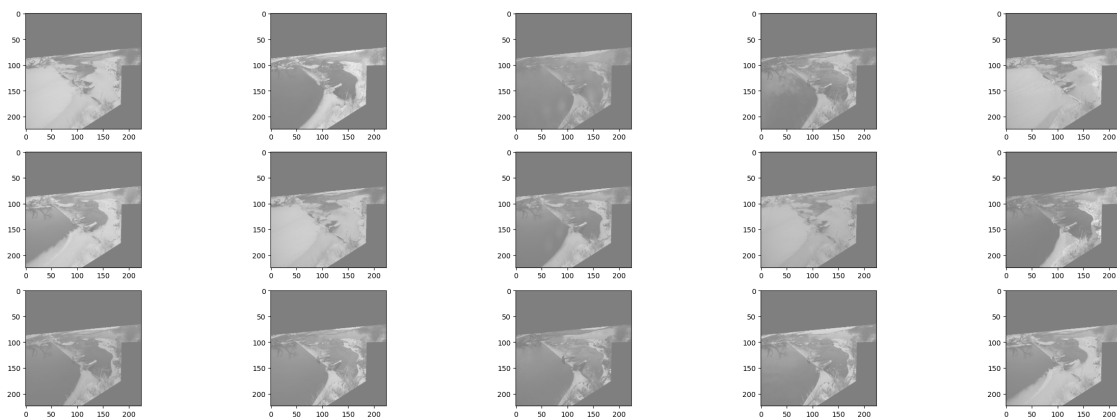
```
(224, 224, 3)
(2,)
```

### 0.3 Check images

```
[ ]: fig, ax = plt.subplots(nrows=3, ncols=5, figsize=(30, 10))

for image, stage_discharge in test_ds.take(1):
    images = image[:15]
    for img, ax in zip(images, ax.flatten()):
        #img = img.numpy()[0]
        img = img.numpy()
        img = img / 2 + 0.5      # unnormalize
        ax.imshow(img)

plt.show()
```



## 0.4 Create model

```
[ ]: def create_model(input_shape, output_shape, option="normal"):
    model = Sequential()

    if option == "transfer":
        base_model = tf.keras.applications.ResNet50V2(include_top=False,
                                                    weights='imagenet',
                                                    input_shape=input_shape)

        base_model.trainable = False
        base_model._name = 'base_model_ResNet50'

        model.add(base_model)
        model.add(Dropout(0.3))
        model.add(GlobalAveragePooling2D())

        model.add(Dense(512, activation='elu'))
        model.add(Dense(512, activation='elu'))
        model.add(Dense(256, activation='elu'))
        model.add(Dense(128, activation='elu'))
    elif option == "normal":
        model.add(Input(shape=input_shape))

        """model.add(Conv2D(16, kernel_size=(3, 3), activation="elu",
        ↪padding='same', kernel_initializer='he_uniform'))
        model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
        model.add(BatchNormalization())

        model.add(Conv2D(32, kernel_size=(3, 3), activation="elu",
        ↪padding='same', kernel_initializer='he_uniform'))
        model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
        model.add(BatchNormalization())

        model.add(Conv2D(32, kernel_size=(3, 3), activation="elu",
        ↪padding='same', kernel_initializer='he_uniform'))
        model.add(Conv2D(32, kernel_size=(3, 3), activation="elu",
        ↪padding='same', kernel_initializer='he_uniform'))
        model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
        model.add(BatchNormalization())

        model.add(Conv2D(64, kernel_size=(4, 4), activation="elu",
        ↪padding='same', kernel_initializer='he_uniform'))
        model.add(Conv2D(64, kernel_size=(4, 4), activation="elu",
        ↪padding='same', kernel_initializer='he_uniform'))
        model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
        model.add(BatchNormalization())
```



```

        model.add(Conv2D(64, kernel_size=(4, 4), activation="elu",
↳padding='same', kernel_initializer='he_uniform'))
        model.add(Conv2D(64, kernel_size=(4, 4), activation="elu",
↳padding='same', kernel_initializer='he_uniform'))
        model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
        model.add(BatchNormalization())

        model.add(Conv2D(64, kernel_size=(3, 3), activation="elu",
↳padding='same', kernel_initializer='he_uniform'))
        model.add(Conv2D(64, kernel_size=(3, 3), activation="elu",
↳padding='same', kernel_initializer='he_uniform'))
        model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
        model.add(BatchNormalization())

        model.add(GlobalAveragePooling2D())

        model.add(Dense(512, activation='elu'))
        model.add(Dropout(0.3))
        model.add(Dense(512, activation='elu'))
        model.add(Dropout(0.3))
        model.add(Dense(256, activation='elu'))
        model.add(Dense(64, activation='elu'))"""

        model.add(Conv2D(32, kernel_size=(4, 4), strides=(2, 2),
↳padding='same', activation="elu"))
        model.add(MaxPooling2D(pool_size=(2, 2)))

        model.add(Conv2D(32, kernel_size=(4, 4), strides=(2, 2),
↳activation="elu", padding='same'))
        model.add(MaxPooling2D(pool_size=(2, 2)))

        model.add(Conv2D(64, kernel_size=(3, 3), activation="elu",
↳padding='same'))
        #model.add(AveragePooling2D(pool_size=(2, 2)))

        model.add(Conv2D(64, kernel_size=(3, 3), activation='elu'))

        model.add(Conv2D(64, kernel_size=(2, 2), activation='elu'))
        model.add(MaxPooling2D(pool_size=(2, 2)))

        model.add(Conv2D(64, kernel_size=(3, 3), activation='elu'))

        model.add(Conv2D(64, kernel_size=(2, 2), activation='elu'))
        model.add(MaxPooling2D(pool_size=(2, 2)))

        model.add(Flatten())

```

```

        model.add(Dense(256, activation='tanh'))
        model.add(Dropout(0.3))
        model.add(Dense(128, activation='tanh'))
        model.add(Dense(64, activation='tanh'))
        model.add(Dense(32, activation='tanh'))
    elif option == "cnn/lstm":
        model.add(Input(shape=input_shape))

        model.add(TimeDistributed(Conv2D(32, kernel_size=(4, 4), strides=(2, 2), padding='same', activation=LeakyReLU()))
        model.add(TimeDistributed(MaxPooling2D(pool_size=(4, 4))))

        model.add(TimeDistributed(Conv2D(32, kernel_size=(4, 4), strides=(2, 2), padding='same', activation=LeakyReLU()))
        model.add(TimeDistributed(MaxPooling2D(pool_size=(2, 2))))

        model.add(TimeDistributed(Conv2D(32, kernel_size=(3, 3), padding='same', activation=LeakyReLU(0.2))))
        model.add(TimeDistributed(AveragePooling2D(pool_size=(2, 2))))

        model.add(TimeDistributed(GlobalAveragePooling2D()))

        model.add(LSTM(10))

        model.add(Dense(64, activation='tanh'))
        model.add(Dense(32, activation='tanh'))

        model.add(Dense(output_shape, activation='linear')) # linear regression
        # output layer

    return model

```

```

[ ]: model = create_model(input_shape, output_shape[0], "normal")
      #model = create_model(input_shape, output_shape[1], "cnn/lstm")

```

```

[ ]: model.summary()

```

Model: "sequential\_6"

Layer (type)	Output Shape	Param #
conv2d_32 (Conv2D)	(None, 112, 112, 32)	1568
max_pooling2d_14 (MaxPooling2D)	(None, 56, 56, 32)	0

conv2d_33 (Conv2D)	(None, 28, 28, 32)	16416
max_pooling2d_15 (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_34 (Conv2D)	(None, 14, 14, 64)	18496
conv2d_35 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_36 (Conv2D)	(None, 11, 11, 64)	16448
max_pooling2d_16 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_37 (Conv2D)	(None, 3, 3, 64)	36928
conv2d_38 (Conv2D)	(None, 2, 2, 64)	16448
max_pooling2d_17 (MaxPooling2D)	(None, 1, 1, 64)	0
flatten_4 (Flatten)	(None, 64)	0
dense_20 (Dense)	(None, 256)	16640
dropout_4 (Dropout)	(None, 256)	0
dense_21 (Dense)	(None, 128)	32896
dense_22 (Dense)	(None, 64)	8256
dense_23 (Dense)	(None, 32)	2080
dense_24 (Dense)	(None, 2)	66

```
=====
Total params: 203,170
Trainable params: 203,170
Non-trainable params: 0
-----
```

```
[ ]: def compile_model(loss_func, optimizer, metrics=["accuracy"]):
      model.compile(loss=loss_func, optimizer=optimizer, metrics=metrics)
```

```
[ ]: sgd = SGD(learning_rate=0.01, decay=1e-3, momentum=0.9, nesterov=True)
      adam = Adam(learning_rate=1e-3, decay=1e-3 / 200)
```

```
compile_model('mse', sgd, [
    'mse', tf.keras.metrics.RootMeanSquaredError(name='rmse'), 'mae',
    'mape'])
```

```
[ ]: def fit_model(training_values, validation_values=None, epochs=10, steps=32,
    val_steps=32, callbacks=[]):
    return model.fit(training_values, validation_data=validation_values,
    epochs=epochs, steps_per_epoch=steps, validation_steps=val_steps,
    callbacks=callbacks)
```

```
[ ]: import datetime

date_actual = datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
log_dir = "logs/fit/" + date_actual
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir,
    histogram_freq=1)

es_callback = tf.keras.callbacks.EarlyStopping(monitor='val_loss', mode='min',
    verbose=1, patience=5)

checkpoint_callback = tf.keras.callbacks.
    ModelCheckpoint(filepath=f"model_weights/{date_actual}_cnn_best_weights.
    hdf5",
                    monitor='val_loss',
                    verbose=1,
                    save_best_only=True)
```

```
[ ]: # batch_size = 0 because we already have batch size in tf dataset
model_h = fit_model(train_ds, val_ds, epochs=60, steps=np.ceil(train_size /
    BATCH_SIZE), val_steps=np.ceil(val_size / BATCH_SIZE),
    callbacks=[tensorboard_callback, checkpoint_callback, es_callback])
```

```
Epoch 1/60
635/635 [=====] - ETA: 0s - loss: 0.1610 - mse: 0.1610
- rmse: 0.4013 - mae: 0.2118 - mape: 68.5414
Epoch 1: val_loss improved from inf to 0.04874, saving model to
model_weights/20221030-202000_cnn_best_weights.hdf5
635/635 [=====] - 31s 48ms/step - loss: 0.1610 - mse:
0.1610 - rmse: 0.4013 - mae: 0.2118 - mape: 68.5414 - val_loss: 0.0487 -
val_mse: 0.0487 - val_rmse: 0.2208 - val_mae: 0.1577 - val_mape: 41.4325
Epoch 2/60
635/635 [=====] - ETA: 0s - loss: 0.0190 - mse: 0.0190
- rmse: 0.1379 - mae: 0.0948 - mape: 38.0835
Epoch 2: val_loss improved from 0.04874 to 0.03504, saving model to
model_weights/20221030-202000_cnn_best_weights.hdf5
635/635 [=====] - 27s 43ms/step - loss: 0.0190 - mse:
0.0190 - rmse: 0.1379 - mae: 0.0948 - mape: 38.0835 - val_loss: 0.0350 -
```

val\_mse: 0.0350 - val\_rmse: 0.1872 - val\_mae: 0.1340 - val\_mape: 39.6013  
Epoch 3/60  
634/635 [=====>.] - ETA: 0s - loss: 0.0135 - mse: 0.0135  
- rmse: 0.1162 - mae: 0.0801 - mape: 31.9182  
Epoch 3: val\_loss improved from 0.03504 to 0.03017, saving model to  
model\_weights/20221030-202000\_cnn\_best\_weights.hdf5  
635/635 [=====] - 26s 41ms/step - loss: 0.0135 - mse:  
0.0135 - rmse: 0.1162 - mae: 0.0800 - mape: 31.9021 - val\_loss: 0.0302 -  
val\_mse: 0.0302 - val\_rmse: 0.1737 - val\_mae: 0.1270 - val\_mape: 39.2433  
Epoch 4/60  
635/635 [=====] - ETA: 0s - loss: 0.0115 - mse: 0.0115  
- rmse: 0.1071 - mae: 0.0739 - mape: 29.4948  
Epoch 4: val\_loss improved from 0.03017 to 0.02752, saving model to  
model\_weights/20221030-202000\_cnn\_best\_weights.hdf5  
635/635 [=====] - 26s 41ms/step - loss: 0.0115 - mse:  
0.0115 - rmse: 0.1071 - mae: 0.0739 - mape: 29.4948 - val\_loss: 0.0275 -  
val\_mse: 0.0275 - val\_rmse: 0.1659 - val\_mae: 0.1206 - val\_mape: 37.2144  
Epoch 5/60  
635/635 [=====] - ETA: 0s - loss: 0.0101 - mse: 0.0101  
- rmse: 0.1005 - mae: 0.0693 - mape: 27.3959  
Epoch 5: val\_loss did not improve from 0.02752  
635/635 [=====] - 26s 40ms/step - loss: 0.0101 - mse:  
0.0101 - rmse: 0.1005 - mae: 0.0693 - mape: 27.3959 - val\_loss: 0.0276 -  
val\_mse: 0.0276 - val\_rmse: 0.1662 - val\_mae: 0.1144 - val\_mape: 32.7270  
Epoch 6/60  
635/635 [=====] - ETA: 0s - loss: 0.0092 - mse: 0.0092  
- rmse: 0.0959 - mae: 0.0664 - mape: 26.6044  
Epoch 6: val\_loss improved from 0.02752 to 0.02501, saving model to  
model\_weights/20221030-202000\_cnn\_best\_weights.hdf5  
635/635 [=====] - 26s 40ms/step - loss: 0.0092 - mse:  
0.0092 - rmse: 0.0959 - mae: 0.0664 - mape: 26.6044 - val\_loss: 0.0250 -  
val\_mse: 0.0250 - val\_rmse: 0.1581 - val\_mae: 0.1112 - val\_mape: 33.5459  
Epoch 7/60  
635/635 [=====] - ETA: 0s - loss: 0.0085 - mse: 0.0085  
- rmse: 0.0921 - mae: 0.0642 - mape: 26.2814  
Epoch 7: val\_loss improved from 0.02501 to 0.02380, saving model to  
model\_weights/20221030-202000\_cnn\_best\_weights.hdf5  
635/635 [=====] - 26s 41ms/step - loss: 0.0085 - mse:  
0.0085 - rmse: 0.0921 - mae: 0.0642 - mape: 26.2814 - val\_loss: 0.0238 -  
val\_mse: 0.0238 - val\_rmse: 0.1543 - val\_mae: 0.1069 - val\_mape: 31.6758  
Epoch 8/60  
635/635 [=====] - ETA: 0s - loss: 0.0083 - mse: 0.0083  
- rmse: 0.0912 - mae: 0.0634 - mape: 25.5699  
Epoch 8: val\_loss improved from 0.02380 to 0.02371, saving model to  
model\_weights/20221030-202000\_cnn\_best\_weights.hdf5  
635/635 [=====] - 26s 41ms/step - loss: 0.0083 - mse:  
0.0083 - rmse: 0.0912 - mae: 0.0634 - mape: 25.5699 - val\_loss: 0.0237 -  
val\_mse: 0.0237 - val\_rmse: 0.1540 - val\_mae: 0.1072 - val\_mape: 31.8466

Epoch 9/60  
633/635 [=====>.] - ETA: 0s - loss: 0.0077 - mse: 0.0077  
- rmse: 0.0879 - mae: 0.0610 - mape: 25.2393  
Epoch 9: val\_loss improved from 0.02371 to 0.02269, saving model to  
model\_weights/20221030-202000\_cnn\_best\_weights.hdf5  
635/635 [=====] - 27s 42ms/step - loss: 0.0077 - mse:  
0.0077 - rmse: 0.0879 - mae: 0.0610 - mape: 25.2183 - val\_loss: 0.0227 -  
val\_mse: 0.0227 - val\_rmse: 0.1506 - val\_mae: 0.1030 - val\_mape: 29.6224  
Epoch 10/60  
635/635 [=====] - ETA: 0s - loss: 0.0077 - mse: 0.0077  
- rmse: 0.0878 - mae: 0.0608 - mape: 24.0997  
Epoch 10: val\_loss improved from 0.02269 to 0.02234, saving model to  
model\_weights/20221030-202000\_cnn\_best\_weights.hdf5  
635/635 [=====] - 26s 41ms/step - loss: 0.0077 - mse:  
0.0077 - rmse: 0.0878 - mae: 0.0608 - mape: 24.0997 - val\_loss: 0.0223 -  
val\_mse: 0.0223 - val\_rmse: 0.1495 - val\_mae: 0.1036 - val\_mape: 30.8130  
Epoch 11/60  
634/635 [=====>.] - ETA: 0s - loss: 0.0070 - mse: 0.0070  
- rmse: 0.0839 - mae: 0.0585 - mape: 23.9554  
Epoch 11: val\_loss improved from 0.02234 to 0.02147, saving model to  
model\_weights/20221030-202000\_cnn\_best\_weights.hdf5  
635/635 [=====] - 26s 41ms/step - loss: 0.0070 - mse:  
0.0070 - rmse: 0.0839 - mae: 0.0585 - mape: 23.9445 - val\_loss: 0.0215 -  
val\_mse: 0.0215 - val\_rmse: 0.1465 - val\_mae: 0.1007 - val\_mape: 29.5933  
Epoch 12/60  
635/635 [=====] - ETA: 0s - loss: 0.0070 - mse: 0.0070  
- rmse: 0.0837 - mae: 0.0579 - mape: 24.0591  
Epoch 12: val\_loss did not improve from 0.02147  
635/635 [=====] - 26s 41ms/step - loss: 0.0070 - mse:  
0.0070 - rmse: 0.0837 - mae: 0.0579 - mape: 24.0591 - val\_loss: 0.0217 -  
val\_mse: 0.0217 - val\_rmse: 0.1474 - val\_mae: 0.1014 - val\_mape: 29.3682  
Epoch 13/60  
635/635 [=====] - ETA: 0s - loss: 0.0068 - mse: 0.0068  
- rmse: 0.0823 - mae: 0.0573 - mape: 24.0166  
Epoch 13: val\_loss did not improve from 0.02147  
635/635 [=====] - 26s 41ms/step - loss: 0.0068 - mse:  
0.0068 - rmse: 0.0823 - mae: 0.0573 - mape: 24.0166 - val\_loss: 0.0220 -  
val\_mse: 0.0220 - val\_rmse: 0.1484 - val\_mae: 0.1023 - val\_mape: 29.7553  
Epoch 14/60  
633/635 [=====>.] - ETA: 0s - loss: 0.0067 - mse: 0.0067  
- rmse: 0.0818 - mae: 0.0569 - mape: 23.8327  
Epoch 14: val\_loss improved from 0.02147 to 0.02129, saving model to  
model\_weights/20221030-202000\_cnn\_best\_weights.hdf5  
635/635 [=====] - 26s 41ms/step - loss: 0.0067 - mse:  
0.0067 - rmse: 0.0819 - mae: 0.0570 - mape: 23.8016 - val\_loss: 0.0213 -  
val\_mse: 0.0213 - val\_rmse: 0.1459 - val\_mae: 0.1006 - val\_mape: 29.6195  
Epoch 15/60  
633/635 [=====>.] - ETA: 0s - loss: 0.0064 - mse: 0.0064

```

- rmse: 0.0801 - mae: 0.0559 - mape: 23.0364
Epoch 15: val_loss improved from 0.02129 to 0.02126, saving model to
model_weights/20221030-202000_cnn_best_weights.hdf5
635/635 [=====] - 27s 42ms/step - loss: 0.0064 - mse:
0.0064 - rmse: 0.0801 - mae: 0.0559 - mape: 23.0177 - val_loss: 0.0213 -
val_mse: 0.0213 - val_rmse: 0.1458 - val_mae: 0.1004 - val_mape: 28.8095
Epoch 16/60
635/635 [=====] - ETA: 0s - loss: 0.0064 - mse: 0.0064
- rmse: 0.0803 - mae: 0.0557 - mape: 22.5464
Epoch 16: val_loss did not improve from 0.02126
635/635 [=====] - 26s 41ms/step - loss: 0.0064 - mse:
0.0064 - rmse: 0.0803 - mae: 0.0557 - mape: 22.5464 - val_loss: 0.0214 -
val_mse: 0.0214 - val_rmse: 0.1464 - val_mae: 0.0997 - val_mape: 27.7094
Epoch 17/60
633/635 [=====>.] - ETA: 0s - loss: 0.0063 - mse: 0.0063
- rmse: 0.0793 - mae: 0.0552 - mape: 23.4435
Epoch 17: val_loss did not improve from 0.02126
635/635 [=====] - 26s 40ms/step - loss: 0.0063 - mse:
0.0063 - rmse: 0.0794 - mae: 0.0552 - mape: 23.4109 - val_loss: 0.0216 -
val_mse: 0.0216 - val_rmse: 0.1471 - val_mae: 0.1006 - val_mape: 28.6674
Epoch 18/60
633/635 [=====>.] - ETA: 0s - loss: 0.0063 - mse: 0.0063
- rmse: 0.0791 - mae: 0.0547 - mape: 22.6994
Epoch 18: val_loss did not improve from 0.02126
635/635 [=====] - 26s 40ms/step - loss: 0.0063 - mse:
0.0063 - rmse: 0.0791 - mae: 0.0547 - mape: 22.6669 - val_loss: 0.0221 -
val_mse: 0.0221 - val_rmse: 0.1486 - val_mae: 0.1013 - val_mape: 27.8553
Epoch 19/60
635/635 [=====] - ETA: 0s - loss: 0.0061 - mse: 0.0061
- rmse: 0.0783 - mae: 0.0546 - mape: 22.7656
Epoch 19: val_loss improved from 0.02126 to 0.02073, saving model to
model_weights/20221030-202000_cnn_best_weights.hdf5
635/635 [=====] - 26s 42ms/step - loss: 0.0061 - mse:
0.0061 - rmse: 0.0783 - mae: 0.0546 - mape: 22.7656 - val_loss: 0.0207 -
val_mse: 0.0207 - val_rmse: 0.1440 - val_mae: 0.0990 - val_mape: 28.4503
Epoch 20/60
632/635 [=====>.] - ETA: 0s - loss: 0.0059 - mse: 0.0059
- rmse: 0.0770 - mae: 0.0535 - mape: 22.1829
Epoch 20: val_loss improved from 0.02073 to 0.02018, saving model to
model_weights/20221030-202000_cnn_best_weights.hdf5
635/635 [=====] - 26s 41ms/step - loss: 0.0059 - mse:
0.0059 - rmse: 0.0770 - mae: 0.0535 - mape: 22.1841 - val_loss: 0.0202 -
val_mse: 0.0202 - val_rmse: 0.1421 - val_mae: 0.0969 - val_mape: 27.7550
Epoch 21/60
635/635 [=====] - ETA: 0s - loss: 0.0060 - mse: 0.0060
- rmse: 0.0773 - mae: 0.0537 - mape: 22.4613
Epoch 21: val_loss did not improve from 0.02018
635/635 [=====] - 26s 41ms/step - loss: 0.0060 - mse:

```

0.0060 - rmse: 0.0773 - mae: 0.0537 - mape: 22.4613 - val\_loss: 0.0209 -  
val\_mse: 0.0209 - val\_rmse: 0.1446 - val\_mae: 0.0990 - val\_mape: 27.7635  
Epoch 22/60  
635/635 [=====] - ETA: 0s - loss: 0.0059 - mse: 0.0059  
- rmse: 0.0766 - mae: 0.0533 - mape: 22.1828  
Epoch 22: val\_loss did not improve from 0.02018  
635/635 [=====] - 26s 41ms/step - loss: 0.0059 - mse:  
0.0059 - rmse: 0.0766 - mae: 0.0533 - mape: 22.1828 - val\_loss: 0.0206 -  
val\_mse: 0.0206 - val\_rmse: 0.1437 - val\_mae: 0.0978 - val\_mape: 27.7665  
Epoch 23/60  
635/635 [=====] - ETA: 0s - loss: 0.0058 - mse: 0.0058  
- rmse: 0.0764 - mae: 0.0530 - mape: 22.2552  
Epoch 23: val\_loss did not improve from 0.02018  
635/635 [=====] - 26s 41ms/step - loss: 0.0058 - mse:  
0.0058 - rmse: 0.0764 - mae: 0.0530 - mape: 22.2552 - val\_loss: 0.0207 -  
val\_mse: 0.0207 - val\_rmse: 0.1440 - val\_mae: 0.0983 - val\_mape: 27.4993  
Epoch 24/60  
635/635 [=====] - ETA: 0s - loss: 0.0058 - mse: 0.0058  
- rmse: 0.0760 - mae: 0.0528 - mape: 21.7594  
Epoch 24: val\_loss improved from 0.02018 to 0.01967, saving model to  
model\_weights/20221030-202000\_cnn\_best\_weights.hdf5  
635/635 [=====] - 26s 41ms/step - loss: 0.0058 - mse:  
0.0058 - rmse: 0.0760 - mae: 0.0528 - mape: 21.7594 - val\_loss: 0.0197 -  
val\_mse: 0.0197 - val\_rmse: 0.1403 - val\_mae: 0.0958 - val\_mape: 27.2156  
Epoch 25/60  
634/635 [=====>.] - ETA: 0s - loss: 0.0058 - mse: 0.0058  
- rmse: 0.0763 - mae: 0.0526 - mape: 21.9381  
Epoch 25: val\_loss did not improve from 0.01967  
635/635 [=====] - 25s 40ms/step - loss: 0.0058 - mse:  
0.0058 - rmse: 0.0763 - mae: 0.0526 - mape: 21.9236 - val\_loss: 0.0201 -  
val\_mse: 0.0201 - val\_rmse: 0.1417 - val\_mae: 0.0969 - val\_mape: 27.2769  
Epoch 26/60  
632/635 [=====>.] - ETA: 0s - loss: 0.0056 - mse: 0.0056  
- rmse: 0.0747 - mae: 0.0520 - mape: 22.2876  
Epoch 26: val\_loss did not improve from 0.01967  
635/635 [=====] - 26s 41ms/step - loss: 0.0056 - mse:  
0.0056 - rmse: 0.0747 - mae: 0.0520 - mape: 22.2408 - val\_loss: 0.0213 -  
val\_mse: 0.0213 - val\_rmse: 0.1461 - val\_mae: 0.0996 - val\_mape: 27.1532  
Epoch 27/60  
635/635 [=====] - ETA: 0s - loss: 0.0056 - mse: 0.0056  
- rmse: 0.0749 - mae: 0.0520 - mape: 22.2381  
Epoch 27: val\_loss did not improve from 0.01967  
635/635 [=====] - 29s 45ms/step - loss: 0.0056 - mse:  
0.0056 - rmse: 0.0749 - mae: 0.0520 - mape: 22.2381 - val\_loss: 0.0201 -  
val\_mse: 0.0201 - val\_rmse: 0.1419 - val\_mae: 0.0970 - val\_mape: 27.4242  
Epoch 28/60  
635/635 [=====] - ETA: 0s - loss: 0.0054 - mse: 0.0054  
- rmse: 0.0737 - mae: 0.0515 - mape: 21.9319



```

Epoch 28: val_loss did not improve from 0.01967
635/635 [=====] - 27s 42ms/step - loss: 0.0054 - mse:
0.0054 - rmse: 0.0737 - mae: 0.0515 - mape: 21.9319 - val_loss: 0.0201 -
val_mse: 0.0201 - val_rmse: 0.1419 - val_mae: 0.0971 - val_mape: 27.1583
Epoch 29/60
635/635 [=====] - ETA: 0s - loss: 0.0055 - mse: 0.0055
- rmse: 0.0743 - mae: 0.0516 - mape: 21.7428
Epoch 29: val_loss did not improve from 0.01967
635/635 [=====] - 26s 41ms/step - loss: 0.0055 - mse:
0.0055 - rmse: 0.0743 - mae: 0.0516 - mape: 21.7428 - val_loss: 0.0203 -
val_mse: 0.0203 - val_rmse: 0.1424 - val_mae: 0.0973 - val_mape: 27.1245
Epoch 29: early stopping

```

## 0.5 Evaluate model

```

[ ]: print(date_actual)

20221030-202000

[ ]: best_model = models.load_model(f'model_weights/{date_actual}_cnn_best_weights.
    ↪hdf5')
    #best_model = models.load_model(f'best_models_weights/cnn_best_weights_v9.hdf5')

[ ]: def evaluate_model(model, test_values, steps):
    score = model.evaluate(test_values, steps=steps)
    return score

[ ]: test_loss, test_mse, test_rmse, test_mae, test_mape = ↪
    ↪evaluate_model(best_model, test_ds, steps=np.ceil(test_size / BATCH_SIZE))

398/398 [=====] - 12s 30ms/step - loss: 0.0505 - mse:
0.0505 - rmse: 0.2246 - mae: 0.1389 - mape: 72.7527

[ ]: #predictions = best_model.predict(test_ds, steps=np.ceil(test_size / ↪
    ↪BATCH_SIZE))

[ ]: for image, stage_discharge in test_ds.take(1):
    predictions = best_model.predict(x=image)

    stage_discharge_test_values = stage_discharge.numpy()
    predictions_values = predictions

    diff = predictions_values.flatten() - stage_discharge_test_values.
    ↪flatten()
    percentDiff = (diff / stage_discharge_test_values.flatten()) * 100
    absPercentDiff = np.abs(percentDiff)
    # compute the mean and standard deviation of the absolute percentage
    # difference

```

```

mean = np.mean(absPercentDiff)
std = np.std(absPercentDiff)
# finally, show some statistics on our model
print(mean)
print(std)

stage_discharge_test_values = stage_discharge[:10]
predictions_values = predictions[:10]

for i in range(len(stage_discharge_test_values.numpy())):
    print(f"pred stage: {scaler.
→inverse_transform(predictions_values)[i][0]}, actual stage: {scaler.
→inverse_transform(stage_discharge_test_values)[i][0]}")
    print(f"pred discharge: {scaler.
→inverse_transform(predictions_values)[i][1]}, actual discharge: {scaler.
→inverse_transform(stage_discharge_test_values)[i][1]}")

```

```

1/1 [=====] - 0s 88ms/step
441.19266182980726
3243.6085036341224
pred stage: 2.074333429336548, actual stage: 2.42
pred discharge: 267.8599853515625, actual discharge: 210.0
pred stage: 2.1716480255126953, actual stage: 2.24
pred discharge: 227.58538818359375, actual discharge: 193.99999999999999
pred stage: 2.097628116607666, actual stage: 2.24
pred discharge: 134.73583984375, actual discharge: 193.99999999999999
pred stage: 1.9582111835479736, actual stage: 2.46
pred discharge: 116.59221649169922, actual discharge: 209.0
pred stage: 2.164323091506958, actual stage: 2.24
pred discharge: 195.36407470703125, actual discharge: 193.99999999999999
pred stage: 2.0977611541748047, actual stage: 2.25
pred discharge: 154.81903076171875, actual discharge: 198.0
pred stage: 2.106117010116577, actual stage: 2.25
pred discharge: 164.81689453125, actual discharge: 199.0
pred stage: 2.0872802734375, actual stage: 2.25
pred discharge: 146.6630859375, actual discharge: 198.0
pred stage: 2.1385679244995117, actual stage: 2.26
pred discharge: 154.3438720703125, actual discharge: 204.0
pred stage: 2.1561810970306396, actual stage: 2.26
pred discharge: 170.2928466796875, actual discharge: 204.0

```

### 0.5.1 Residual analysis

```

[ ]: y_predictions = np.empty(shape=(1, 2))
     y_real = np.empty(shape=(1, 2))

"""for image, stage_discharge in test_ds.take(100):

```

```

y_predictions = np.concatenate((y_predictions, best_model.predict(x=image)))
y_real = np.concatenate((y_real, stage_discharge.numpy()))"""

```

```

[ ]: 'for image, stage_discharge in test_ds.take(100):\n    y_predictions =
np.concatenate((y_predictions, best_model.predict(x=image)))\n    y_real =
np.concatenate((y_real, stage_discharge.numpy()))'

```

```

[ ]: residuals = y_real - y_predictions
residuals_std = residuals/residuals.std()

y_real_stage = np.array([i[0] for i in y_real])
residual_stage = np.array([i[0] for i in residuals])

y_real_discharge = np.array([i[1] for i in y_real])
residual_discharge = np.array([i[1] for i in residuals])

plt.scatter(y_real_stage, residual_stage / residual_stage.std(), label="stage_
↳residuals")
plt.scatter(y_real_discharge, residual_discharge / residual_discharge.std(),
↳label="discharge residuals")
plt.axhline(y=0.0, color='r', linestyle='-')
plt.xlabel("Fitted values")
plt.ylabel("Standarized residuals")

plt.legend()
plt.show()

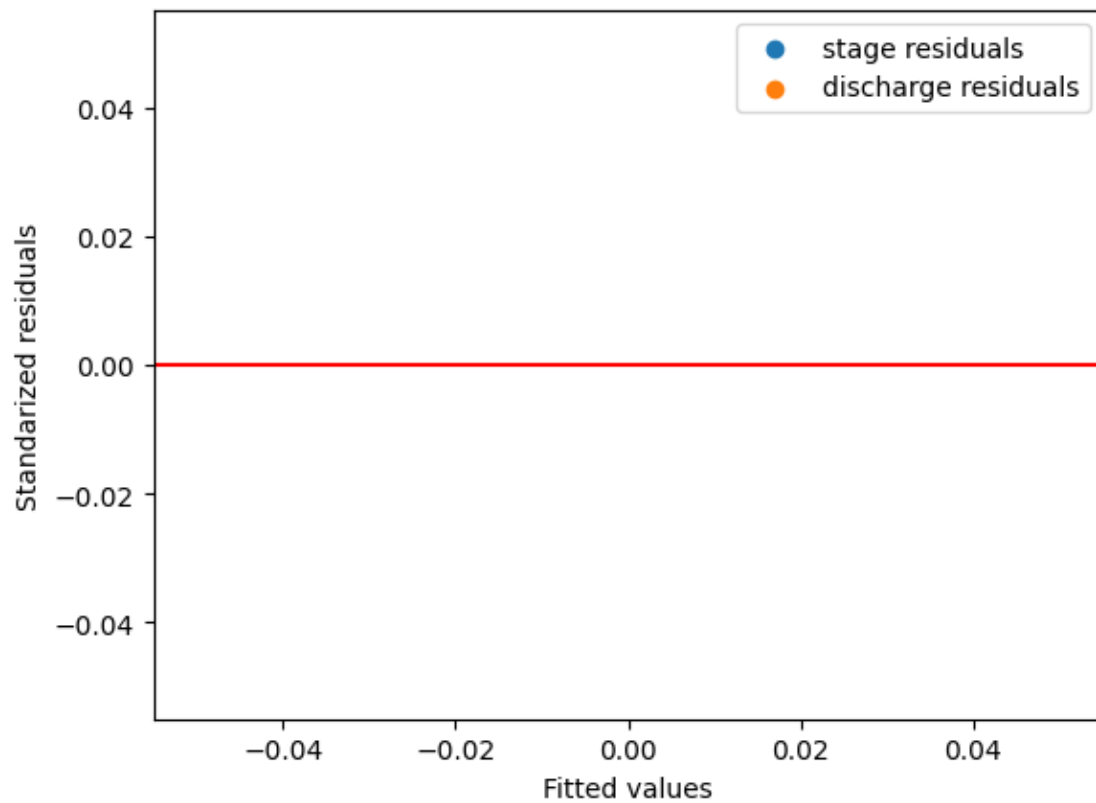
```

/tmp/ipykernel\_23436/3264406076.py:11: RuntimeWarning: divide by zero encountered in divide

```


plt.scatter(y_real_discharge, residual_discharge / residual_discharge.std(),
label="discharge residuals")

```



```
[ ]: import statsmodels.api as sm
      from statsmodels.stats.diagnostic import normal_ad


      figure = sm.qqplot(residual_stage / residual_stage.std(), line='45',
        ↪label='stage')
      plt.show()
```



```
[ ]: figure = sm.qqplot(residual_discharge / residual_discharge.std(), line='45',  
    ↪label='discharge')  
plt.show()
```

/tmp/ipykernel\_23436/1192211388.py:1: RuntimeWarning: divide by zero encountered  
in divide

```
    figure = sm.qqplot(residual_discharge / residual_discharge.std(), line='45',  
label='discharge')
```



```
[ ]: import seaborn as sns

#sns.histplot(residuals, kde=True, bins = 10)
```

```
[ ]: stat, pval = normal_ad(residual_stage)
print("p-value:", pval)

if pval<0.05:
    print("Hay evidencia de que los residuos no provienen de una distribución ↵
    ↵normal.")
else:
    print("No hay evidencia para rechazar la hipótesis de que los residuos ↵
    ↵vienen de una distribución normal.")
```

p-value: 0.0

Hay evidencia de que los residuos no provienen de una distribución normal.

/home/nkspartan/miniconda3/envs/tf-gpu/lib/python3.10/site-packages/numpy/core/\_methods.py:265: RuntimeWarning: Degrees of freedom <= 0 for slice

```
ret = _var(a, axis=axis, dtype=dtype, out=out, ddof=ddof,
```

```
[ ]: stat, pval = normal_ad(residual_discharge)
print("p-value:", pval)

if pval < 0.05:
    print("Hay evidencia de que los residuos no provienen de una distribución
    ↪normal.")
else:
    print("No hay evidencia para rechazar la hipótesis de que los residuos
    ↪vienen de una distribución normal.")
```

p-value: 0.0

Hay evidencia de que los residuos no provienen de una distribución normal.

```
/home/nkspartan/miniconda3/envs/tf-gpu/lib/python3.10/site-
packages/numpy/core/_methods.py:257: RuntimeWarning: invalid value encountered
in double_scalars
    ret = ret.dtype.type(ret / rcount)
```

## 0.6 Visualize layers

```
[ ]: layer_outputs = [layer.output for layer in best_model.layers[:12]]
# Extracts the outputs of the top 12 layers
activation_model = models.Model(inputs=best_model.input, outputs=layer_outputs)
    ↪# Creates a model that will return these outputs, given the model input
```

```
[ ]: activations = activation_model.predict(test_ds.take(1))
```

1/1 [=====] - 0s 212ms/step

```
[ ]: import matplotlib.pyplot as plt

layer_names = []
for layer in best_model.layers[:12]:
    layer_names.append(layer.name) # Names of the layers, so you can have them
    ↪as part of your plot

images_per_row = 16

for layer_name, layer_activation in zip(layer_names, activations): # Displays
    ↪the feature maps
    n_features = layer_activation.shape[-1] # Number of features in the feature
    ↪map
    size = layer_activation.shape[1] #The feature map has shape (1, size, size,
    ↪n_features).
    n_cols = n_features // images_per_row # Tiles the activation channels in
    ↪this matrix
    display_grid = np.zeros((size * n_cols, images_per_row * size))
```

```

print(layer_name)
if "flatten" in layer_name or "dense" in layer_name: break

for col in range(n_cols): # Tiles each filter into a big horizontal grid
    for row in range(images_per_row):
        channel_image = layer_activation[0,
                                         :, :,
                                         col * images_per_row + row]
        channel_image -= channel_image.mean() # Post-processes the feature
↳ to make it visually palatable
        channel_image /= channel_image.std()
        channel_image *= 64
        channel_image += 128
        channel_image = np.clip(channel_image, 0, 255).astype('uint8')
        display_grid[col * size : (col + 1) * size, # Displays the grid
                     row * size : (row + 1) * size] = channel_image

scale = 1. / size
plt.figure(figsize=(scale * display_grid.shape[1],
                    scale * display_grid.shape[0]))
plt.title(layer_name)
plt.grid(False)
plt.imshow(display_grid, aspect='auto', cmap='viridis')

```

```

conv2d_32
max_pooling2d_14
conv2d_33
max_pooling2d_15
conv2d_34
conv2d_35
conv2d_36
max_pooling2d_16
conv2d_37
conv2d_38
max_pooling2d_17
flatten_4

```

/tmp/ipykernel\_23436/1031377702.py:24: RuntimeWarning: invalid value encountered in divide

```
channel_image /= channel_image.std()
```

