

cnn_v12

November 4, 2022

```
[ ]: %env LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$CONDA_PREFIX/lib/  
      #%env TF_GPU_ALLOCATOR=cuda_malloc_async
```

```
env: LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$CONDA_PREFIX/lib/
```

```
[ ]: import os  
      print(os.environ["LD_LIBRARY_PATH"])
```

```
$LD_LIBRARY_PATH:$CONDA_PREFIX/lib/
```

```
[ ]: import tensorflow as tf  
      import numpy as np  
      import pandas as pd  
      import os  
      import matplotlib.pyplot as plt  
  
      from tensorflow.keras import Sequential, models, Input  
      from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D,  
      ↪Dropout, LeakyReLU, AveragePooling2D, GlobalAveragePooling2D,  
      ↪BatchNormalization, TimeDistributed, LSTM, SpatialDropout2D  
      from tensorflow.keras.optimizers import SGD, Adam
```

```
2022-11-03 22:13:21.588312: I tensorflow/core/platform/cpu_feature_guard.cc:193]  
This TensorFlow binary is optimized with oneAPI Deep Neural Network Library  
(oneDNN) to use the following CPU instructions in performance-critical  
operations:  AVX2 FMA
```

```
To enable them in other operations, rebuild TensorFlow with the appropriate  
compiler flags.
```

```
2022-11-03 22:13:22.142775: E tensorflow/stream_executor/cuda/cuda_blas.cc:2981]  
Unable to register cuBLAS factory: Attempting to register factory for plugin  
cuBLAS when one has already been registered
```

```
2022-11-03 22:13:23.284245: W
```

```
tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load  
dynamic library 'libnvinfer.so.7'; dlopen: libnvinfer.so.7: cannot open shared  
object file: No such file or directory; LD_LIBRARY_PATH:  
:/home/nkspartan/miniconda3/envs/tf-gpu/lib/:/home/nkspartan/miniconda3/envs/tf-  
gpu/lib/
```

```
2022-11-03 22:13:23.284487: W
```

```
tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load
dynamic library 'libnvinfer_plugin.so.7'; dlopen: libnvinfer_plugin.so.7:
cannot open shared object file: No such file or directory; LD_LIBRARY_PATH:
:/home/nkspartan/miniconda3/envs/tf-gpu/lib/;/home/nkspartan/miniconda3/envs/tf-
gpu/lib/
2022-11-03 22:13:23.284497: W
tensorflow/compiler/tf2tensorrt/utils/py_utils.cc:38] TF-TRT Warning: Cannot
dlopen some TensorRT libraries. If you would like to use Nvidia GPU with
TensorRT, please make sure the missing libraries mentioned above are installed
properly.
```

```
[ ]: from tensorflow.python.client import device_lib

print('Default GPU Device: {}'.format(tf.test.gpu_device_name()))
```

```
Default GPU Device: /device:GPU:0
```

```
2022-11-03 22:13:24.971054: I tensorflow/core/platform/cpu_feature_guard.cc:193]
This TensorFlow binary is optimized with oneAPI Deep Neural Network Library
(oneDNN) to use the following CPU instructions in performance-critical
operations:  AVX2 FMA
```

```
To enable them in other operations, rebuild TensorFlow with the appropriate
compiler flags.
```

```
2022-11-03 22:13:25.006076: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
```

```
2022-11-03 22:13:25.050066: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
```

```
2022-11-03 22:13:25.050275: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
```

```
2022-11-03 22:13:25.934775: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
```

```
2022-11-03 22:13:25.935857: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
```

```
2022-11-03 22:13:25.936081: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
```

```
2022-11-03 22:13:25.936259: I
```

```
tensorflow/core/common_runtime/gpu/gpu_device.cc:1616] Created device
/device:GPU:0 with 4008 MB memory: -> device: 0, name: NVIDIA GeForce RTX 2060,
pci bus id: 0000:08:00.0, compute capability: 7.5
```

```
[ ]: physical_devices = tf.config.list_physical_devices('GPU')
tf.config.experimental.set_memory_growth(physical_devices[0], True)
```

```
2022-11-03 22:13:25.991236: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-11-03 22:13:25.992002: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-11-03 22:13:25.992181: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
```

0.1 Read the csv dataset to get the values for stage and discharge of the images

```
[ ]: df = pd.read_csv("../dataset/2012_2019_PlatteRiverWeir_features_merged_all.
    ↪ csv")
df.head()
```

```
[ ]: Unnamed: 0      SensorTime      CaptureTime \
0          0  2012-06-09 13:15:00  2012-06-09T13:09:07
1          1  2012-06-09 13:15:00  2012-06-09T13:10:29
2          2  2012-06-09 13:45:00  2012-06-09T13:44:01
3          3  2012-06-09 14:45:00  2012-06-09T14:44:30
4          4  2012-06-09 15:45:00  2012-06-09T15:44:59

      Filename Agency  SiteNumber TimeZone  Stage \
0  StateLineWeir_20120609_Farrell_001.jpg  USGS    6674500    MDT    2.99
1  StateLineWeir_20120609_Farrell_002.jpg  USGS    6674500    MDT    2.99
2  StateLineWeir_20120609_Farrell_003.jpg  USGS    6674500    MDT    2.96
3  StateLineWeir_20120609_Farrell_004.jpg  USGS    6674500    MDT    2.94
4  StateLineWeir_20120609_Farrell_005.jpg  USGS    6674500    MDT    2.94

      Discharge      CalcTimestamp  ...  WeirPt2X  WeirPt2Y  WwRawLineMin  \
0      916.0  2020-03-11T16:58:28  ...      -1      -1            0.0
1      916.0  2020-03-11T16:58:33  ...      -1      -1            0.0
2      873.0  2020-03-11T16:58:40  ...      -1      -1            0.0
3      846.0  2020-03-11T16:58:47  ...      -1      -1            0.0
4      846.0  2020-03-11T16:58:55  ...      -1      -1            0.0
```

	WwRawLineMax	WwRawLineMean	WwRawLineSigma	WwCurveLineMin	\
0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	

	WwCurveLineMax	WwCurveLineMean	WwCurveLineSigma
0	0.0	0.0	0.0
1	0.0	0.0	0.0
2	0.0	0.0	0.0
3	0.0	0.0	0.0
4	0.0	0.0	0.0

[5 rows x 60 columns]

```
[ ]: df = df[["Filename", "Stage", "Discharge", 'SensorTime']]
```

```
[ ]: df['SensorTime'] = pd.to_datetime(df['SensorTime'])
df['Year'] = df['SensorTime'].dt.year
df.head()
```

	Filename	Stage	Discharge	\
0	StateLineWeir_20120609_Farrell_001.jpg	2.99	916.0	
1	StateLineWeir_20120609_Farrell_002.jpg	2.99	916.0	
2	StateLineWeir_20120609_Farrell_003.jpg	2.96	873.0	
3	StateLineWeir_20120609_Farrell_004.jpg	2.94	846.0	
4	StateLineWeir_20120609_Farrell_005.jpg	2.94	846.0	

	SensorTime	Year
0	2012-06-09 13:15:00	2012
1	2012-06-09 13:15:00	2012
2	2012-06-09 13:45:00	2012
3	2012-06-09 14:45:00	2012
4	2012-06-09 15:45:00	2012

```
[ ]: df = df.sort_values(by="SensorTime", ascending=True)
df.head()
```

	Filename	Stage	Discharge	\
0	StateLineWeir_20120609_Farrell_001.jpg	2.99	916.0	
1	StateLineWeir_20120609_Farrell_002.jpg	2.99	916.0	
2	StateLineWeir_20120609_Farrell_003.jpg	2.96	873.0	
3	StateLineWeir_20120609_Farrell_004.jpg	2.94	846.0	
4	StateLineWeir_20120609_Farrell_005.jpg	2.94	846.0	

	SensorTime	Year
0	2012-06-09 13:15:00	2012
1	2012-06-09 13:15:00	2012
2	2012-06-09 13:45:00	2012
3	2012-06-09 14:45:00	2012
4	2012-06-09 15:45:00	2012

```

0 2012-06-09 13:15:00 2012
1 2012-06-09 13:15:00 2012
2 2012-06-09 13:45:00 2012
3 2012-06-09 14:45:00 2012
4 2012-06-09 15:45:00 2012

```

0.1.1 Remove outliers

```
[ ]: df = df[df.Stage > 0]
df = df[df.Discharge > 0]
```

We consider values equal to 0 as outliers because from the photos it doesn't seem that it would be possible that at this time we would have a value of 0 for stage or discharge

```
[ ]: df.shape
```

```
[ ]: (40148, 5)
```

0.1.2 Scale the data

```
[ ]: from sklearn.preprocessing import StandardScaler
from joblib import load

scaler = StandardScaler()
#scaler = load('std_scaler.joblib') # scaler with all the 42059 observations
```

Scale the data based only on the training dataset (in this case the training dataset is from 2012 to 2016)

```
[ ]: data_to_scale_fit = df[(df["Year"] >= 2012) & (df["Year"] <= 2016)][["Stage",
↪ "Discharge"]]
data_to_scale_fit
```

```
[ ]:
      Stage  Discharge
0      2.99      916.0
1      2.99      916.0
2      2.96      873.0
3      2.94      846.0
4      2.94      846.0
...
21416  2.38      279.0
21417  2.38      279.0
21418  2.38      279.0
21419  2.38      279.0
21420  2.38      279.0
```

```
[20304 rows x 2 columns]
```

```
[ ]: scaler.fit(data_to_scale_fit)
```

```
[ ]: StandardScaler()
```

```
[ ]: df[["Stage", "Discharge"]] = scaler.transform(df[["Stage", "Discharge"]])
df
```

```
[ ]:
```

	Filename	Stage	Discharge	\
0	StateLineWeir_20120609_Farrell_001.jpg	0.077964	-0.136077	
1	StateLineWeir_20120609_Farrell_002.jpg	0.077964	-0.136077	
2	StateLineWeir_20120609_Farrell_003.jpg	0.045759	-0.165451	
3	StateLineWeir_20120609_Farrell_004.jpg	0.024290	-0.183894	
4	StateLineWeir_20120609_Farrell_005.jpg	0.024290	-0.183894	
...	
42054	StateLineWeir_20191011_Farrell_409.jpg	-0.405103	-0.465332	
42055	StateLineWeir_20191011_Farrell_410.jpg	-0.405103	-0.465332	
42056	StateLineWeir_20191011_Farrell_411.jpg	-0.405103	-0.465332	
42057	StateLineWeir_20191011_Farrell_412.jpg	-0.405103	-0.465332	
42058	StateLineWeir_20191011_Farrell_413.jpg	-0.405103	-0.465332	

	SensorTime	Year
0	2012-06-09 13:15:00	2012
1	2012-06-09 13:15:00	2012
2	2012-06-09 13:45:00	2012
3	2012-06-09 14:45:00	2012
4	2012-06-09 15:45:00	2012
...
42054	2019-10-11 09:00:00	2019
42055	2019-10-11 10:00:00	2019
42056	2019-10-11 11:00:00	2019
42057	2019-10-11 12:00:00	2019
42058	2019-10-11 12:45:00	2019

[40148 rows x 5 columns]

```
[ ]: from joblib import dump
#dump(scaler, 'std_scaler_train_value_0_outliers.joblib')
```

0.2 Create the dataset pipeline

```
[ ]: IMG_SIZE = 224
#IMG_SIZE = 512
BATCH_SIZE = 32
FRAMES = 10
```

```
[ ]: from dataset_transformer import make_dataset
```

```
[ ]: path = "../../dataset/images_tmp_draw"

with tf.device("/gpu:0"):
    train_ds, train_size, val_ds, val_size, test_ds, test_size = \
        make_dataset(path, BATCH_SIZE, IMG_SIZE, FRAMES, df, 10, True, "cnn")
```

```
2022-11-03 22:13:27.304760: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-11-03 22:13:27.304964: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-11-03 22:13:27.305105: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-11-03 22:13:27.305466: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-11-03 22:13:27.305620: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-11-03 22:13:27.305737: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:1616] Created device
/job:localhost/replica:0/task:0/device:GPU:0 with 4008 MB memory: -> device: 0,
name: NVIDIA GeForce RTX 2060, pci bus id: 0000:08:00.0, compute capability: 7.5

20304
7117
12727
```

```
[ ]: input_shape = 0
output_shape = 0

for image, stage_discharge in train_ds.take(1):
    print(image.numpy().shape)
    print(stage_discharge.numpy().shape)

    input_shape = image.numpy().shape[1:]
    output_shape = stage_discharge.numpy().shape[1:]
```

```
(32, 224, 224, 3)
(32, 2)
```

```
[ ]: print(input_shape)
      print(output_shape)
```


```
(224, 224, 3)
(2,)
```

0.3 Check images

```
[ ]: fig, ax = plt.subplots(nrows=3, ncols=5, figsize=(30, 15))

for image, stage_discharge in test_ds.take(1):
    images = image[:15]
    for img, ax in zip(images, ax.flatten()):
        img = img.numpy()[0]
        #img = img.numpy()
        img = img / 2 + 0.5      # unnormalize
        ax.imshow(img)

plt.show()
```



0.4 Create model

```
[ ]: def create_model(input_shape, output_shape, option="normal"):
    model = Sequential()

    if option == "transfer":
        base_model = tf.keras.applications.ResNet50V2(include_top=False,
                                                    weights='imagenet',
                                                    input_shape=input_shape)

        for layer in base_model.layers:
            layer.trainable = False

        base_model._name = 'base_model_ResNet50'

        model.add(base_model)
        model.add(Dropout(0.3))
        model.add(GlobalAveragePooling2D())

        model.add(Dense(1024, activation='elu'))
        model.add(Dropout(0.3))
        model.add(Dense(512, activation='elu'))
        model.add(Dropout(0.3))
        model.add(Dense(256, activation='elu'))
        model.add(Dense(128, activation='elu'))
    elif option == "normal":
        model.add(Input(shape=input_shape))

        """model.add(Conv2D(16, kernel_size=(3, 3), activation="elu",
→padding='same', kernel_initializer='he_uniform'))
        model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
        model.add(BatchNormalization())

        model.add(Conv2D(32, kernel_size=(3, 3), activation="elu",
→padding='same', kernel_initializer='he_uniform'))
        model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
        model.add(BatchNormalization())

        model.add(Conv2D(32, kernel_size=(3, 3), activation="elu",
→padding='same', kernel_initializer='he_uniform'))
        model.add(Conv2D(32, kernel_size=(3, 3), activation="elu",
→padding='same', kernel_initializer='he_uniform'))
        model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
        model.add(BatchNormalization())

        model.add(Conv2D(64, kernel_size=(4, 4), activation="elu",
→padding='same', kernel_initializer='he_uniform'))
```

```

        model.add(Conv2D(64, kernel_size=(4, 4), activation="elu",
↳padding='same', kernel_initializer='he_uniform'))
        model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
        model.add(BatchNormalization())

        model.add(Conv2D(64, kernel_size=(4, 4), activation="elu",
↳padding='same', kernel_initializer='he_uniform'))
        model.add(Conv2D(64, kernel_size=(4, 4), activation="elu",
↳padding='same', kernel_initializer='he_uniform'))
        model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
        model.add(BatchNormalization())

        model.add(Conv2D(64, kernel_size=(3, 3), activation="elu",
↳padding='same', kernel_initializer='he_uniform'))
        model.add(Conv2D(64, kernel_size=(3, 3), activation="elu",
↳padding='same', kernel_initializer='he_uniform'))
        model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
        model.add(BatchNormalization())

        model.add(GlobalAveragePooling2D())

        model.add(Dense(512, activation='elu'))
        model.add(Dropout(0.3))
        model.add(Dense(512, activation='elu'))
        model.add(Dropout(0.3))
        model.add(Dense(256, activation='elu'))
        model.add(Dense(64, activation='elu'))"""

        model.add(Conv2D(32, kernel_size=(4, 4), strides=(2, 2),
↳padding='same', activation="elu"))
        model.add(MaxPooling2D(pool_size=(2, 2)))

        model.add(Conv2D(32, kernel_size=(4, 4), strides=(2, 2),
↳activation="elu", padding='same'))
        model.add(MaxPooling2D(pool_size=(2, 2)))

        model.add(Conv2D(64, kernel_size=(3, 3), activation="elu",
↳padding='same'))
        #model.add(AveragePooling2D(pool_size=(2, 2)))

        model.add(Conv2D(64, kernel_size=(3, 3), activation='elu'))

        model.add(Conv2D(64, kernel_size=(2, 2), activation='elu'))
        model.add(MaxPooling2D(pool_size=(2, 2)))

        model.add(Conv2D(64, kernel_size=(3, 3), activation='elu'))

```

```

model.add(Conv2D(64, kernel_size=(2, 2), activation='elu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())
model.add(Dense(256, activation='tanh'))
model.add(Dropout(0.3))
model.add(Dense(128, activation='tanh'))
model.add(Dense(64, activation='tanh'))
model.add(Dense(32, activation='tanh'))
elif option == "cnn/lstm":
    base_model = tf.keras.applications.ResNet50V2(include_top=False,
                                                    weights='imagenet',
                                                    input_shape=(224, 224, 3))

    base_model.trainable = False
    base_model._name = 'base_model_ResNet50'
    model.add(Input(shape=input_shape))
    model.add(TimeDistributed(base_model))
    model.add(TimeDistributed(Dropout(0.3)))
    model.add(TimeDistributed(GlobalAveragePooling2D()))

    """model.add(Input(shape=input_shape))

    model.add(TimeDistributed(Conv2D(16, kernel_size=(4, 4), strides=(2,
→2), padding='same', activation='elu'))))
    model.add(TimeDistributed(MaxPooling2D(pool_size=(3, 3))))

    model.add(TimeDistributed(Conv2D(32, kernel_size=(4, 4), strides=(2,
→2), activation='elu', padding='same'))))
    model.add(TimeDistributed(MaxPooling2D(pool_size=(2, 2))))

    model.add(TimeDistributed(Conv2D(32, kernel_size=(3, 3),
→activation='elu', padding='same'))))
    model.add(TimeDistributed(MaxPooling2D(pool_size=(2, 2))))

    model.add(TimeDistributed(Conv2D(32, kernel_size=(3, 3),
→activation='elu', padding='same'))))
    model.add(TimeDistributed(MaxPooling2D(pool_size=(2, 2))))

    model.add(TimeDistributed(Flatten()))"""

    model.add(LSTM(10, return_sequences=True))
    model.add(LSTM(15))

    model.add(Dense(512, activation='elu'))
    model.add(Dense(256, activation='elu'))

```

```

        model.add(Dense(128, activation='elu'))
        model.add(Dense(128, activation='elu'))

        model.add(Dense(output_shape, activation='linear')) # linear regression
        ↪ output layer

    return model

```

```
[ ]: model = create_model(input_shape, output_shape[0], "transfer")
```

```
[ ]: model.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
base_model_ResNet50 (Functional)	(None, 7, 7, 2048)	23564800
dropout_6 (Dropout)	(None, 7, 7, 2048)	0
global_average_pooling2d_2 (GlobalAveragePooling2D)	(None, 2048)	0
dense_10 (Dense)	(None, 1024)	2098176
dropout_7 (Dropout)	(None, 1024)	0
dense_11 (Dense)	(None, 512)	524800
dropout_8 (Dropout)	(None, 512)	0
dense_12 (Dense)	(None, 256)	131328
dense_13 (Dense)	(None, 128)	32896
dense_14 (Dense)	(None, 2)	258

=====
 Total params: 26,352,258
 Trainable params: 2,787,458
 Non-trainable params: 23,564,800
 =====

```
[ ]: def compile_model(loss_func, optimizer, metrics=["accuracy"]):
    model.compile(loss=loss_func, optimizer=optimizer, metrics=metrics)
```

```
[ ]: sgd = SGD(learning_rate=0.01, decay=1e-3, momentum=0.9, nesterov=True)
      #adam = Adam(learning_rate=1e-3, decay=1e-3 / 200)
      adam = Adam(learning_rate=1e-3, decay=1e-3 / 200)

      compile_model('mse', adam, [
          'mse', tf.keras.metrics.RootMeanSquaredError(name='rmse'), 'mae',
          ↪ 'mape'])
```

```
[ ]: def fit_model(training_values, validation_values=None, epochs=10, steps=32,
      ↪ val_steps=32, callbacks=[]):
      return model.fit(training_values, validation_data=validation_values,
      ↪ epochs=epochs, steps_per_epoch=steps, validation_steps=val_steps,
      ↪ callbacks=callbacks)
```

```
[ ]: import datetime

      date_actual = datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
      log_dir = "logs/fit/" + date_actual
      tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir,
      ↪ histogram_freq=1)

      es_callback = tf.keras.callbacks.EarlyStopping(monitor='val_loss', mode='min',
      ↪ verbose=1, patience=8)

      checkpoint_callback = tf.keras.callbacks.
      ↪ ModelCheckpoint(filepath=f"model_weights/{date_actual}_cnn_best_weights.
      ↪ hdf5",
                        monitor='val_loss',
                        verbose=1,
                        save_best_only=True)
```

```
[ ]: # batch_size = 0 because we already have batch size in tf dataset
      with tf.device("/gpu:0"):
          model_h = fit_model(train_ds, val_ds, epochs=60, steps=np.ceil(train_size /
      ↪ BATCH_SIZE), val_steps=np.ceil(val_size / BATCH_SIZE),
      ↪ callbacks=[tensorboard_callback, checkpoint_callback, es_callback])
```

```
Epoch 1/60
634/635 [=====>.] - ETA: 0s - loss: 0.2345 - mse: 0.2345
- rmse: 0.4842 - mae: 0.2784 - mape: 140.3885
Epoch 1: val_loss improved from inf to 0.08667, saving model to
model_weights/20221103-230002_cnn_best_weights.hdf5
635/635 [=====] - 80s 122ms/step - loss: 0.2343 - mse:
0.2343 - rmse: 0.4841 - mae: 0.2784 - mape: 140.3092 - val_loss: 0.0867 -
val_mse: 0.0867 - val_rmse: 0.2944 - val_mae: 0.2105 - val_mape: 65.5157
Epoch 2/60
634/635 [=====>.] - ETA: 0s - loss: 0.0419 - mse: 0.0419
```

- rmse: 0.2046 - mae: 0.1448 - mape: 78.3303
Epoch 2: val_loss improved from 0.08667 to 0.07906, saving model to
model_weights/20221103-230002_cnn_best_weights.hdf5
635/635 [=====] - 76s 120ms/step - loss: 0.0419 - mse:
0.0419 - rmse: 0.2046 - mae: 0.1448 - mape: 78.2950 - val_loss: 0.0791 -
val_mse: 0.0791 - val_rmse: 0.2812 - val_mae: 0.1927 - val_mape: 64.4429
Epoch 3/60
634/635 [=====>.] - ETA: 0s - loss: 0.0305 - mse: 0.0305
- rmse: 0.1748 - mae: 0.1226 - mape: 70.2164
Epoch 3: val_loss improved from 0.07906 to 0.07415, saving model to
model_weights/20221103-230002_cnn_best_weights.hdf5
635/635 [=====] - 76s 119ms/step - loss: 0.0305 - mse:
0.0305 - rmse: 0.1747 - mae: 0.1225 - mape: 70.1723 - val_loss: 0.0742 -
val_mse: 0.0742 - val_rmse: 0.2723 - val_mae: 0.1817 - val_mape: 60.1087
Epoch 4/60
634/635 [=====>.] - ETA: 0s - loss: 0.0270 - mse: 0.0270
- rmse: 0.1642 - mae: 0.1144 - mape: 63.5036
Epoch 4: val_loss improved from 0.07415 to 0.06344, saving model to
model_weights/20221103-230002_cnn_best_weights.hdf5
635/635 [=====] - 76s 119ms/step - loss: 0.0269 - mse:
0.0269 - rmse: 0.1641 - mae: 0.1144 - mape: 63.4689 - val_loss: 0.0634 -
val_mse: 0.0634 - val_rmse: 0.2519 - val_mae: 0.1574 - val_mape: 49.7525
Epoch 5/60
634/635 [=====>.] - ETA: 0s - loss: 0.0224 - mse: 0.0224
- rmse: 0.1496 - mae: 0.1036 - mape: 57.7936
Epoch 5: val_loss did not improve from 0.06344
635/635 [=====] - 75s 118ms/step - loss: 0.0224 - mse:
0.0224 - rmse: 0.1496 - mae: 0.1036 - mape: 57.7635 - val_loss: 0.0652 -
val_mse: 0.0652 - val_rmse: 0.2553 - val_mae: 0.1577 - val_mape: 49.2678
Epoch 6/60
634/635 [=====>.] - ETA: 0s - loss: 0.0219 - mse: 0.0219
- rmse: 0.1479 - mae: 0.1024 - mape: 59.2665
Epoch 6: val_loss did not improve from 0.06344
635/635 [=====] - 75s 118ms/step - loss: 0.0219 - mse:
0.0219 - rmse: 0.1479 - mae: 0.1024 - mape: 59.2348 - val_loss: 0.0666 -
val_mse: 0.0666 - val_rmse: 0.2582 - val_mae: 0.1632 - val_mape: 50.4255
Epoch 7/60
634/635 [=====>.] - ETA: 0s - loss: 0.0192 - mse: 0.0192
- rmse: 0.1387 - mae: 0.0957 - mape: 55.0968
Epoch 7: val_loss did not improve from 0.06344
635/635 [=====] - 75s 119ms/step - loss: 0.0192 - mse:
0.0192 - rmse: 0.1387 - mae: 0.0957 - mape: 55.0644 - val_loss: 0.0690 -
val_mse: 0.0690 - val_rmse: 0.2627 - val_mae: 0.1642 - val_mape: 59.2706
Epoch 8/60
634/635 [=====>.] - ETA: 0s - loss: 0.0247 - mse: 0.0247
- rmse: 0.1572 - mae: 0.1047 - mape: 58.8439
Epoch 8: val_loss did not improve from 0.06344
635/635 [=====] - 75s 119ms/step - loss: 0.0247 - mse:

0.0247 - rmse: 0.1571 - mae: 0.1047 - mape: 58.8113 - val_loss: 0.0781 - val_mse: 0.0781 - val_rmse: 0.2794 - val_mae: 0.1724 - val_mape: 55.7035

Epoch 9/60

634/635 [=====>.] - ETA: 0s - loss: 0.0246 - mse: 0.0246 - rmse: 0.1568 - mae: 0.1041 - mape: 59.2239

Epoch 9: val_loss did not improve from 0.06344

635/635 [=====] - 75s 118ms/step - loss: 0.0246 - mse: 0.0246 - rmse: 0.1568 - mae: 0.1041 - mape: 59.1929 - val_loss: 0.0802 - val_mse: 0.0802 - val_rmse: 0.2831 - val_mae: 0.1682 - val_mape: 48.9895

Epoch 10/60

634/635 [=====>.] - ETA: 0s - loss: 0.0170 - mse: 0.0170 - rmse: 0.1302 - mae: 0.0882 - mape: 51.5287

Epoch 10: val_loss improved from 0.06344 to 0.06236, saving model to model_weights/20221103-230002_cnn_best_weights.hdf5

635/635 [=====] - 76s 119ms/step - loss: 0.0170 - mse: 0.0170 - rmse: 0.1302 - mae: 0.0881 - mape: 51.4980 - val_loss: 0.0624 - val_mse: 0.0624 - val_rmse: 0.2497 - val_mae: 0.1509 - val_mape: 58.8778

Epoch 11/60

634/635 [=====>.] - ETA: 0s - loss: 0.0145 - mse: 0.0145 - rmse: 0.1205 - mae: 0.0821 - mape: 50.1194

Epoch 11: val_loss improved from 0.06236 to 0.05849, saving model to model_weights/20221103-230002_cnn_best_weights.hdf5

635/635 [=====] - 76s 119ms/step - loss: 0.0145 - mse: 0.0145 - rmse: 0.1205 - mae: 0.0821 - mape: 50.0889 - val_loss: 0.0585 - val_mse: 0.0585 - val_rmse: 0.2419 - val_mae: 0.1437 - val_mape: 44.7324

Epoch 12/60

634/635 [=====>.] - ETA: 0s - loss: 0.0147 - mse: 0.0147 - rmse: 0.1211 - mae: 0.0813 - mape: 47.9710

Epoch 12: val_loss did not improve from 0.05849

635/635 [=====] - 75s 119ms/step - loss: 0.0147 - mse: 0.0147 - rmse: 0.1211 - mae: 0.0813 - mape: 47.9434 - val_loss: 0.1153 - val_mse: 0.1153 - val_rmse: 0.3395 - val_mae: 0.2035 - val_mape: 49.4566

Epoch 13/60

634/635 [=====>.] - ETA: 0s - loss: 0.0145 - mse: 0.0145 - rmse: 0.1203 - mae: 0.0804 - mape: 48.8194

Epoch 13: val_loss did not improve from 0.05849

635/635 [=====] - 72s 113ms/step - loss: 0.0145 - mse: 0.0145 - rmse: 0.1202 - mae: 0.0804 - mape: 48.7864 - val_loss: 0.0618 - val_mse: 0.0618 - val_rmse: 0.2487 - val_mae: 0.1561 - val_mape: 63.2799

Epoch 14/60

634/635 [=====>.] - ETA: 0s - loss: 0.0129 - mse: 0.0129 - rmse: 0.1137 - mae: 0.0762 - mape: 43.7422

Epoch 14: val_loss did not improve from 0.05849

635/635 [=====] - 72s 113ms/step - loss: 0.0129 - mse: 0.0129 - rmse: 0.1137 - mae: 0.0762 - mape: 43.7214 - val_loss: 0.0601 - val_mse: 0.0601 - val_rmse: 0.2451 - val_mae: 0.1567 - val_mape: 58.1463

Epoch 15/60

634/635 [=====>.] - ETA: 0s - loss: 0.0118 - mse: 0.0118

```

- rmse: 0.1088 - mae: 0.0739 - mape: 43.2454
Epoch 15: val_loss did not improve from 0.05849
635/635 [=====] - 72s 113ms/step - loss: 0.0118 - mse:
0.0118 - rmse: 0.1088 - mae: 0.0739 - mape: 43.2178 - val_loss: 0.0681 -
val_mse: 0.0681 - val_rmse: 0.2610 - val_mae: 0.1638 - val_mape: 59.8237
Epoch 16/60
634/635 [=====>.] - ETA: 0s - loss: 0.0209 - mse: 0.0209
- rmse: 0.1445 - mae: 0.0915 - mape: 52.1147
Epoch 16: val_loss did not improve from 0.05849
635/635 [=====] - 72s 113ms/step - loss: 0.0209 - mse:
0.0209 - rmse: 0.1445 - mae: 0.0915 - mape: 52.0831 - val_loss: 0.0629 -
val_mse: 0.0629 - val_rmse: 0.2509 - val_mae: 0.1527 - val_mape: 51.7619
Epoch 17/60
634/635 [=====>.] - ETA: 0s - loss: 0.0119 - mse: 0.0119
- rmse: 0.1091 - mae: 0.0733 - mape: 43.4485
Epoch 17: val_loss did not improve from 0.05849
635/635 [=====] - 71s 112ms/step - loss: 0.0119 - mse:
0.0119 - rmse: 0.1090 - mae: 0.0733 - mape: 43.4247 - val_loss: 0.0628 -
val_mse: 0.0628 - val_rmse: 0.2505 - val_mae: 0.1556 - val_mape: 52.8115
Epoch 18/60
634/635 [=====>.] - ETA: 0s - loss: 0.0134 - mse: 0.0134
- rmse: 0.1158 - mae: 0.0766 - mape: 43.6014
Epoch 18: val_loss did not improve from 0.05849
635/635 [=====] - 72s 113ms/step - loss: 0.0134 - mse:
0.0134 - rmse: 0.1160 - mae: 0.0767 - mape: 43.5898 - val_loss: 0.0890 -
val_mse: 0.0890 - val_rmse: 0.2984 - val_mae: 0.1766 - val_mape: 57.7655
Epoch 19/60
634/635 [=====>.] - ETA: 0s - loss: 0.0213 - mse: 0.0213
- rmse: 0.1460 - mae: 0.0934 - mape: 55.5282
Epoch 19: val_loss did not improve from 0.05849
635/635 [=====] - 72s 113ms/step - loss: 0.0213 - mse:
0.0213 - rmse: 0.1460 - mae: 0.0934 - mape: 55.4964 - val_loss: 0.0698 -
val_mse: 0.0698 - val_rmse: 0.2643 - val_mae: 0.1614 - val_mape: 55.8318
Epoch 19: early stopping

```

0.5 Evaluate model

```
[ ]: print(date_actual)
```

```
20221103-230002
```

```
[ ]: best_model = models.load_model(f'model_weights/{date_actual}_cnn_best_weights.
↳hdf5')
#best_model = models.load_model(f'best_models_weights/cnn_best_weights_v9.hdf5')
```

```
[ ]: def evaluate_model(model, test_values, steps):
    score = model.evaluate(test_values, steps=steps)
```



```
return score
```

```
[ ]: test_loss, test_mse, test_rmse, test_mae, test_mape =   
      ↪ evaluate_model(best_model, test_ds, steps=np.ceil(test_size / BATCH_SIZE))
```

```
398/398 [=====] - 32s 78ms/step - loss: 0.0540 - mse: 0.0540 - rmse: 0.2324 - mae: 0.1440 - mape: 68.1599
```

```
[ ]: #predictions = best_model.predict(test_ds, steps=np.ceil(test_size /   
      ↪ BATCH_SIZE))
```

```
[ ]: for image, stage_discharge in test_ds.take(1):  
      predictions = best_model.predict(x=image)  
  
      stage_discharge_test_values = stage_discharge.numpy()  
      predictions_values = predictions  
  
      diff = predictions_values.flatten() - stage_discharge_test_values.  
      ↪ flatten()  
      percentDiff = (diff / stage_discharge_test_values.flatten()) * 100  
      absPercentDiff = np.abs(percentDiff)  
      # compute the mean and standard deviation of the absolute percentage  
      # difference  
      mean = np.mean(absPercentDiff)  
      std = np.std(absPercentDiff)  
      # finally, show some statistics on our model  
      print(mean)  
      print(std)  
  
      stage_discharge_test_values = stage_discharge[:10]  
      stage_discharge_test_values = stage_discharge_test_values.numpy().  
      ↪ reshape(10, 2)  
      predictions_values = predictions[:10]  
  
      for i in range(len(stage_discharge_test_values)):  
          print(f"pred stage: {scaler.  
      ↪ inverse_transform(predictions_values)[i][0]}, actual stage: {scaler.  
      ↪ inverse_transform(stage_discharge_test_values)[i][0]}")  
          print(f"pred discharge: {scaler.  
      ↪ inverse_transform(predictions_values)[i][1]}, actual discharge: {scaler.  
      ↪ inverse_transform(stage_discharge_test_values)[i][1]}")
```

```
1/1 [=====] - 0s 495ms/step  
79.2600525051517  
218.17216149864743  
pred stage: 2.4097249507904053, actual stage: 2.31  
pred discharge: 284.9043884277344, actual discharge: 210.0000000000001
```

```

pred stage: 2.463397979736328, actual stage: 2.41
pred discharge: 332.1559753417969, actual discharge: 210.00000000000001
pred stage: 3.3087246417999268, actual stage: 2.67
pred discharge: 1338.3272705078125, actual discharge: 210.00000000000001
pred stage: 2.3510382175445557, actual stage: 2.43
pred discharge: 257.7505187988281, actual discharge: 210.00000000000001
pred stage: 2.4264814853668213, actual stage: 2.42
pred discharge: 329.0776672363281, actual discharge: 210.00000000000001
pred stage: 2.7900302410125732, actual stage: 2.71
pred discharge: 804.0298461914062, actual discharge: 210.00000000000001
pred stage: 2.40978741645813, actual stage: 2.71
pred discharge: 290.5561828613281, actual discharge: 210.00000000000001
pred stage: 2.2910003662109375, actual stage: 2.76
pred discharge: 218.1156768798828, actual discharge: 210.00000000000001
pred stage: 2.5312607288360596, actual stage: 2.89
pred discharge: 354.0725402832031, actual discharge: 210.00000000000001
pred stage: 2.907449245452881, actual stage: 2.82
pred discharge: 1025.2611083984375, actual discharge: 210.00000000000001

```

0.5.1 Residual analysis

```

[ ]: y_predictions = np.empty(shape=(1, 2)).astype('float32')
     y_real = np.empty(shape=(1, 2)).astype('float32')

     for image, stage_discharge in test_ds.take(100):
         y_predictions = np.concatenate((y_predictions, best_model.predict(x=image)))
         y_real = np.concatenate((y_real, stage_discharge.numpy()))

```

```

1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 22ms/step

```

1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 20ms/step
1/1 [=====] - 0s 21ms/step

[illegible]

```
[ ]: residuals = y_real - y_predictions
residuals_std = residuals/residuals.std()

y_real_stage = np.array([i[0] for i in y_real])
residual_stage = np.array([i[0] for i in residuals])

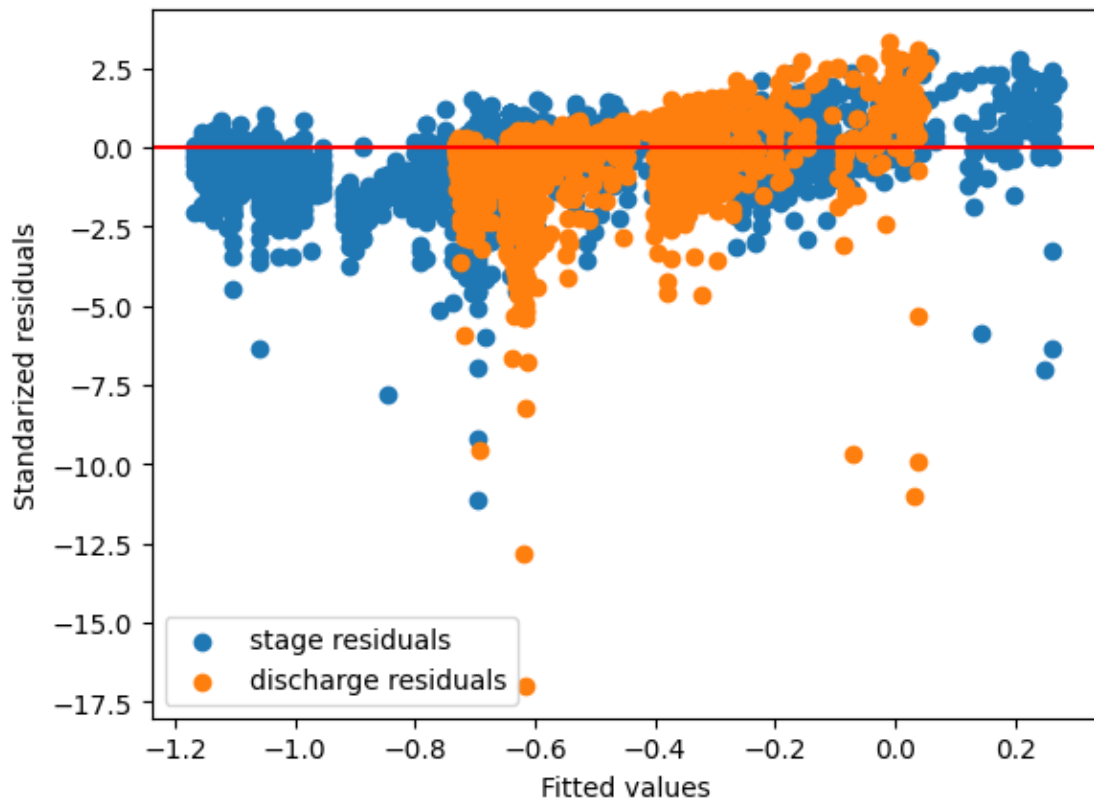
y_real_discharge = np.array([i[1] for i in y_real])
residual_discharge = np.array([i[1] for i in residuals])

print(residual_stage.std())


plt.scatter(y_real_stage, residual_stage / residual_stage.std(), label="stage_1
↪residuals")
```

```
plt.scatter(y_real_discharge, residual_discharge / residual_discharge.std(),  
            ↪label="discharge residuals")  
plt.axhline(y=0.0, color='r', linestyle='-')  
plt.xlabel("Fitted values")  
plt.ylabel("Standardized residuals")  
  
plt.legend()  
plt.show()
```


0.21593582108796294




```
[ ]: import statsmodels.api as sm  
from statsmodels.stats.diagnostic import normal_ad  
  
figure = sm.qqplot(residual_stage / residual_stage.std(), line='45',  
                    ↪label='stage')  
plt.show()
```




```
[ ]: figure = sm.qqplot(residual_discharge / residual_discharge.std(), line='45',
    ↪label='discharge')
plt.show()
```



```
[ ]: plt.hist(residual_stage, density=True, bins = 50)
plt.ylabel('Count')
plt.xlabel('Residual stage');
plt.show()
```



```
[ ]: plt.hist(residual_discharge, density=True, bins = 50)
plt.ylabel('Count')
plt.xlabel('Residual discharge');
plt.show()
```

```
[ ]: stat, pval = normal_ad(residual_stage)
print("p-value:", pval)

if pval<0.05:
    print("Hay evidencia de que los residuos no provienen de una distribución
    ↪normal.")
else:
    print("No hay evidencia para rechazar la hipótesis de que los residuos
    ↪vienen de una distribución normal.")
```

p-value: 0.0

Hay evidencia de que los residuos no provienen de una distribución normal.

```
[ ]: stat, pval = normal_ad(residual_discharge)
print("p-value:", pval)

if pval < 0.05:
    print("Hay evidencia de que los residuos no provienen de una distribución
    ↪normal.")
else:
```

```
print("No hay evidencia para rechazar la hipótesis de que los residuos_
↪vienen de una distribución normal.")
```

p-value: 0.0

Hay evidencia de que los residuos no provienen de una distribución normal.

0.6 Visualize layers

```
[ ]: layer_outputs = [layer.output for layer in best_model.layers[:12]]
# Extracts the outputs of the top 12 layers
activation_model = models.Model(inputs=best_model.input, outputs=layer_outputs)_
↪# Creates a model that will return these outputs, given the model input
```

```
-----
ValueError                                Traceback (most recent call last)
Cell In [95], line 3
      1 layer_outputs = [layer.output for layer in best_model.layers[:12]]
      2 # Extracts the outputs of the top 12 layers
----> 3 activation_model = models.Model(inputs=best_model.input,
↪outputs=layer_outputs)

File ~/miniconda3/envs/tf-gpu/lib/python3.10/site-packages/tensorflow/python/
↪trackable/base.py:205, in no_automatic_dependency_tracking.<locals>._
↪method_wrapper(self, *args, **kwargs)
    203 self._self_setattr_tracking = False # pylint: disable=protected-access
    204 try:
--> 205     result = method(self, *args, **kwargs)
    206 finally:
    207     self._self_setattr_tracking = previous_value # pylint:
↪disable=protected-access

File ~/miniconda3/envs/tf-gpu/lib/python3.10/site-packages/keras/engine/
↪functional.py:165, in Functional.__init__(self, inputs, outputs, name,
↪trainable, **kwargs)
    156     if not all(
    157         [
    158             functional_utils.is_input_keras_tensor(t)
    159             for t in tf.nest.flatten(inputs)
    160         ]
    161     ):
    162         inputs, outputs = functional_utils.clone_graph_nodes(
    163             inputs, outputs
    164         )
--> 165 self._init_graph_network(inputs, outputs)

File ~/miniconda3/envs/tf-gpu/lib/python3.10/site-packages/tensorflow/python/
↪trackable/base.py:205, in no_automatic_dependency_tracking.<locals>._
↪method_wrapper(self, *args, **kwargs)
```

```

    203 self._self_setattr_tracking = False # pylint: disable=protected-access
    204 try:
--> 205     result = method(self, *args, **kwargs)
    206 finally:
    207     self._self_setattr_tracking = previous_value # pylint:
↳ disable=protected-access

```

```

File ~/miniconda3/envs/tf-gpu/lib/python3.10/site-packages/keras/engine/
↳ functional.py:264, in Functional._init_graph_network(self, inputs, outputs)
    261     self._input_coordinates.append((layer, node_index, tensor_index))
    263 # Keep track of the network's nodes and layers.
--> 264 nodes, nodes_by_depth, layers, _ = _map_graph_network(
    265     self.inputs, self.outputs
    266 )
    267 self._network_nodes = nodes
    268 self._nodes_by_depth = nodes_by_depth

```

```

File ~/miniconda3/envs/tf-gpu/lib/python3.10/site-packages/keras/engine/
↳ functional.py:1128, in _map_graph_network(inputs, outputs)
    1126 for x in tf.nest.flatten(node.keras_inputs):
    1127     if id(x) not in computable_tensors:
-> 1128         raise ValueError(
    1129             f"Graph disconnected: cannot obtain value for "
    1130             f'tensor {x} at layer "{layer.name}". '
    1131             "The following previous layers were accessed "
    1132             f"without issue: {layers_with_complete_input}"
    1133         )
    1134 for x in tf.nest.flatten(node.outputs):
    1135     computable_tensors.add(id(x))

```

```

ValueError: Graph disconnected: cannot obtain value for tensor_
↳ KerasTensor(type_spec=TensorSpec(shape=(None, 224, 224, 3), dtype=tf.float32,
↳ name='input_4'), name='input_4', description="created by layer 'input_4'") at
↳ layer "conv1_pad". The following previous layers were accessed without issue:
↳ []

```

```
[ ]: activations = activation_model.predict(test_ds.take(1))
```

Running cells with 'Python 3.9.13 ('tf-metal')' requires ipykernel package.

Run the following command to install 'ipykernel' into the Python environment.

Command: 'conda install -n tf-metal ipykernel --update-deps --force-reinstall'

```
[ ]: import matplotlib.pyplot as plt

layer_names = []
for layer in best_model.layers[:12]:
    layer_names.append(layer.name) # Names of the layers, so you can have them
    ↪ as part of your plot

images_per_row = 16

for layer_name, layer_activation in zip(layer_names, activations): # Displays
    ↪ the feature maps
    n_features = layer_activation.shape[-1] # Number of features in the feature
    ↪ map
    size = layer_activation.shape[1] # The feature map has shape (1, size, size,
    ↪ n_features).
    n_cols = n_features // images_per_row # Tiles the activation channels in
    ↪ this matrix
    display_grid = np.zeros((size * n_cols, images_per_row * size))

    print(layer_name)
    if "flatten" in layer_name or "dense" in layer_name: break

    for col in range(n_cols): # Tiles each filter into a big horizontal grid
        for row in range(images_per_row):
            channel_image = layer_activation[0,
                                           :, :,
                                           col * images_per_row + row]
            channel_image -= channel_image.mean() # Post-processes the feature
            ↪ to make it visually palatable
            channel_image /= channel_image.std()
            channel_image *= 64
            channel_image += 128
            channel_image = np.clip(channel_image, 0, 255).astype('uint8')
            display_grid[col * size : (col + 1) * size, # Displays the grid
                          row * size : (row + 1) * size] = channel_image

    scale = 1. / size
    plt.figure(figsize=(scale * display_grid.shape[1],
                        scale * display_grid.shape[0]))
    plt.title(layer_name)
    plt.grid(False)
    plt.imshow(display_grid, aspect='auto', cmap='viridis')
```

Running cells with 'Python 3.9.13 ('tf-metal')' requires ipykernel package.

Run the following command to install 'ipykernel' into the Python environment.

```
Command: 'conda install -n tf-metal ipykernel --update-deps --force-reinstall'
```