

# cnn\_v1\_stage

November 25, 2022

```
[ ]: %env LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$CONDA_PREFIX/lib/  
      #%env TF_GPU_ALLOCATOR=cuda_malloc_async
```

```
env: LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$CONDA_PREFIX/lib/
```

```
[ ]: import os  
      print(os.environ["LD_LIBRARY_PATH"])
```

```
$LD_LIBRARY_PATH:$CONDA_PREFIX/lib/
```

```
[ ]: import tensorflow as tf  
      import numpy as np  
      import pandas as pd  
      import os  
      import matplotlib.pyplot as plt  
  
      from tensorflow.keras import Sequential, models, Input  
      from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D,  
      ↪Dropout, LeakyReLU, AveragePooling2D, GlobalAveragePooling2D,  
      ↪BatchNormalization, TimeDistributed, LSTM, SpatialDropout2D, concatenate  
      from tensorflow.keras.optimizers import SGD, Adam
```

```
2022-11-25 11:53:21.459527: I tensorflow/core/platform/cpu_feature_guard.cc:193]  
This TensorFlow binary is optimized with oneAPI Deep Neural Network Library  
(oneDNN) to use the following CPU instructions in performance-critical  
operations:  AVX2 FMA
```

```
To enable them in other operations, rebuild TensorFlow with the appropriate  
compiler flags.
```

```
2022-11-25 11:53:22.497099: E tensorflow/stream_executor/cuda/cuda_blas.cc:2981]  
Unable to register cuBLAS factory: Attempting to register factory for plugin  
cuBLAS when one has already been registered
```

```
2022-11-25 11:53:23.505095: W
```

```
tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load  
dynamic library 'libnvinfer.so.7'; dLError: libnvinfer.so.7: cannot open shared  
object file: No such file or directory; LD_LIBRARY_PATH:  
:/home/nkspartan/miniconda3/envs/tf-gpu/lib/
```

```
2022-11-25 11:53:23.505179: W
```

```
tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load
```

dynamic library 'libnvinfer\_plugin.so.7'; dLError: libnvinfer\_plugin.so.7:  
cannot open shared object file: No such file or directory; LD\_LIBRARY\_PATH:  
:/home/nkspartan/miniconda3/envs/tf-gpu/lib/  
2022-11-25 11:53:23.505185: W  
tensorflow/compiler/tf2tensorrt/utils/py\_utils.cc:38] TF-TRT Warning: Cannot  
dlopen some TensorRT libraries. If you would like to use Nvidia GPU with  
TensorRT, please make sure the missing libraries mentioned above are installed  
properly.

```
[ ]: gpus = tf.config.experimental.list_physical_devices('GPU')  
      for gpu in gpus:  
          tf.config.experimental.set_memory_growth(gpu, True)
```

2022-11-25 11:53:24.976824: I  
tensorflow/stream\_executor/cuda/cuda\_gpu\_executor.cc:980] successful NUMA node  
read from SysFS had negative value (-1), but there must be at least one NUMA  
node, so returning NUMA node zero  
2022-11-25 11:53:25.027810: I  
tensorflow/stream\_executor/cuda/cuda\_gpu\_executor.cc:980] successful NUMA node  
read from SysFS had negative value (-1), but there must be at least one NUMA  
node, so returning NUMA node zero  
2022-11-25 11:53:25.028090: I  
tensorflow/stream\_executor/cuda/cuda\_gpu\_executor.cc:980] successful NUMA node  
read from SysFS had negative value (-1), but there must be at least one NUMA  
node, so returning NUMA node zero

```
[ ]: from tensorflow.python.client import device_lib  
  
      print('Default GPU Device: {}'.format(tf.test.gpu_device_name()))
```

2022-11-25 11:53:25.085015: I tensorflow/core/platform/cpu\_feature\_guard.cc:193]  
This TensorFlow binary is optimized with oneAPI Deep Neural Network Library  
(oneDNN) to use the following CPU instructions in performance-critical  
operations: AVX2 FMA  
To enable them in other operations, rebuild TensorFlow with the appropriate  
compiler flags.  
2022-11-25 11:53:25.086183: I  
tensorflow/stream\_executor/cuda/cuda\_gpu\_executor.cc:980] successful NUMA node  
read from SysFS had negative value (-1), but there must be at least one NUMA  
node, so returning NUMA node zero  
2022-11-25 11:53:25.086397: I  
tensorflow/stream\_executor/cuda/cuda\_gpu\_executor.cc:980] successful NUMA node  
read from SysFS had negative value (-1), but there must be at least one NUMA  
node, so returning NUMA node zero  
2022-11-25 11:53:25.086539: I  
tensorflow/stream\_executor/cuda/cuda\_gpu\_executor.cc:980] successful NUMA node  
read from SysFS had negative value (-1), but there must be at least one NUMA  
node, so returning NUMA node zero

```

2022-11-25 11:53:25.929895: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-11-25 11:53:25.930508: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from S

Default GPU Device: /device:GPU:0

ysFS had negative value (-1), but there must be at least one NUMA node, so
returning NUMA node zero
2022-11-25 11:53:25.930819: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-11-25 11:53:25.931063: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:1616] Created device
/device:GPU:0 with 4063 MB memory: -> device: 0, name: NVIDIA GeForce RTX 2060,
pci bus id: 0000:08:00.0, compute capability: 7.5

```

## 0.1 Read the csv dataset to get the values for stage and discharge of the images

```

[ ]: df = pd.read_csv("../dataset/2012_2019_PlatteRiverWeir_features_merged_all.
    ↪ csv")
df.head()

```

```

[ ]:
   Unnamed: 0  SensorTime  CaptureTime \
0           0  2012-06-09 13:15:00  2012-06-09T13:09:07
1           1  2012-06-09 13:15:00  2012-06-09T13:10:29
2           2  2012-06-09 13:45:00  2012-06-09T13:44:01
3           3  2012-06-09 14:45:00  2012-06-09T14:44:30
4           4  2012-06-09 15:45:00  2012-06-09T15:44:59

   Filename Agency  SiteNumber  TimeZone  Stage \
0  StateLineWeir_20120609_Farrell_001.jpg  USGS    6674500      MDT    2.99
1  StateLineWeir_20120609_Farrell_002.jpg  USGS    6674500      MDT    2.99
2  StateLineWeir_20120609_Farrell_003.jpg  USGS    6674500      MDT    2.96
3  StateLineWeir_20120609_Farrell_004.jpg  USGS    6674500      MDT    2.94
4  StateLineWeir_20120609_Farrell_005.jpg  USGS    6674500      MDT    2.94

   Discharge  CalcTimestamp  ...  WeirPt2X  WeirPt2Y  WwRawLineMin  \
0      916.0  2020-03-11T16:58:28  ...      -1      -1           0.0
1      916.0  2020-03-11T16:58:33  ...      -1      -1           0.0
2      873.0  2020-03-11T16:58:40  ...      -1      -1           0.0
3      846.0  2020-03-11T16:58:47  ...      -1      -1           0.0
4      846.0  2020-03-11T16:58:55  ...      -1      -1           0.0

```

	WwRawLineMax	WwRawLineMean	WwRawLineSigma	WwCurveLineMin	\
0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	

	WwCurveLineMax	WwCurveLineMean	WwCurveLineSigma
0	0.0	0.0	0.0
1	0.0	0.0	0.0
2	0.0	0.0	0.0
3	0.0	0.0	0.0
4	0.0	0.0	0.0

[5 rows x 60 columns]

```
[ ]: #df = df[["Filename", "Stage", "Discharge", 'SensorTime', "RiverArea",
↪ "RiverWidth"]]
df = df[["Filename", "Stage", "Discharge", 'SensorTime']]
```

```
[ ]: df['SensorTime'] = pd.to_datetime(df['SensorTime'])
df['Year'] = df['SensorTime'].dt.year
df.head()
```

	Filename	Stage	Discharge	\
0	StateLineWeir_20120609_Farrell_001.jpg	2.99	916.0	
1	StateLineWeir_20120609_Farrell_002.jpg	2.99	916.0	
2	StateLineWeir_20120609_Farrell_003.jpg	2.96	873.0	
3	StateLineWeir_20120609_Farrell_004.jpg	2.94	846.0	
4	StateLineWeir_20120609_Farrell_005.jpg	2.94	846.0	

	SensorTime	Year
0	2012-06-09 13:15:00	2012
1	2012-06-09 13:15:00	2012
2	2012-06-09 13:45:00	2012
3	2012-06-09 14:45:00	2012
4	2012-06-09 15:45:00	2012

```
[ ]: df = df.sort_values(by="SensorTime", ascending=True)
df.head()
```

	Filename	Stage	Discharge	\
0	StateLineWeir_20120609_Farrell_001.jpg	2.99	916.0	
1	StateLineWeir_20120609_Farrell_002.jpg	2.99	916.0	
2	StateLineWeir_20120609_Farrell_003.jpg	2.96	873.0	
3	StateLineWeir_20120609_Farrell_004.jpg	2.94	846.0	
4	StateLineWeir_20120609_Farrell_005.jpg	2.94	846.0	

	SensorTime	Year
0	2012-06-09 13:15:00	2012
1	2012-06-09 13:15:00	2012
2	2012-06-09 13:45:00	2012
3	2012-06-09 14:45:00	2012
4	2012-06-09 15:45:00	2012

### 0.1.1 Remove outliers

```
[ ]: df = df[df.Stage > 0]
df = df[df.Discharge > 0]
```

We consider values equal to 0 as outliers because from the photos it doesn't seem that it would be possible that at this time we would have a value of 0 for stage or discharge

```
[ ]: df.shape
```

```
[ ]: (40148, 5)
```

### 0.1.2 Scale the data

```
[ ]: from sklearn.preprocessing import StandardScaler, MinMaxScaler, RobustScaler
from joblib import load

scaler = RobustScaler()
#scaler = load('std_scaler.joblib') # scaler with all the 42059 observations
```

Scale the data based only on the training dataset (in this case the training dataset is from 2012 to 2016)

```
[ ]: #data_to_scale_fit = df[(df["Year"] >= 2012) & (df["Year"] <= 2016)][["Stage",
↪ "Discharge"]]
data_to_scale_fit = df[(df["Year"] >= 2012) & (df["Year"] <=
↪ 2016)][["Discharge"]]
data_to_scale_fit
```

```
[ ]: Discharge
0      916.0
1      916.0
2      873.0
3      846.0
4      846.0
...      ...
21416    279.0
21417    279.0
21418    279.0
```

```
21419      279.0
21420      279.0
```

```
[20304 rows x 1 columns]
```

```
[ ]: scaler.fit(data_to_scale_fit)
```

```
[ ]: RobustScaler()
```

```
[ ]: #df[["Stage", "Discharge"]] = scaler.transform(df[["Stage", "Discharge"]])
#df[["Discharge"]] = scaler.transform(df[["Discharge"]])
df
```

```
[ ]:
```

	Filename	Stage	Discharge	\
0	StateLineWeir_20120609_Farrell_001.jpg	2.99	916.0	
1	StateLineWeir_20120609_Farrell_002.jpg	2.99	916.0	
2	StateLineWeir_20120609_Farrell_003.jpg	2.96	873.0	
3	StateLineWeir_20120609_Farrell_004.jpg	2.94	846.0	
4	StateLineWeir_20120609_Farrell_005.jpg	2.94	846.0	
...	...	...	...	
42054	StateLineWeir_20191011_Farrell_409.jpg	2.54	434.0	
42055	StateLineWeir_20191011_Farrell_410.jpg	2.54	434.0	
42056	StateLineWeir_20191011_Farrell_411.jpg	2.54	434.0	
42057	StateLineWeir_20191011_Farrell_412.jpg	2.54	434.0	
42058	StateLineWeir_20191011_Farrell_413.jpg	2.54	434.0	

	SensorTime	Year
0	2012-06-09 13:15:00	2012
1	2012-06-09 13:15:00	2012
2	2012-06-09 13:45:00	2012
3	2012-06-09 14:45:00	2012
4	2012-06-09 15:45:00	2012
...	...	...
42054	2019-10-11 09:00:00	2019
42055	2019-10-11 10:00:00	2019
42056	2019-10-11 11:00:00	2019
42057	2019-10-11 12:00:00	2019
42058	2019-10-11 12:45:00	2019

```
[40148 rows x 5 columns]
```

```
[ ]: df.describe()
```

```
[ ]:
```

	Stage	Discharge	Year
count	40148.000000	40148.000000	40148.000000
mean	2.903601	1017.063288	2016.168228
std	0.814612	1200.944046	1.997968

min	1.370000	6.730000	2012.000000
25%	2.280000	226.000000	2015.000000
50%	2.600000	451.500000	2016.000000
75%	3.320000	1390.000000	2018.000000
max	6.490000	7920.000000	2019.000000

```
[ ]: from joblib import dump
      #dump(scaler, 'std_scaler_train_value_0_outliers.joblib')
```

## 0.2 Create the dataset pipeline

```
[ ]: #IMG_SIZE = 224
      IMG_SIZE = 320
      BATCH_SIZE = 8
      FRAMES = 10
```

```
[ ]: from dataset_transformer import make_dataset_with_time, make_dataset,
      ↪make_dataset_and_time

      from dataset_transformer_2 import Dataset, DataLoader
```

```
[ ]: path = "../../../dataset/images_tmp_draw_sky"

      #train_ds, train_size, val_ds, val_size, test_ds, test_size =
      ↪make_dataset_and_time(path, BATCH_SIZE, IMG_SIZE, FRAMES, df, 10, True,
      ↪"cnn")

      train_dataset = Dataset(
          path,
          df[(df.Year >= 2012) & (df.Year <= 2016)],
          classes=[0, 1],
      )

      # Dataset for validation images
      val_dataset = Dataset(
          path,
          df[(df.Year >= 2017) & (df.Year <= 2017)],
          classes=[0, 1],
      )

      test_dataset = Dataset(
          path,
          df[(df.Year >= 2018) & (df.Year <= 2019)],
          classes=[0, 1],
          shuffle=False
      )
```

```

train_ds = Dataloader(train_dataset, batch_size=BATCH_SIZE, shuffle=True)
val_ds = Dataloader(val_dataset, batch_size=BATCH_SIZE, shuffle=False)
test_ds = Dataloader(test_dataset, batch_size=1, shuffle=False)

```

```

[ ]: input_shape = 0
output_shape = 0

for i in range(0, 1):
    image_time, stage_discharge = test_ds[i]
    #print(np.array(image_time).shape)

    image = image_time["input_1"]
    stage_discharge = stage_discharge
    #print(stage_discharge.shape)

    print(stage_discharge.shape)

    input_shape = image.shape[1:]
    #input_shape = image_time.numpy().shape[1:]
    output_shape = stage_discharge.shape[1:]

```

(8, 1)

```

[ ]: print(input_shape)
print(output_shape)

```

(320, 320, 3)  
(1,)

### 0.3 Check images

```

[ ]: fig, ax = plt.subplots(nrows=3, ncols=5, figsize=(30, 15))

for i in range(0, 1):
    image_time, stage_discharge = test_ds[i]
    images = image_time["input_1"]

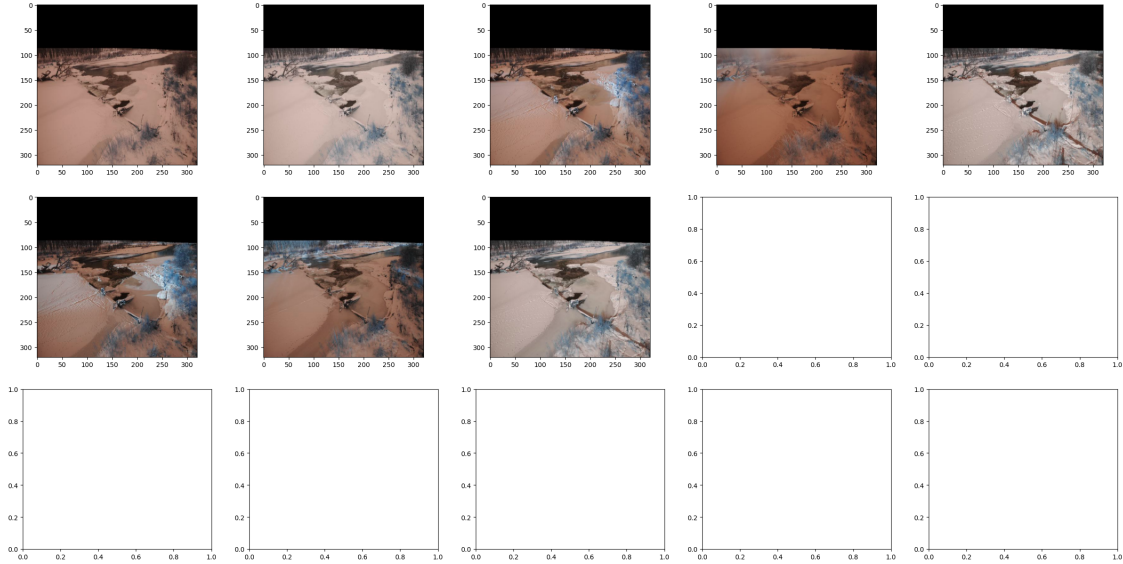
    for img, ax in zip(images, ax.flatten()):
        #print(img.numpy()[ :, :, 3])
        #img = img.numpy()[ :, :, 3]
        #img = img / 2 + 0.5      # unnormalize

        #print(img)
        ax.imshow(img)

plt.show()

```





## 0.4 Create model

```
[ ]: from scipy import ndimage as ndi
from skimage.filters import gabor_kernel

def sobel_kernel(shape, dtype=None):
    #print(shape)
    sobel_x = tf.constant(
        [
            [-5, -4, 0, 4, 5],
            [-8, -10, 0, 10, 8],
            [-10, -20, 0, 20, 10],
            [-8, -10, 0, 10, 8],
            [-5, -4, 0, 4, 5]
        ], dtype=dtype )
    #create the missing dims.
    sobel_x = tf.reshape(sobel_x, (5, 5, 1, 1))

    #print(tf.shape(sobel_x))
    #tile the last 2 axis to get the expected dims.
    sobel_x = tf.tile(sobel_x, (1, 1, shape[-2], shape[-1]))

    #print(tf.shape(sobel_x))
    return sobel_x

def gfb_filter(shape, size=3, tlist=[1,2,3], slist=[2,5], flist=[0.01,0.25,0.
↪5], dtype=None):
    print(shape)
```

```

fsize=np.ones([size,size])
kernels = []
for theta in tlist:
    theta = theta / 4. * np.pi
    for sigma in slist:
        for frequency in flist:
            kernel = np.real(gabor_kernel(frequency,
→theta=theta,sigma_x=sigma, sigma_y=sigma))
            kernels.append(kernel)
gfblist = []
for k, kernel in enumerate(kernels):
    ck=ndi.convolve(fsize, kernel, mode='wrap')
    gfblist.append(ck)

gfblist = np.asarray(gfblist).reshape(size,size,1,len(gfblist))
gfblist = np.repeat(gfblist[:, :, :, :], gfblist.shape[1], axis=2)
print(gfblist.shape)
return tf.keras.backend.variable(gfblist, dtype='float32')

```

```

[ ]: import segmentation_models as sm

seg_model = sm.Unet("resnet50", classes=1, activation="sigmoid")

seg_model.load_weights(
    f'model_weights/seg_model_resnet_50_1.hdf5')

```

```

[ ]: from classification_models.keras import Classifiers

def create_model(input_shape, output_shape, option="normal"):
    model = Sequential()

    if option == "transfer":
        # Inputs
        input_base = Input(shape=input_shape, name="input_1")
        #time_area_input = Input(shape=(2), name="input_2")
        time_input = Input(shape=(1), name="input_2")

        base_model = tf.keras.applications.ResNet50V2(include_top=False,
                                                    weights='imagenet',
                                                    input_shape=input_shape)

        #base_model = ResNet34(include_top=False, weights='imagenet',
→input_shape=input_shape)

        for layer in base_model.layers:
            layer.trainable = False

```

```

base_model._name = 'base_model_ResNet50V2'

cnn_model = base_model(input_base)

cnn_model = Dropout(0.5)(cnn_model)
cnn_model = GlobalAveragePooling2D()(cnn_model)
cnn_model = Dense(1024, activation="relu")(cnn_model)
cnn_model = Dense(512, activation="relu")(cnn_model)

#edge_detection = Conv2D(4, kernel_size=(5, 5),
↳kernel_initializer=sobel_kernel, strides=(2, 2), activation='relu',
↳trainable=False)

#edge_detection = edge_detection(input_base)
#edge_detection.trainable = False
#edge_detection = GlobalAveragePooling2D()(edge_detection)
#edge_detection = Dense(512, activation="elu")(edge_detection)
#edge_detection = Dense(512, activation="elu")(edge_detection)

"""
gfb = Conv2D(filters=18, kernel_size=3, kernel_initializer=gfb_filter,
↳strides=1, padding='valid', trainable=False, name="Gabor_filter")
gfb = gfb(input_base)
gfb.trainable = False
gfb = GlobalAveragePooling2D()(gfb)
gfb = Dense(1024, activation="relu")(gfb)
gfb = Dense(512, activation="relu")(gfb)
"""

combined = concatenate([cnn_model, time_input], name="combined_model")
#combined = BatchNormalization()(combined)

cnn_time = Dense(513, activation="elu")(combined)
#cnn_time = Dropout(0.3)(cnn_time)
cnn_time = Dense(256, activation="elu")(cnn_time)
#cnn_time = Dropout(0.3)(cnn_time)
cnn_time = Dense(128, activation="elu")(cnn_time)
cnn_time = Dense(64, activation="elu")(cnn_time)
output = Dense(output_shape, activation='linear')(cnn_time)

model = tf.keras.Model([input_base, time_input], output,
↳name="cnn_segmentation")
elif option == "normal":
    model.add(Input(shape=input_shape))

```

```

        """model.add(Conv2D(16, kernel_size=(3, 3), activation="elu",
↳padding='same', kernel_initializer='he_uniform'))
        model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
        model.add(BatchNormalization())

        model.add(Conv2D(32, kernel_size=(3, 3), activation="elu",
↳padding='same', kernel_initializer='he_uniform'))
        model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
        model.add(BatchNormalization())

        model.add(Conv2D(32, kernel_size=(3, 3), activation="elu",
↳padding='same', kernel_initializer='he_uniform'))
        model.add(Conv2D(32, kernel_size=(3, 3), activation="elu",
↳padding='same', kernel_initializer='he_uniform'))
        model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
        model.add(BatchNormalization())

        model.add(Conv2D(64, kernel_size=(4, 4), activation="elu",
↳padding='same', kernel_initializer='he_uniform'))
        model.add(Conv2D(64, kernel_size=(4, 4), activation="elu",
↳padding='same', kernel_initializer='he_uniform'))
        model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
        model.add(BatchNormalization())

        model.add(Conv2D(64, kernel_size=(4, 4), activation="elu",
↳padding='same', kernel_initializer='he_uniform'))
        model.add(Conv2D(64, kernel_size=(4, 4), activation="elu",
↳padding='same', kernel_initializer='he_uniform'))
        model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
        model.add(BatchNormalization())

        model.add(Conv2D(64, kernel_size=(3, 3), activation="elu",
↳padding='same', kernel_initializer='he_uniform'))
        model.add(Conv2D(64, kernel_size=(3, 3), activation="elu",
↳padding='same', kernel_initializer='he_uniform'))
        model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
        model.add(BatchNormalization())

        model.add(GlobalAveragePooling2D())

        model.add(Dense(512, activation='elu'))
        model.add(Dropout(0.3))
        model.add(Dense(512, activation='elu'))
        model.add(Dropout(0.3))
        model.add(Dense(256, activation='elu'))

```

```

        model.add(Dense(64, activation='elu'))"""

        model.add(Conv2D(32, kernel_size=(4, 4), strides=(2, 2),
        ↳padding='same', activation="elu"))
        model.add(MaxPooling2D(pool_size=(2, 2)))

        model.add(Conv2D(32, kernel_size=(3, 3), strides=(2, 2),
        ↳activation="elu", padding='same'))
        model.add(MaxPooling2D(pool_size=(2, 2)))

        model.add(Conv2D(32, kernel_size=(3, 3), activation="elu",
        ↳padding='same'))
        #model.add(MaxPooling2D(pool_size=(2, 2)))

        model.add(Conv2D(32, kernel_size=(3, 3), activation='elu'))

        model.add(Conv2D(32, kernel_size=(3, 3), activation='elu'))
        model.add(MaxPooling2D(pool_size=(2, 2)))

        model.add(Conv2D(64, kernel_size=(2, 2), activation='elu'))

        model.add(Conv2D(64, kernel_size=(2, 2), activation='elu'))
        model.add(MaxPooling2D(pool_size=(2, 2)))

        model.add(Flatten())
        model.add(Dense(2048, activation='relu'))
        model.add(Dense(2048, activation='relu'))
        model.add(Dense(1024, activation='relu'))
        model.add(Dense(1024, activation='relu'))

        #model.add(Dense(output_shape, activation='linear')) # linear regression
        ↳output layer

    return model

```

```
[ ]: model = create_model(input_shape, output_shape[0], "transfer")
```

```
[ ]: model.summary()
```

Model: "cnn\_segmentation"

```

-----
Layer (type)                Output Shape              Param #   Connected to
=====
input_1 (InputLayer)        [(None, 320, 320, 3  0   []
                             )]

```

```

base_model_ResNet50V2 (Functional) (None, 10, 10, 2048) 23564800
['input_1[0][0]']
)

dropout (Dropout) (None, 10, 10, 2048) 0
['base_model_ResNet50V2[0][0]']
)

global_average_pooling2d (GlobalAveragePooling2D) (None, 2048) 0
['dropout[0][0]']

dense_5 (Dense) (None, 1024) 2098176
['global_average_pooling2d[0][0]']

dense_6 (Dense) (None, 512) 524800
['dense_5[0][0]']

input_2 (InputLayer) [(None, 1)] 0 []

combined_model (Concatenate) (None, 513) 0
['dense_6[0][0]',
'input_2[0][0]']

dense_7 (Dense) (None, 513) 263682
['combined_model[0][0]']

dense_8 (Dense) (None, 256) 131584
['dense_7[0][0]']

dense_9 (Dense) (None, 128) 32896
['dense_8[0][0]']

dense_10 (Dense) (None, 64) 8256
['dense_9[0][0]']

dense_11 (Dense) (None, 1) 65
['dense_10[0][0]']

```

```

=====
Total params: 26,624,259
Trainable params: 3,059,459
Non-trainable params: 23,564,800
-----

```

```
[ ]: def compile_model(loss_func, optimizer, metrics=["accuracy"]):
    model.compile(loss=loss_func, optimizer=optimizer, metrics=metrics)

[ ]: import tensorflow_addons as tfa

[ ]: sgd = SGD(learning_rate=0.01, decay=1e-3, momentum=0.9, nesterov=True)
adam = Adam(learning_rate=1e-3, decay=1e-3 / 200)

"""compile_model(tfa.losses.PinballLoss(tau=.6), adam, [
    'mse', tf.keras.metrics.RootMeanSquaredError(name='rmse'), 'mae',
    'mape'])"""

"""compile_model(tf.keras.losses.Huber(), adam, [
    'mse', tf.keras.metrics.RootMeanSquaredError(name='rmse'), 'mae',
    'mape'])"""

compile_model("mse", adam, [
    'mse', tf.keras.metrics.RootMeanSquaredError(name='rmse'), 'mae',
    'mape'])

[ ]: def fit_model(training_values, validation_values=None, epochs=10, steps=32,
    val_steps=32, callbacks=[]):
    return model.fit(training_values, validation_data=validation_values,
    epochs=epochs, steps_per_epoch=steps, validation_steps=val_steps,
    callbacks=callbacks)

[ ]: import datetime

date_actual = datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
log_dir = "logs/fit/" + date_actual
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir,
    histogram_freq=1)

es_callback = tf.keras.callbacks.EarlyStopping(monitor='val_loss', mode='min',
    verbose=1, patience=15)

checkpoint_callback = tf.keras.callbacks.
    ModelCheckpoint(filepath=f"model_weights/{date_actual}_cnn_best_weights.
    hdf5",
                    monitor='val_loss',
                    verbose=1,
                    save_best_only=True)

[ ]: # batch_size = 0 because we already have batch size in tf dataset
```

```

model_h = fit_model(train_ds, val_ds, epochs=100, steps=len(train_ds),
    ↳ val_steps=len(val_ds), callbacks=[tensorboard_callback, checkpoint_callback,
    ↳ es_callback, tf.keras.callbacks.ReduceLROnPlateau(patience=5)])

#model.fit(train_ds, validation_data=val_ds, epochs=60,
    ↳ steps_per_epoch=len(train_ds), validation_steps=len(val_ds),
    ↳ callbacks=[tensorboard_callback, checkpoint_callback, es_callback, tf.keras.
    ↳ callbacks.ReduceLROnPlateau(patience=5)])

```

Epoch 1/100

```

2022-11-24 21:21:53.905353: I tensorflow/stream_executor/cuda/cuda_dnn.cc:384]
Loaded cuDNN version 8100
2022-11-24 21:21:55.199009: I
tensorflow/core/platform/default/subprocess.cc:304] Start cannot spawn child
process: No such file or directory
2022-11-24 21:21:55.200399: I
tensorflow/core/platform/default/subprocess.cc:304] Start cannot spawn child
process: No such file or directory
2022-11-24 21:21:55.200480: W tensorflow/stream_executor/gpu/asm_compiler.cc:80]
Couldn't get ptxas version string: INTERNAL: Couldn't invoke ptxas --version
2022-11-24 21:21:55.202030: I
tensorflow/core/platform/default/subprocess.cc:304] Start cannot spawn child
process: No such file or directory
2022-11-24 21:21:55.202173: W
tensorflow/stream_executor/gpu/redzone_allocator.cc:314] INTERNAL: Failed to
launch ptxas
Relying on driver to perform ptx compilation.
Modify $PATH to customize ptxas location.
This message will be only logged once.
2022-11-24 21:21:56.387001: W
tensorflow/core/common_runtime/bfc_allocator.cc:290] Allocator (GPU_0_bfc) ran
out of memory trying to allocate 4.21GiB with freed_by_count=0. The caller
indicates that this is not a failure, but this may mean that there could be
performance gains if more memory were available.
2022-11-24 21:21:56.387040: W
tensorflow/core/common_runtime/bfc_allocator.cc:290] Allocator (GPU_0_bfc) ran
out of memory trying to allocate 4.21GiB with freed_by_count=0. The caller
indicates that this is not a failure, but this may mean that there could be
performance gains if more memory were available.
2022-11-24 21:21:56.596477: W
tensorflow/core/common_runtime/bfc_allocator.cc:290] Allocator (GPU_0_bfc) ran
out of memory trying to allocate 4.30GiB with freed_by_count=0. The caller
indicates that this is not a failure, but this may mean that there could be
performance gains if more memory were available.
2022-11-24 21:21:56.596527: W
tensorflow/core/common_runtime/bfc_allocator.cc:290] Allocator (GPU_0_bfc) ran
out of memory trying to allocate 4.30GiB with freed_by_count=0. The caller

```



indicates that this is not a failure, but this may mean that there could be performance gains if more memory were available.

2022-11-24 21:21:56.841150: W  
tensorflow/core/common\_runtime/bfc\_allocator.cc:290] Allocator (GPU\_0\_bfc) ran out of memory trying to allocate 4.54GiB with freed\_by\_count=0. The caller indicates that this is not a failure, but this may mean that there could be performance gains if more memory were available.

2022-11-24 21:21:56.841186: W  
tensorflow/core/common\_runtime/bfc\_allocator.cc:290] Allocator (GPU\_0\_bfc) ran out of memory trying to allocate 4.54GiB with freed\_by\_count=0. The caller indicates that this is not a failure, but this may mean that there could be performance gains if more memory were available.

2538/2538 [=====] - ETA: 0s - loss: 3.1358 - mse: 3.1358 - rmse: 1.7708 - mae: 0.6604 - mape: 23.3414

2022-11-24 21:26:27.013064: W  
tensorflow/core/framework/cpu\_allocator\_impl.cc:82] Allocation of 566231040 exceeds 10% of free system memory.

2022-11-24 21:26:27.218105: W  
tensorflow/core/framework/cpu\_allocator\_impl.cc:82] Allocation of 503316480 exceeds 10% of free system memory.

2022-11-24 21:26:27.565439: W  
tensorflow/core/framework/cpu\_allocator\_impl.cc:82] Allocation of 566231040 exceeds 10% of free system memory.

2022-11-24 21:26:27.940428: W  
tensorflow/core/framework/cpu\_allocator\_impl.cc:82] Allocation of 566231040 exceeds 10% of free system memory.

2022-11-24 21:26:28.237012: W  
tensorflow/core/framework/cpu\_allocator\_impl.cc:82] Allocation of 503316480 exceeds 10% of free system memory.

Epoch 1: val\_loss improved from inf to 0.14867, saving model to model\_weights/20221124-212150\_cnn\_best\_weights.hdf5

2538/2538 [=====] - 278s 107ms/step - loss: 3.1358 - mse: 3.1358 - rmse: 1.7708 - mae: 0.6604 - mape: 23.3414 - val\_loss: 0.1487 - val\_mse: 0.1487 - val\_rmse: 0.3856 - val\_mae: 0.2858 - val\_mape: 8.6893 - lr: 0.0010

Epoch 2/100

2538/2538 [=====] - ETA: 0s - loss: 0.2386 - mse: 0.2386 - rmse: 0.4884 - mae: 0.3452 - mape: 11.8577

Epoch 2: val\_loss did not improve from 0.14867

2538/2538 [=====] - 261s 103ms/step - loss: 0.2386 - mse: 0.2386 - rmse: 0.4884 - mae: 0.3452 - mape: 11.8577 - val\_loss: 0.7105 - val\_mse: 0.7105 - val\_rmse: 0.8429 - val\_mae: 0.7684 - val\_mape: 24.2654 - lr: 0.0010

Epoch 3/100

2538/2538 [=====] - ETA: 0s - loss: 0.1895 - mse:

0.1895 - rmse: 0.4353 - mae: 0.2996 - mape: 10.2114  
Epoch 3: val\_loss did not improve from 0.14867  
2538/2538 [=====] - 260s 102ms/step - loss: 0.1895 - mse: 0.1895 - rmse: 0.4353 - mae: 0.2996 - mape: 10.2114 - val\_loss: 0.3360 - val\_mse: 0.3360 - val\_rmse: 0.5797 - val\_mae: 0.5015 - val\_mape: 15.4744 - lr: 0.0010  
Epoch 4/100  
2538/2538 [=====] - ETA: 0s - loss: 0.1931 - mse: 0.1931 - rmse: 0.4395 - mae: 0.2998 - mape: 10.2581  
Epoch 4: val\_loss improved from 0.14867 to 0.08927, saving model to model\_weights/20221124-212150\_cnn\_best\_weights.hdf5  
2538/2538 [=====] - 262s 103ms/step - loss: 0.1931 - mse: 0.1931 - rmse: 0.4395 - mae: 0.2998 - mape: 10.2581 - val\_loss: 0.0893 - val\_mse: 0.0893 - val\_rmse: 0.2988 - val\_mae: 0.2171 - val\_mape: 6.7587 - lr: 0.0010  
Epoch 5/100  
2538/2538 [=====] - ETA: 0s - loss: 0.1488 - mse: 0.1488 - rmse: 0.3857 - mae: 0.2629 - mape: 8.9841  
Epoch 5: val\_loss did not improve from 0.08927  
2538/2538 [=====] - 263s 104ms/step - loss: 0.1488 - mse: 0.1488 - rmse: 0.3857 - mae: 0.2629 - mape: 8.9841 - val\_loss: 0.1522 - val\_mse: 0.1522 - val\_rmse: 0.3901 - val\_mae: 0.3079 - val\_mape: 9.8317 - lr: 0.0010  
Epoch 6/100  
2538/2538 [=====] - ETA: 0s - loss: 0.1329 - mse: 0.1329 - rmse: 0.3645 - mae: 0.2444 - mape: 8.3617  
Epoch 6: val\_loss did not improve from 0.08927  
2538/2538 [=====] - 262s 103ms/step - loss: 0.1329 - mse: 0.1329 - rmse: 0.3645 - mae: 0.2444 - mape: 8.3617 - val\_loss: 0.2256 - val\_mse: 0.2256 - val\_rmse: 0.4750 - val\_mae: 0.3723 - val\_mape: 11.6466 - lr: 0.0010  
Epoch 7/100  
2538/2538 [=====] - ETA: 0s - loss: 0.1189 - mse: 0.1189 - rmse: 0.3449 - mae: 0.2273 - mape: 7.7810  
Epoch 7: val\_loss did not improve from 0.08927  
2538/2538 [=====] - 263s 104ms/step - loss: 0.1189 - mse: 0.1189 - rmse: 0.3449 - mae: 0.2273 - mape: 7.7810 - val\_loss: 0.1959 - val\_mse: 0.1959 - val\_rmse: 0.4426 - val\_mae: 0.3009 - val\_mape: 9.2265 - lr: 0.0010  
Epoch 8/100  
2538/2538 [=====] - ETA: 0s - loss: 0.1197 - mse: 0.1197 - rmse: 0.3460 - mae: 0.2302 - mape: 7.9117  
Epoch 8: val\_loss did not improve from 0.08927  
2538/2538 [=====] - 265s 104ms/step - loss: 0.1197 - mse: 0.1197 - rmse: 0.3460 - mae: 0.2302 - mape: 7.9117 - val\_loss: 0.1973 - val\_mse: 0.1973 - val\_rmse: 0.4441 - val\_mae: 0.3051 - val\_mape: 9.2742 - lr: 0.0010  
Epoch 9/100

2538/2538 [=====] - ETA: 0s - loss: 0.7158 - mse: 0.7158 - rmse: 0.8461 - mae: 0.3721 - mape: 12.1265  
Epoch 9: val\_loss did not improve from 0.08927  
2538/2538 [=====] - 264s 104ms/step - loss: 0.7158 - mse: 0.7158 - rmse: 0.8461 - mae: 0.3721 - mape: 12.1265 - val\_loss: 0.2933 - val\_mse: 0.2933 - val\_rmse: 0.5416 - val\_mae: 0.4273 - val\_mape: 12.3974 - lr: 0.0010  
Epoch 10/100  
2538/2538 [=====] - ETA: 0s - loss: 0.1877 - mse: 0.1877 - rmse: 0.4333 - mae: 0.2713 - mape: 8.7439  
Epoch 10: val\_loss did not improve from 0.08927  
2538/2538 [=====] - 267s 105ms/step - loss: 0.1877 - mse: 0.1877 - rmse: 0.4333 - mae: 0.2713 - mape: 8.7439 - val\_loss: 0.1482 - val\_mse: 0.1482 - val\_rmse: 0.3850 - val\_mae: 0.2738 - val\_mape: 8.4133 - lr: 1.0000e-04  
Epoch 11/100  
2538/2538 [=====] - ETA: 0s - loss: 0.1694 - mse: 0.1694 - rmse: 0.4116 - mae: 0.2556 - mape: 8.2853  
Epoch 11: val\_loss did not improve from 0.08927  
2538/2538 [=====] - 266s 105ms/step - loss: 0.1694 - mse: 0.1694 - rmse: 0.4116 - mae: 0.2556 - mape: 8.2853 - val\_loss: 0.1605 - val\_mse: 0.1605 - val\_rmse: 0.4006 - val\_mae: 0.2836 - val\_mape: 8.5856 - lr: 1.0000e-04  
Epoch 12/100  
2538/2538 [=====] - ETA: 0s - loss: 0.1631 - mse: 0.1631 - rmse: 0.4039 - mae: 0.2495 - mape: 8.0556  
Epoch 12: val\_loss did not improve from 0.08927  
2538/2538 [=====] - 267s 105ms/step - loss: 0.1631 - mse: 0.1631 - rmse: 0.4039 - mae: 0.2495 - mape: 8.0556 - val\_loss: 0.1635 - val\_mse: 0.1635 - val\_rmse: 0.4043 - val\_mae: 0.2950 - val\_mape: 9.1206 - lr: 1.0000e-04  
Epoch 13/100  
2538/2538 [=====] - ETA: 0s - loss: 0.1572 - mse: 0.1572 - rmse: 0.3965 - mae: 0.2448 - mape: 7.9186  
Epoch 13: val\_loss did not improve from 0.08927  
2538/2538 [=====] - 266s 105ms/step - loss: 0.1572 - mse: 0.1572 - rmse: 0.3965 - mae: 0.2448 - mape: 7.9186 - val\_loss: 0.2036 - val\_mse: 0.2036 - val\_rmse: 0.4512 - val\_mae: 0.3240 - val\_mape: 9.9381 - lr: 1.0000e-04  
Epoch 14/100  
2538/2538 [=====] - ETA: 0s - loss: 0.1391 - mse: 0.1391 - rmse: 0.3729 - mae: 0.2317 - mape: 7.5976  
Epoch 14: val\_loss did not improve from 0.08927  
2538/2538 [=====] - 266s 105ms/step - loss: 0.1391 - mse: 0.1391 - rmse: 0.3729 - mae: 0.2317 - mape: 7.5976 - val\_loss: 0.1761 - val\_mse: 0.1761 - val\_rmse: 0.4197 - val\_mae: 0.3157 - val\_mape: 9.8832 - lr: 1.0000e-04  
Epoch 15/100

```

2538/2538 [=====] - ETA: 0s - loss: 0.1208 - mse:
0.1208 - rmse: 0.3476 - mae: 0.2166 - mape: 7.1439
Epoch 15: val_loss did not improve from 0.08927
2538/2538 [=====] - 271s 107ms/step - loss: 0.1208 -
mse: 0.1208 - rmse: 0.3476 - mae: 0.2166 - mape: 7.1439 - val_loss: 0.1676 -
val_mse: 0.1676 - val_rmse: 0.4093 - val_mae: 0.3035 - val_mape: 9.4201 - lr:
1.0000e-05
Epoch 16/100
2538/2538 [=====] - ETA: 0s - loss: 0.1204 - mse:
0.1204 - rmse: 0.3470 - mae: 0.2162 - mape: 7.1619
Epoch 16: val_loss did not improve from 0.08927
2538/2538 [=====] - 269s 106ms/step - loss: 0.1204 -
mse: 0.1204 - rmse: 0.3470 - mae: 0.2162 - mape: 7.1619 - val_loss: 0.1801 -
val_mse: 0.1801 - val_rmse: 0.4244 - val_mae: 0.3100 - val_mape: 9.6213 - lr:
1.0000e-05
Epoch 17/100
2538/2538 [=====] - ETA: 0s - loss: 0.1172 - mse:
0.1172 - rmse: 0.3423 - mae: 0.2144 - mape: 7.1048
Epoch 17: val_loss did not improve from 0.08927
2538/2538 [=====] - 265s 104ms/step - loss: 0.1172 -
mse: 0.1172 - rmse: 0.3423 - mae: 0.2144 - mape: 7.1048 - val_loss: 0.1724 -
val_mse: 0.1724 - val_rmse: 0.4152 - val_mae: 0.3039 - val_mape: 9.4628 - lr:
1.0000e-05
Epoch 18/100
2538/2538 [=====] - ETA: 0s - loss: 0.1164 - mse:
0.1164 - rmse: 0.3411 - mae: 0.2135 - mape: 7.0689
Epoch 18: val_loss did not improve from 0.08927
2538/2538 [=====] - 265s 104ms/step - loss: 0.1164 -
mse: 0.1164 - rmse: 0.3411 - mae: 0.2135 - mape: 7.0689 - val_loss: 0.1649 -
val_mse: 0.1649 - val_rmse: 0.4061 - val_mae: 0.2997 - val_mape: 9.3374 - lr:
1.0000e-05
Epoch 19/100
2538/2538 [=====] - ETA: 0s - loss: 0.1156 - mse:
0.1156 - rmse: 0.3400 - mae: 0.2132 - mape: 7.0579
Epoch 19: val_loss did not improve from 0.08927
2538/2538 [=====] - 266s 105ms/step - loss: 0.1156 -
mse: 0.1156 - rmse: 0.3400 - mae: 0.2132 - mape: 7.0579 - val_loss: 0.1653 -
val_mse: 0.1653 - val_rmse: 0.4066 - val_mae: 0.3037 - val_mape: 9.4544 - lr:
1.0000e-05
Epoch 19: early stopping

```

## 0.5 Evaluate model

```
[ ]: print(date_actual)
```

```
20221124-212150
```

```
[ ]: #best_model = models.load_model(f'model_weights/{date_actual}_cnn_best_weights.
↳hdf5', custom_objects={"gfb_filter": gfb_filter, "sobel_kernel":
↳sobel_kernel})

best_model = models.load_model(f'model_weights/20221124-212150_cnn_best_weights.
↳hdf5', custom_objects={"gfb_filter": gfb_filter, "sobel_kernel":
↳sobel_kernel})

#best_model = models.load_model(f'model_weights/{date_actual}_cnn_best_weights.
↳hdf5', custom_objects={"gfb_filter": gfb_filter, "sobel_kernel":
↳sobel_kernel})

#best_model = models.load_model(f'best_models_weights/cnn_best_weights_v9.hdf5')
```

```
2022-11-25 11:54:36.549396: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-11-25 11:54:36.549613: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-11-25 11:54:36.549808: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-11-25 11:54:36.549994: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-11-25 11:54:36.550140: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-11-25 11:54:36.550258: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:1616] Created device
/job:localhost/replica:0/task:0/device:GPU:0 with 4063 MB memory: -> device: 0,
name: NVIDIA GeForce RTX 2060, pci bus id: 0000:08:00.0, compute capability: 7.5
```

```
[ ]: def evaluate_model(model, test_values, steps):
    score = model.evaluate(test_values, steps=steps)
    return score
```

```
[ ]: test_loss, test_mse, test_rmse, test_mae, test_mape =
↳evaluate_model(best_model, test_ds, steps=len(test_ds))
```

```
2022-11-25 11:55:43.191337: W
tensorflow/core/common_runtime/bfc_allocator.cc:290] Allocator (GPU_0_bfc) ran
```

out of memory trying to allocate 4.15GiB with freed\_by\_count=0. The caller indicates that this is not a failure, but this may mean that there could be performance gains if more memory were available.

2022-11-25 11:55:43.191387: W  
tensorflow/core/common\_runtime/bfc\_allocator.cc:290] Allocator (GPU\_0\_bfc) ran out of memory trying to allocate 4.15GiB with freed\_by\_count=0. The caller indicates that this is not a failure, but this may mean that there could be performance gains if more memory were available.

2022-11-25 11:55:43.388164: W  
tensorflow/core/common\_runtime/bfc\_allocator.cc:290] Allocator (GPU\_0\_bfc) ran out of memory trying to allocate 4.27GiB with freed\_by\_count=0. The caller indicates that this is not a failure, but this may mean that there could be performance gains if more memory were available.

2022-11-25 11:55:43.388201: W  
tensorflow/core/common\_runtime/bfc\_allocator.cc:290] Allocator (GPU\_0\_bfc) ran out of memory trying to allocate 4.27GiB with freed\_by\_count=0. The caller indicates that this is not a failure, but this may mean that there could be performance gains if more memory were available.

12727/12727 [=====] - 163s 13ms/step - loss: 0.0764 - mse: 0.0764 - rmse: 0.2765 - mae: 0.1920 - mape: 7.2217

```
[ ]: #predictions = best_model.predict(test_ds, steps=np.ceil(test_size /  
    ↪BATCH_SIZE))  
y_pred = best_model.predict(test_ds, steps=len(test_ds))
```

12727/12727 [=====] - 151s 12ms/step

```
[ ]: y_real_test = []
```

```
[ ]: for i in range(len(test_ds)):  
    image_time, stage_discharge = test_ds[i]  
    images = image_time["input_1"]  
    stage = stage_discharge  
  
    #print(stage[0])  
  
    y_real_test.append(stage[0])
```

```
[ ]: y_real_test = np.array(y_real_test)
```

```
[ ]: from sklearn.metrics import r2_score, mean_absolute_percentage_error,  
    ↪mean_absolute_error, mean_squared_error  
from statsmodels.tools.eval_measures import stde
```

```
[ ]: print("R^2: ", r2_score(y_real_test, y_pred))  
print("mse: ", mean_squared_error(y_real_test, y_pred))  
print("rmse: ", mean_squared_error(y_real_test, y_pred, squared=False))
```

```

print("mae: ", mean_absolute_error(y_real_test, y_pred))
print("mape: ", mean_absolute_percentage_error(y_real_test, y_pred))
print("Error estandar: ", stde(y_real_test.squeeze(),
                               y_pred.squeeze(), ddof=2))

```

```

R^2: 0.8043055230971803
mse: 0.07642631697574952
rmse: 0.2764531008611579
mae: 0.19199198128694336
mape: 0.07221717524291128
Error estandar: 0.2659437455875203

```

### 0.5.1 Residual analysis

```

[ ]: residuals = y_real_test - y_pred
residuals_std = residuals/residuals.std()

y_real_stage = y_real_test
residual_stage = residuals

#y_real_discharge = np.array([i[-1] for i in y_test])
#residual_discharge = np.array([i[-1] for i in residuals])

figure, ax = plt.subplots(ncols=2, figsize=(20, 8), dpi=80)


ax[1].scatter(y_real_stage, residual_stage / residual_stage.std(), label="stage_
↳residuals")
#ax[0].scatter(y_real_discharge, residual_discharge / residual_discharge.std(),
↳label="discharge residuals")
ax[1].axhline(y=0.0, color='r', linestyle='-')
ax[0].axhline(y=0.0, color='r', linestyle='-')

ax[1].set_title("Stage residuals")
ax[0].set_title("Discharge residuals")

ax[1].set_xlabel("Fitted values")
ax[0].set_xlabel("Fitted values")
ax[1].set_ylabel("Standarized residuals")
ax[0].set_ylabel("Standarized residuals")


plt.legend()
plt.show()

```



### Check residuals


```
[ ]: plt.hist(residual_stage / residual_stage.std(), density=True, bins = 60)
plt.ylabel('Count')
plt.xlabel('Residual stage');
plt.show()
```






```
[ ]: import statsmodels.api as sm
      from statsmodels.stats.diagnostic import normal_ad
```

```
[ ]: figure = sm.qqplot(residual_stage / residual_stage.std(), line='45',
      ↪label='discharge')
      plt.show()
```



```
[ ]: plt.figure(figsize=(8, 6), dpi=80)
      plt.plot(np.arange(len(y_real_test)), y_real_test, label="Stage real")
```

```
[ ]: [<matplotlib.lines.Line2D at 0x7f828c3ad210>]
```



```
[ ]: figure, ax = plt.subplots(ncols=2, figsize=(20, 8), dpi=80)


ax[0].plot(np.arange(len(y_real_test)), y_real_test, label="Stage real")
ax[0].plot(np.arange(len(y_real_test)), y_pred, label="Stage pred")

ax[0].set_title("Stage predictions")
ax[1].set_title("Discharge predictions")

ax[1].set_ylabel("Values")
ax[0].set_ylabel("Values")
ax[1].set_xlabel("Time")
ax[0].set_xlabel("Time")

ax[0].legend()
ax[1].legend()
plt.show()
```

No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.



## 0.6 Visualize layers

```
[ ]: layer_outputs = [layer.output for layer in best_model.layers[:12]]
# Extracts the outputs of the top 12 layers
activation_model = models.Model(inputs=best_model.input, outputs=layer_outputs)
↪ # Creates a model that will return these outputs, given the model input
```

```
-----
ValueError                                Traceback (most recent call last)
Cell In [115], line 3
      1 layer_outputs = [layer.output for layer in best_model.layers[:12]]
      2 # Extracts the outputs of the top 12 layers
----> 3 activation_model = models.Model(inputs=best_model.input,
    ↪ outputs=layer_outputs)

File ~/miniconda3/envs/tf-gpu/lib/python3.10/site-packages/tensorflow/python/
↪ trackable/base.py:205, in no_automatic_dependency_tracking.<locals>._
↪ method_wrapper(self, *args, **kwargs)
    203 self._self_setattr_tracking = False # pylint: disable=protected-access
    204 try:
--> 205     result = method(self, *args, **kwargs)
    206 finally:
    207     self._self_setattr_tracking = previous_value # pylint:
    ↪ disable=protected-access

File ~/miniconda3/envs/tf-gpu/lib/python3.10/site-packages/keras/engine/
↪ functional.py:165, in Functional.__init__(self, inputs, outputs, name,
↪ trainable, **kwargs)
    156     if not all(
    157         [
    158             functional_utils.is_input_keras_tensor(t)
```

```

159         for t in tf.nest.flatten(inputs)
160     ]
161 ):
162     inputs, outputs = functional_utils.clone_graph_nodes(
163         inputs, outputs
164     )
--> 165 self._init_graph_network(inputs, outputs)

File ~/miniconda3/envs/tf-gpu/lib/python3.10/site-packages/tensorflow/python/
↳ trackable/base.py:205, in no_automatic_dependency_tracking.<locals>.
↳ _method_wrapper(self, *args, **kwargs)
    203 self._self_setattr_tracking = False # pylint: disable=protected-access
    204 try:
--> 205     result = method(self, *args, **kwargs)
    206 finally:
    207     self._self_setattr_tracking = previous_value # pylint:
↳ disable=protected-access

```

```

File ~/miniconda3/envs/tf-gpu/lib/python3.10/site-packages/keras/engine/
↳ functional.py:264, in Functional._init_graph_network(self, inputs, outputs)
    261     self._input_coordinates.append((layer, node_index, tensor_index))
    263 # Keep track of the network's nodes and layers.
--> 264 nodes, nodes_by_depth, layers, _ = _map_graph_network(
    265     self.inputs, self.outputs
    266 )
    267 self._network_nodes = nodes
    268 self._nodes_by_depth = nodes_by_depth

```

```

File ~/miniconda3/envs/tf-gpu/lib/python3.10/site-packages/keras/engine/
↳ functional.py:1128, in _map_graph_network(inputs, outputs)
    1126 for x in tf.nest.flatten(node.keras_inputs):
    1127     if id(x) not in computable_tensors:
-> 1128         raise ValueError(
    1129             f"Graph disconnected: cannot obtain value for "
    1130             f'tensor {x} at layer "{layer.name}". '
    1131             "The following previous layers were accessed "
    1132             f"without issue: {layers_with_complete_input}"
    1133         )
    1134 for x in tf.nest.flatten(node.outputs):
    1135     computable_tensors.add(id(x))

```

```

ValueError: Graph disconnected: cannot obtain value for tensor
↳ KerasTensor(type_spec=TensorSpec(shape=(None, 320, 320, 3), dtype=tf.float32,
↳ name='input_1'), name='input_1', description="created by layer 'input_1'" at
↳ layer "conv1_pad". The following previous layers were accessed without issue:
↳ []

```

```
[ ]: activations = activation_model.predict(test_ds.take(1))
```

1/1 [=====] - 0s 369ms/step

```
[ ]: import matplotlib.pyplot as plt

layer_names = []
for layer in best_model.layers[:12]:
    layer_names.append(layer.name) # Names of the layers, so you can have them
    ↪ as part of your plot

images_per_row = 16

for layer_name, layer_activation in zip(layer_names, activations): # Displays
    ↪ the feature maps
    n_features = layer_activation.shape[-1] # Number of features in the feature
    ↪ map
    size = layer_activation.shape[1] # The feature map has shape (1, size, size,
    ↪ n_features).
    n_cols = n_features // images_per_row # Tiles the activation channels in
    ↪ this matrix
    display_grid = np.zeros((size * n_cols, images_per_row * size))

    print(layer_name)
    if "flatten" in layer_name or "dense" in layer_name: break

    for col in range(n_cols): # Tiles each filter into a big horizontal grid
        for row in range(images_per_row):
            channel_image = layer_activation[0,
                                           :, :,
                                           col * images_per_row + row]

            channel_image -= channel_image.mean() # Post-processes the feature
            ↪ to make it visually palatable
            channel_image /= channel_image.std()
            channel_image *= 64
            channel_image += 128
            channel_image = np.clip(channel_image, 0, 255).astype('uint8')
            display_grid[col * size : (col + 1) * size, # Displays the grid
                          row * size : (row + 1) * size] = channel_image

    scale = 1. / size
    plt.figure(figsize=(scale * display_grid.shape[1],
                        scale * display_grid.shape[0]))
    plt.title(layer_name)
    plt.grid(False)
    plt.imshow(display_grid, aspect='auto', cmap='viridis')
```

conv2d  
max\_pooling2d  
conv2d\_1

```

max_pooling2d_1
conv2d_2
conv2d_3
conv2d_4
max_pooling2d_2
conv2d_5
conv2d_6
max_pooling2d_3

```

```

-----
MemoryError                                Traceback (most recent call last)
Cell In [50], line 13
     11 size = layer_activation.shape[1] #The feature map has shape (1, size,
↪size, n_features).
     12 n_cols = n_features // images_per_row # Tiles the activation channels i
↪this matrix
--> 13 display_grid = np.zeros((size * n_cols, images_per_row * size))
     15 print(layer_name)
     16 if "flatten" in layer_name or "dense" in layer_name: break

MemoryError: Unable to allocate 91.1 GiB for an array with shape (331776, 36864
↪and data type float64

```

