

linear_regression_v1_7

October 21, 2022

1 Linear regression

```
[ ]: import numpy as np
import pandas as pd

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split, cross_val_score, KFold
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LinearRegression
from sklearn.feature_selection import SelectFromModel
from sklearn.metrics import r2_score, mean_absolute_percentage_error, \
    mean_absolute_error, mean_squared_error
from statsmodels.tools.eval_measures import stde
```

1.1 Read the etl info results

```
[ ]: df_info = pd.read_csv('../dataset_clean/options_csv_v1_etl.csv')
df_info
```

```
[ ]: remove_time_features generic_features remove_atypical_values \
0 False True False

feature_combination remove_feature_selection \
0 False False

remove_invalid_correlated_features
0 False
```

1.2 Read the dataset

```
[ ]: df = pd.read_csv('../dataset_clean/PlatteRiverWeir_features_v1_clean.csv')
df
```

```
[ ]: SensorTime CaptureTime Stage Discharge grayMean \
0 2012-06-09 13:15:00 2012-06-09T13:09:07 2.99 916.0 97.405096
```

1	2012-06-09	13:15:00	2012-06-09T13:10:29	2.99	916.0	104.066757
2	2012-06-09	13:45:00	2012-06-09T13:44:01	2.96	873.0	105.636831
3	2012-06-09	14:45:00	2012-06-09T14:44:30	2.94	846.0	104.418949
4	2012-06-09	15:45:00	2012-06-09T15:44:59	2.94	846.0	106.763541
...
42054	2019-10-11	09:00:00	2019-10-11T08:59:53	2.54	434.0	82.872720
42055	2019-10-11	10:00:00	2019-10-11T09:59:52	2.54	434.0	89.028383
42056	2019-10-11	11:00:00	2019-10-11T10:59:52	2.54	434.0	94.722097
42057	2019-10-11	12:00:00	2019-10-11T11:59:53	2.54	434.0	96.693270
42058	2019-10-11	12:45:00	2019-10-11T12:59:52	2.54	434.0	98.738399

	graySigma	entropyMean	entropySigma	hMean	hSigma	\
0	39.623303	0.203417	0.979825	105.368375	41.572939	
1	40.179745	0.206835	1.002624	112.399458	41.795584	
2	40.533218	0.204756	0.994246	114.021526	42.145582	
3	41.752678	0.202428	0.983170	112.612830	43.575351	
4	44.442097	0.202661	0.989625	114.839424	46.302008	
...
42054	57.702652	0.221708	1.076393	87.260572	61.485334	
42055	55.840861	0.233168	1.124774	94.175906	59.006132	
42056	54.355753	0.240722	1.151833	100.534577	56.921028	
42057	52.787629	0.244789	1.171987	102.891159	55.083532	
42058	52.025453	0.252812	1.213278	105.292067	53.994155	

	sMean	sSigma	vMean	vSigma
0	124.520218	4.111846	132.405971	14.983367
1	124.317679	4.270429	133.070221	15.334166
2	124.304621	4.310293	133.294541	15.502448
3	124.369736	4.120586	133.458381	15.190064
4	124.283191	4.088480	133.573595	14.801143
...
42054	127.807813	2.564157	124.073149	13.757842
42055	127.336000	2.585121	124.882812	13.234735
42056	126.958768	2.774867	126.145409	13.408480
42057	126.679956	2.998683	127.508063	13.863205
42058	126.328075	3.258103	128.788256	14.353808

[42059 rows x 14 columns]

```
[ ]: df['SensorTime'] = pd.to_datetime(df['SensorTime'])
df['Year'] = df['SensorTime'].dt.year
```

```
[ ]: df_train = df[(df.Year >= 2012) & (df.Year <= 2017)]
df_test = df[(df.Year >= 2018) & (df.Year <= 2019)]
```

```
[ ]: df_train = df_train.drop(columns=["Year", "SensorTime", "CaptureTime"])
df_test = df_test.drop(columns=["Year", "SensorTime", "CaptureTime"])
```

1.3 Divide dataset to X and Y

```
[ ]: y_train = df_train[["Stage", "Discharge"]]
X_train = df_train.drop(columns=["Stage", "Discharge"])
y_test = df_test[["Stage", "Discharge"]]
X_test = df_test.drop(columns=["Stage", "Discharge"])

[ ]: #X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33,
↳random_state=0)
```

1.4 Train model

```
[ ]: pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('clf', LinearRegression())
])

folds = KFold(n_splits = 5, shuffle = True, random_state = 100)
clf = cross_val_score(pipeline, X_train, y_train, scoring='r2', cv=folds)

[ ]: clf

[ ]: array([0.35480337, 0.34735783, 0.35186074, 0.3379347 , 0.35390495])

[ ]: pipeline.fit(X_train, y_train)

[ ]: Pipeline(steps=[('scaler', StandardScaler()), ('clf', LinearRegression())])
```

1.5 Test Model

```
[ ]: y_pred = pipeline.predict(X_test)

[ ]: print("R^2: ", r2_score(y_test, y_pred))
print("mse: ", mean_squared_error(y_test, y_pred))
print("rmse: ", mean_squared_error(y_test, y_pred, squared=False))
print("mae: ", mean_absolute_error(y_test, y_pred))
print("mape: ", mean_absolute_percentage_error(y_test, y_pred))
print("Error estandar: ", stde(y_test.squeeze(),
    y_pred.squeeze(), ddof=len(X_train.columns) + 1))
```

```
R^2: 0.20765578843687355
mse: 250149.23256003836
rmse: 353.9215663151034
mae: 289.14866835681914
mape: 3.631923146666917e+16
Error estandar: [4.53599216e-01 6.20381370e+02]
```

```
[ ]: residuals = y_test - y_pred
residuals
```

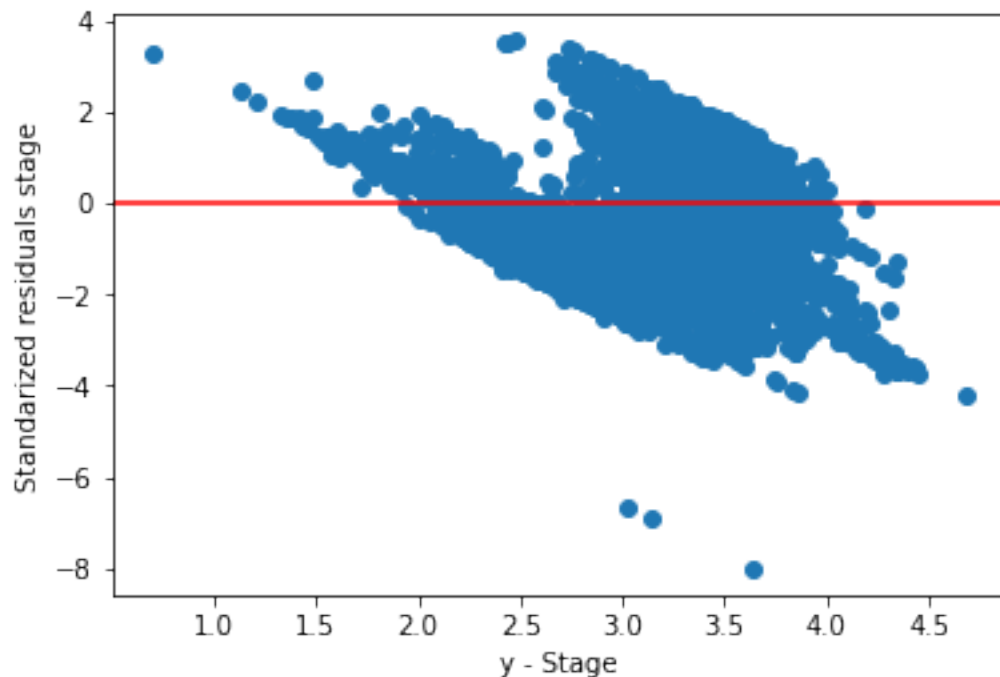
```
[ ]:      Stage    Discharge
28811 -0.634446 -1035.924075
28812  0.860892   845.054275
28813  0.651571   615.180653
28814  0.651607   653.184582
28815 -0.261186  -480.739190
...
42054  0.022499    53.364330
42055  0.052229    80.122600
42056  0.024390    20.047555
42057 -0.007539   -45.129367
42058 -0.007057   -72.281490
```

[13248 rows x 2 columns]

```
[ ]: resid = np.array(residuals["Stage"])
norm_resid = resid / resid.std()

plt.scatter([i[0] for i in y_pred], norm_resid)
plt.axhline(y = 0.0, color = 'r', linestyle = '-')
plt.xlabel("y - Stage")
plt.ylabel("Standardized residuals stage")
```


```
[ ]: Text(0, 0.5, 'Standardized residuals stage')
```



```
[ ]: resid = np.array(residuals["Discharge"])
norm_resid = resid / resid.std()

plt.scatter([i[1] for i in y_pred], norm_resid)
plt.axhline(y = 0.0, color = 'r', linestyle = '-')
plt.xlabel("y - Discharge")
plt.ylabel("Standardized residuals Discharge")
```

```
[ ]: Text(0, 0.5, 'Standardized residuals Discharge')
```



```
[ ]:
```