

linear_regression_v1_8

October 21, 2022

1 Linear regression

```
[ ]: import numpy as np
import pandas as pd

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split, cross_val_score, KFold
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LinearRegression
from sklearn.feature_selection import SelectFromModel
from sklearn.metrics import r2_score, mean_absolute_percentage_error, \
    mean_absolute_error, mean_squared_error
from statsmodels.tools.eval_measures import stde
```

1.1 Read the etl info results

```
[ ]: df_info = pd.read_csv('../dataset_clean/options_csv_v1_etl.csv')
df_info
```

```
[ ]: remove_time_features generic_features remove_atypical_values \
0 False True False

feature_combination remove_feature_selection \
0 False Lasso

remove_invalid_correlated_features
0 False
```

1.2 Read the dataset

```
[ ]: df = pd.read_csv('../dataset_clean/PlatteRiverWeir_features_v1_clean.csv')
df
```

```
[ ]:
      SensorTime      CaptureTime Stage Discharge grayMean \
0 2012-06-09 13:15:00 2012-06-09T13:09:07 2.99 916.0 97.405096
```

1	2012-06-09	13:15:00	2012-06-09T13:10:29	2.99	916.0	104.066757
2	2012-06-09	13:45:00	2012-06-09T13:44:01	2.96	873.0	105.636831
3	2012-06-09	14:45:00	2012-06-09T14:44:30	2.94	846.0	104.418949
4	2012-06-09	15:45:00	2012-06-09T15:44:59	2.94	846.0	106.763541
...
42054	2019-10-11	09:00:00	2019-10-11T08:59:53	2.54	434.0	82.872720
42055	2019-10-11	10:00:00	2019-10-11T09:59:52	2.54	434.0	89.028383
42056	2019-10-11	11:00:00	2019-10-11T10:59:52	2.54	434.0	94.722097
42057	2019-10-11	12:00:00	2019-10-11T11:59:53	2.54	434.0	96.693270
42058	2019-10-11	12:45:00	2019-10-11T12:59:52	2.54	434.0	98.738399

	graySigma	hMean	hSigma
0	39.623303	105.368375	41.572939
1	40.179745	112.399458	41.795584
2	40.533218	114.021526	42.145582
3	41.752678	112.612830	43.575351
4	44.442097	114.839424	46.302008
...
42054	57.702652	87.260572	61.485334
42055	55.840861	94.175906	59.006132
42056	54.355753	100.534577	56.921028
42057	52.787629	102.891159	55.083532
42058	52.025453	105.292067	53.994155

[42059 rows x 8 columns]

```
[ ]: df['SensorTime'] = pd.to_datetime(df['SensorTime'])
df['Year'] = df['SensorTime'].dt.year
```

```
[ ]: df_train = df[(df.Year >= 2012) & (df.Year <= 2017)]
df_test = df[(df.Year >= 2018) & (df.Year <= 2019)]
```

```
[ ]: df_train = df_train.drop(columns=["Year", "SensorTime", "CaptureTime"])
df_test = df_test.drop(columns=["Year", "SensorTime", "CaptureTime"])
```

1.3 Divide dataset to X and Y

```
[ ]: y_train = df_train[["Stage", "Discharge"]]
X_train = df_train.drop(columns=["Stage", "Discharge"])
y_test = df_test[["Stage", "Discharge"]]
X_test = df_test.drop(columns=["Stage", "Discharge"])
```

```
[ ]: #X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33,
↳ random_state=0)
```

1.4 Train model

```
[ ]: pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('clf', LinearRegression())
])

folds = KFold(n_splits = 5, shuffle = True, random_state = 100)
clf = cross_val_score(pipeline, X_train, y_train, scoring='r2', cv=folds)

[ ]: clf

[ ]: array([0.1593821 , 0.15823139, 0.15456937, 0.15749307, 0.17837916])

[ ]: pipeline.fit(X_train, y_train)

[ ]: Pipeline(steps=[('scaler', StandardScaler()), ('clf', LinearRegression())])
```

1.5 Test Model

```
[ ]: y_pred = pipeline.predict(X_test)

[ ]: print("R^2: ", r2_score(y_test, y_pred))
print("mse: ", mean_squared_error(y_test, y_pred))
print("rmse: ", mean_squared_error(y_test, y_pred, squared=False))
print("mae: ", mean_absolute_error(y_test, y_pred))
print("mape: ", mean_absolute_percentage_error(y_test, y_pred))
print("Error estandar: ", stde(y_test.squeeze(),
    y_pred.squeeze(), ddof=len(X_train.columns) + 1))

R^2:  0.11826044004406694
mse:  260848.0265376896
rmse:  361.4301195288088
mae:  304.6086277309594
mape:  6.8141953706081656e+16
Error estandar:  [5.16774324e-01 6.29999958e+02]

[ ]: residuals = y_test - y_pred
residuals

[ ]:
      Stage  Discharge
28811  0.106292 -210.874905
28812 -0.580412 -957.347858
28813 -0.451868 -784.631805
28814  0.071584 -87.625149
28815 -0.424021 -711.479812
...      ...      ...
42054 -0.324914 -471.784825
```

```

42055 -0.413741 -614.366084
42056 -0.497440 -750.010401
42057 -0.561150 -846.566605
42058 -0.656136 -989.026133

```

```
[13248 rows x 2 columns]
```


```

[ ]: resid = np.array(residuals["Stage"])
    norm_resid = resid / resid.std()

    plt.scatter([i[0] for i in y_pred], norm_resid)
    plt.axhline(y = 0.0, color = 'r', linestyle = '-')
    plt.xlabel("y - Stage")
    plt.ylabel("Standardized residuals stage")

```

```
[ ]: Text(0, 0.5, 'Standardized residuals stage')
```



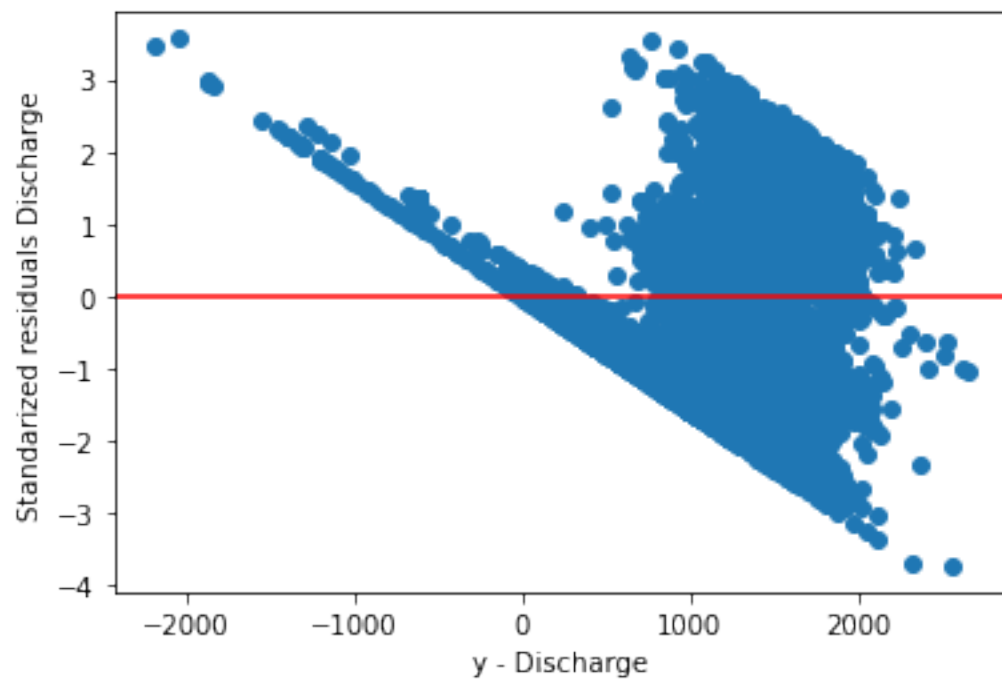
```

[ ]: resid = np.array(residuals["Discharge"])
    norm_resid = resid / resid.std()

    plt.scatter([i[1] for i in y_pred], norm_resid)
    plt.axhline(y = 0.0, color = 'r', linestyle = '-')
    plt.xlabel("y - Discharge")
    plt.ylabel("Standardized residuals Discharge")

```

```
[ ]: Text(0, 0.5, 'Standarized residuals Discharge')
```



```
[ ]:
```