

# MLPRegressor\_v2\_1

November 7, 2022

## 1 MLPRegressor

```
[ ]: import numpy as np
import pandas as pd

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split, RandomizedSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.neural_network import MLPRegressor
from sklearn.feature_selection import SelectFromModel
from sklearn.metrics import r2_score, mean_absolute_percentage_error, \
    mean_absolute_error, mean_squared_error
from statsmodels.tools.eval_measures import stde

import kerastuner as kt
import tensorflow as tf
```

```
[ ]: from tensorflow.python.client import device_lib

print('Default GPU Device: {}'.format(tf.test.gpu_device_name()))
```

Default GPU Device: /device:GPU:0

```
2022-11-06 21:12:45.962030: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-11-06 21:12:45.962248: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-11-06 21:12:45.962403: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-11-06 21:12:45.962598: I
```

```
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-11-06 21:12:45.962753: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2022-11-06 21:12:45.962874: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:1616] Created device
/device:GPU:0 with 4051 MB memory: -> device: 0, name: NVIDIA GeForce RTX 2060,
pci bus id: 0000:08:00.0, compute capability: 7.5
```

## 1.1 Read the etl info results

```
[ ]: df_info = pd.read_csv('../dataset_clean/options_csv_v1_etl.csv')
df_info
```

```
[ ]:      remove_time_features  generic_features  remove_atypical_values  \
0                False                False                False

      feature_combination  remove_feature_selection  \
0                False                False

      remove_invalid_correlated_features
0                False
```

## 1.2 Read the dataset

```
[ ]: df = pd.read_csv('../dataset_clean/PlatteRiverWeir_features_v1_clean.csv')
df
```

```
[ ]:      SensorTime      CaptureTime  Stage  Discharge  grayMean  \
0      2012-06-09 13:15:00  2012-06-09T13:09:07  2.99      916.0  97.405096
1      2012-06-09 13:15:00  2012-06-09T13:10:29  2.99      916.0  104.066757
2      2012-06-09 13:45:00  2012-06-09T13:44:01  2.96      873.0  105.636831
3      2012-06-09 14:45:00  2012-06-09T14:44:30  2.94      846.0  104.418949
4      2012-06-09 15:45:00  2012-06-09T15:44:59  2.94      846.0  106.763541
...      ...      ...      ...      ...      ...
42054  2019-10-11 09:00:00  2019-10-11T08:59:53  2.54      434.0  82.872720
42055  2019-10-11 10:00:00  2019-10-11T09:59:52  2.54      434.0  89.028383
42056  2019-10-11 11:00:00  2019-10-11T10:59:52  2.54      434.0  94.722097
42057  2019-10-11 12:00:00  2019-10-11T11:59:53  2.54      434.0  96.693270
42058  2019-10-11 12:45:00  2019-10-11T12:59:52  2.54      434.0  98.738399

      graySigma  entropyMean  entropySigma      hMean      hSigma  ...  \
0      39.623303      0.203417      0.979825  105.368375  41.572939  ...
1      40.179745      0.206835      1.002624  112.399458  41.795584  ...
```

2	40.533218	0.204756	0.994246	114.021526	42.145582	...
3	41.752678	0.202428	0.983170	112.612830	43.575351	...
4	44.442097	0.202661	0.989625	114.839424	46.302008	...
...	...	...	...	...	...	...
42054	57.702652	0.221708	1.076393	87.260572	61.485334	...
42055	55.840861	0.233168	1.124774	94.175906	59.006132	...
42056	54.355753	0.240722	1.151833	100.534577	56.921028	...
42057	52.787629	0.244789	1.171987	102.891159	55.083532	...
42058	52.025453	0.252812	1.213278	105.292067	53.994155	...

	WeirPt2X	WeirPt2Y	WwRawLineMin	WwRawLineMax	WwRawLineMean	\
0	-1	-1	0.0	0.0	0.000000	
1	-1	-1	0.0	0.0	0.000000	
2	-1	-1	0.0	0.0	0.000000	
3	-1	-1	0.0	0.0	0.000000	
4	-1	-1	0.0	0.0	0.000000	
...	...	...	...	...	...	
42054	2446	1900	9284.0	77521.0	38385.370066	
42055	2440	1900	10092.0	74614.0	40162.989292	
42056	2447	1900	7067.0	83260.0	42095.946590	
42057	2443	1900	6283.0	83045.0	45345.490954	
42058	2436	1900	7375.0	89813.0	47877.870782	

	WwRawLineSigma	WwCurveLineMin	WwCurveLineMax	WwCurveLineMean	\
0	0.000000	0.0	0.0	0.000000	
1	0.000000	0.0	0.0	0.000000	
2	0.000000	0.0	0.0	0.000000	
3	0.000000	0.0	0.0	0.000000	
4	0.000000	0.0	0.0	0.000000	
...	...	...	...	...	
42054	15952.029728	0.0	70085.0	37550.894823	
42055	15467.708856	0.0	70061.0	39397.339095	
42056	16770.357949	0.0	76335.0	41350.006568	
42057	17498.432849	0.0	78882.0	44553.920296	
42058	19963.166359	0.0	82630.0	47280.270559	

	WwCurveLineSigma
0	0.000000
1	0.000000
2	0.000000
3	0.000000
4	0.000000
...	...
42054	16444.401209
42055	16009.008049
42056	17489.374617
42057	18268.294896

42058          20559.358767

[42059 rows x 48 columns]

```
[ ]: df['SensorTime'] = pd.to_datetime(df['SensorTime'])
     df['Year'] = df['SensorTime'].dt.year
```

```
[ ]: df.dtypes
```

```
[ ]: SensorTime      datetime64[ns]
     CaptureTime      object
     Stage            float64
     Discharge        float64
     grayMean         float64
     graySigma        float64
     entropyMean      float64
     entropySigma     float64
     hMean            float64
     hSigma           float64
     sMean            float64
     sSigma           float64
     vMean            float64
     vSigma           float64
     areaFeatCount    int64
     grayMean0        float64
     graySigma0       float64
     entropyMean0     float64
     entropySigma0    float64
     hMean0           float64
     hSigma0          float64
     sMean0           float64
     sSigma0          float64
     vMean0           float64
     vSigma0          float64
     grayMean1        float64
     graySigma1       float64
     entropyMean1     float64
     entropySigma1    float64
     hMean1           float64
     hSigma1          float64
     sMean1           float64
     sSigma1          float64
     vMean1           float64
     vSigma1          float64
     WeirAngle        float64
     WeirPt1X         int64
     WeirPt1Y         int64
```

```

WeirPt2X                int64
WeirPt2Y                int64
WwRawLineMin            float64
WwRawLineMax            float64
WwRawLineMean           float64
WwRawLineSigma          float64
WwCurveLineMin          float64
WwCurveLineMax          float64
WwCurveLineMean         float64
WwCurveLineSigma        float64
Year                    int64
dtype: object

```

### 1.3 Divide dataset to X and Y

```
[ ]: df_train = df[(df.Year >= 2012) & (df.Year <= 2017)]
df_test = df[(df.Year >= 2018) & (df.Year <= 2019)]
```

```
[ ]: df_train = df_train.drop(columns=["Year", "SensorTime", "CaptureTime"])
df_test = df_test.drop(columns=["Year", "SensorTime", "CaptureTime"])
```

```
[ ]: np.random.seed(0)

y_train = df_train[["Stage", "Discharge"]]
X_train = df_train.drop(columns=["Stage", "Discharge"])

temp = list(zip(X_train.values, y_train.values))
np.random.shuffle(temp)
X_train, y_train = zip(*temp)
X_train, y_train = np.array(X_train), np.array(y_train)

y_test = df_test[["Stage", "Discharge"]]
X_test = df_test.drop(columns=["Stage", "Discharge"])

temp = list(zip(X_test.values, y_test.values))
np.random.shuffle(temp)
X_test, y_test = zip(*temp)
X_test, y_test = np.array(X_test), np.array(y_test)
```

```
[ ]: print(X_train.shape)
print(y_train.shape)
```

```

(28811, 44)
(28811, 2)

```

```
[ ]: input_shape = X_train.shape[1]
output_shape = y_train.shape[1]
```

## 1.4 Train model

```
[ ]: def model_builder(hp):
    model = tf.keras.Sequential()
    model.add(tf.keras.Input(shape=input_shape))

    # Tune the number of units in the first Dense layer
    # Choose an optimal value between 32-512
    hp_units = hp.Int('units', min_value = 128, max_value = 1024, step = 32)
    model.add(tf.keras.layers.Dense(units = hp_units, activation = 'relu'))
    hp_units_2 = hp.Int('units_2', min_value = 128, max_value = 512, step = 32)
    model.add(tf.keras.layers.Dense(units = hp_units_2, activation = 'relu'))
    hp_units_3 = hp.Int('units_3', min_value = 64, max_value = 512, step = 32)
    model.add(tf.keras.layers.Dense(units = hp_units_3, activation = 'relu'))
    hp_units_4 = hp.Int('units_4', min_value = 64, max_value = 512, step = 32)
    model.add(tf.keras.layers.Dense(units = hp_units_4, activation = 'relu'))
    """
    hp_units_5 = hp.Int('units_5', min_value = 128, max_value = 512, step = 32)
    model.add(tf.keras.layers.Dense(units = hp_units_5, activation = 'relu'))
    hp_units_6 = hp.Int('units_6', min_value = 64, max_value = 512, step = 32)
    model.add(tf.keras.layers.Dense(units = hp_units_6, activation = 'relu'))
    hp_units_7 = hp.Int('units_7', min_value = 64, max_value = 512, step = 32)
    model.add(tf.keras.layers.Dense(units = hp_units_7, activation = 'relu'))
    hp_units_8 = hp.Int('units_8', min_value = 32, max_value = 512, step = 32)
    model.add(tf.keras.layers.Dense(units = hp_units_8, activation = 'relu'))"""

    model.add(tf.keras.layers.Dense(output_shape))

    # Tune the learning rate for the optimizer
    # Choose an optimal value from 0.01, 0.001, or 0.0001
    hp_learning_rate = hp.Choice('learning_rate', values = [1e-2, 1e-3, 1e-4])

    model.compile(optimizer = tf.keras.optimizers.Adam(learning_rate = hp_learning_rate), loss = 'mse', metrics = ['mse', tf.keras.metrics.
    ↳RootMeanSquaredError(name='rmse'), 'mae', 'mape'])

    return model
```

```
[ ]: tuner = kt.RandomSearch(model_builder,
                             objective = 'val_loss',
                             max_trials = 10,
                             #directory = 'random_search_starter',
                             project_name = 'MLPtf')
```

```
[ ]: tuner.search(X_train, y_train, epochs = 40, validation_data = (X_test, y_test))
```

Trial 10 Complete [00h 01m 18s]  
val\_loss: 109976.1796875

```
Best val_loss So Far: 104425.8203125
Total elapsed time: 00h 13m 04s
INFO:tensorflow:Oracle triggered exit
```

```
[ ]: best_model = tuner.get_best_models(1)[0]
```

```
[ ]: best_hyperparameters = tuner.get_best_hyperparameters(1)[0]
```

## 1.5 Test model

```
[ ]: best_model.evaluate(X_test, y_test)
```

```
414/414 [=====] - 1s 1ms/step - loss: 104425.8203 -
mse: 104425.8203 - rmse: 323.1498 - mae: 144.4191 - mape: 5165232640.0000
```

```
[ ]: [104425.8203125,
      104425.8203125,
      323.14984130859375,
      144.4191131591797,
      5165232640.0]
```

```
[ ]: y_pred = best_model.predict(X_test)
```

```
414/414 [=====] - 0s 601us/step
```

```
[ ]: print("R^2: ", r2_score(y_test, y_pred))
      print("mse: ", mean_squared_error(y_test, y_pred))
      print("rmse: ", mean_squared_error(y_test, y_pred, squared=False))
      print("mae: ", mean_absolute_error(y_test, y_pred))
      print("mape: ", mean_absolute_percentage_error(y_test, y_pred))
      print("Error estandar: ", stde(y_test.squeeze(),
                                     y_pred.squeeze(), ddof=2))
```

```
R^2: -1.5115277087651755
mse: 104425.83337379005
rmse: 229.17087825115055
mae: 144.41915749977616
mape: 2.3262143641528052e+16
Error estandar: [ 1.21269654 431.11023937]
```

```
[ ]: residuals = y_test - y_pred
      residuals_std = residuals/residuals.std()

      y_real_stage = np.array([i[-1] for i in y_test])
      residual_stage = np.array([i[-1] for i in residuals])

      y_real_discharge = np.array([i[0] for i in y_test])
```

```

residual_discharge = np.array([i[0] for i in residuals])

figure, ax = plt.subplots(ncols=2, figsize=(20, 8), dpi=80)


ax[1].scatter(y_real_stage, residual_stage / residual_stage.std(), label="stage_
↳residuals")
ax[0].scatter(y_real_discharge, residual_discharge / residual_discharge.std(),
↳label="discharge residuals")
ax[1].axhline(y=0.0, color='r', linestyle='-')
ax[0].axhline(y=0.0, color='r', linestyle='-')

ax[1].set_title("Stage residuals")
ax[0].set_title("Discharge residuals")

ax[1].set_xlabel("Fitted values")
ax[0].set_xlabel("Fitted values")
ax[1].set_ylabel("Standarized residuals")
ax[0].set_ylabel("Standarized residuals")

plt.legend()
plt.show()

```




```

[ ]: plt.hist(residual_stage / residual_stage.std(), density=True, bins = 60)
plt.ylabel('Count')
plt.xlabel('Residual discharge');
plt.show()

```





```
[ ]: plt.hist(residual_discharge / residual_discharge.std(), density=True, bins = 60)
plt.ylabel('Count')
plt.xlabel('Residual discharge');
plt.show()
```

