

RandomForestRegressor_v1_seg_4

November 24, 2022

1 Random Forest regressor

```
[ ]: import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, RandomizedSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.neural_network import MLPRegressor
from sklearn.feature_selection import SelectFromModel
from sklearn.metrics import r2_score, mean_absolute_percentage_error, \
    mean_absolute_error, mean_squared_error

from sklearn.ensemble import RandomForestRegressor

from statsmodels.tools.eval_measures import stde
```

1.1 Read the etl info results

```
[ ]: df_info = pd.read_csv('../dataset_clean/options_csv_v1_etl.csv')
df_info

[ ]: remove_time_features generic_features remove_atypical_values \
0 False True False

feature_combination remove_feature_selection \
0 False Lasso

remove_invalid_correlated_features
0 False
```

1.2 Read the dataset

```
[ ]: df = pd.read_csv('../dataset_clean/PlatteRiverWeir_features_v1_clean.csv')
df
```

```
[ ]:
      SensorTime      CaptureTime  Stage  Discharge  grayMean \
0      2012-06-09 13:15:00  2012-06-09T13:09:07  2.99      916.0  97.405096
1      2012-06-09 13:15:00  2012-06-09T13:10:29  2.99      916.0  104.066757
2      2012-06-09 13:45:00  2012-06-09T13:44:01  2.96      873.0  105.636831
3      2012-06-09 14:45:00  2012-06-09T14:44:30  2.94      846.0  104.418949
4      2012-06-09 15:45:00  2012-06-09T15:44:59  2.94      846.0  106.763541
...
42054  2019-10-11 09:00:00  2019-10-11T08:59:53  2.54      434.0  82.872720
42055  2019-10-11 10:00:00  2019-10-11T09:59:52  2.54      434.0  89.028383
42056  2019-10-11 11:00:00  2019-10-11T10:59:52  2.54      434.0  94.722097
42057  2019-10-11 12:00:00  2019-10-11T11:59:53  2.54      434.0  96.693270
42058  2019-10-11 12:45:00  2019-10-11T12:59:52  2.54      434.0  98.738399
```

```
      graySigma      hMean      hSigma
0      39.623303  105.368375  41.572939
1      40.179745  112.399458  41.795584
2      40.533218  114.021526  42.145582
3      41.752678  112.612830  43.575351
4      44.442097  114.839424  46.302008
...
42054  57.702652   87.260572  61.485334
42055  55.840861   94.175906  59.006132
42056  54.355753  100.534577  56.921028
42057  52.787629  102.891159  55.083532
42058  52.025453  105.292067  53.994155
```

[42059 rows x 8 columns]

```
[ ]: df['SensorTime'] = pd.to_datetime(df['SensorTime'])
df['Year'] = df['SensorTime'].dt.year
df['Month'] = df['SensorTime'].dt.month
```

```
[ ]: df.dtypes
```

```
[ ]: SensorTime      datetime64[ns]
CaptureTime         object
Stage               float64
Discharge           float64
grayMean            float64
graySigma           float64
hMean               float64
hSigma              float64
Year                int64
```

```
Month                int64
dtype: object
```

```
[ ]: df = df[(df.Stage > 0) & (df.Discharge > 0)]
```

```
[ ]: df.isna().sum()
```

```
[ ]: SensorTime      0
      CaptureTime    0
      Stage          0
      Discharge      0
      grayMean       0
      graySigma      0
      hMean          0
      hSigma         0
      Year           0
      Month          0
      dtype: int64
```

1.3 Divide dataset to X and Y

```
[ ]: np.random.seed(0)

      df_train = df[(df.Year >= 2012) & (df.Year <= 2017)]
      df_train = df_train.iloc[np.random.permutation(len(df_train))]

      df_test = df[(df.Year >= 2018) & (df.Year <= 2019)]
```

```
[ ]: df_train = df_train.drop(columns=["Year", "SensorTime", "CaptureTime"])
      #df_val = df_val.drop(columns=["Year", "SensorTime", "CaptureTime"])
      df_test = df_test.drop(columns=["Year", "SensorTime", "CaptureTime"])
```

```
[ ]: y_train = df_train["Stage"]
      X_train = df_train.drop(columns=["Stage", "Discharge"])

      y_test = df_test["Stage"]
      X_test = df_test.drop(columns=["Stage", "Discharge"])
```

```
[ ]: print(X_train.shape)
      print(y_train.shape)
```

```
(27421, 5)
(27421,)
```

```
[ ]: input_shape = X_train.shape
      output_shape = y_train.shape
```

```
print(input_shape, output_shape)
```

```
(27421, 5) (27421,)
```

1.4 Train model

```
[ ]: pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('clf', RandomForestRegressor(random_state=0))
])

param_grid = {'clf__n_estimators': np.arange(50, 300, 1), 'clf__max_features':
    ["sqrt", 1.0, "log2"]}

clf = RandomizedSearchCV(pipeline, param_distributions=param_grid, n_iter=20,
    n_jobs=8, verbose=3, scoring="neg_mean_squared_error")
```

```
[ ]: clf.fit(X_train, y_train)
```

Fitting 5 folds for each of 20 candidates, totalling 100 fits

```
[CV 1/5] END clf__max_features=log2, clf__n_estimators=193;, score=-0.380 total
time= 7.1s
[CV 5/5] END clf__max_features=log2, clf__n_estimators=193;, score=-0.381 total
time= 7.7s
[CV 3/5] END clf__max_features=log2, clf__n_estimators=193;, score=-0.375 total
time= 7.7s
[CV 3/5] END clf__max_features=log2, clf__n_estimators=206;, score=-0.374 total
time= 7.8s
[CV 4/5] END clf__max_features=log2, clf__n_estimators=193;, score=-0.390 total
time= 8.2s
[CV 2/5] END clf__max_features=log2, clf__n_estimators=193;, score=-0.401 total
time= 8.3s
[CV 2/5] END clf__max_features=log2, clf__n_estimators=206;, score=-0.400 total
time= 8.7s
[CV 1/5] END clf__max_features=log2, clf__n_estimators=206;, score=-0.380 total
time= 8.9s
[CV 2/5] END clf__max_features=log2, clf__n_estimators=118;, score=-0.403 total
time= 4.8s
[CV 1/5] END clf__max_features=log2, clf__n_estimators=118;, score=-0.380 total
time= 5.1s
[CV 3/5] END clf__max_features=log2, clf__n_estimators=118;, score=-0.376 total
time= 4.6s
[CV 4/5] END clf__max_features=log2, clf__n_estimators=118;, score=-0.392 total
time= 4.9s
[CV 5/5] END clf__max_features=log2, clf__n_estimators=118;, score=-0.383 total
time= 4.7s
[CV 4/5] END clf__max_features=log2, clf__n_estimators=206;, score=-0.390 total
time= 8.3s
```

[CV 5/5] END clf__max_features=log2, clf__n_estimators=206;; score=-0.381 total
 time= 8.8s
 [CV 1/5] END clf__max_features=log2, clf__n_estimators=91;; score=-0.380 total
 time= 3.5s
 [CV 2/5] END clf__max_features=log2, clf__n_estimators=91;; score=-0.403 total
 time= 3.9s
 [CV 3/5] END clf__max_features=log2, clf__n_estimators=91;; score=-0.377 total
 time= 3.4s
 [CV 4/5] END clf__max_features=log2, clf__n_estimators=91;; score=-0.393 total
 time= 3.3s
 [CV 1/5] END clf__max_features=1.0, clf__n_estimators=162;; score=-0.383 total
 time= 13.2s
 [CV 5/5] END clf__max_features=log2, clf__n_estimators=91;; score=-0.384 total
 time= 3.7s
 [CV 2/5] END clf__max_features=1.0, clf__n_estimators=162;; score=-0.402 total
 time= 12.9s
 [CV 5/5] END clf__max_features=1.0, clf__n_estimators=162;; score=-0.380 total
 time= 13.1s
 [CV 3/5] END clf__max_features=1.0, clf__n_estimators=162;; score=-0.375 total
 time= 14.2s
 [CV 4/5] END clf__max_features=1.0, clf__n_estimators=162;; score=-0.399 total
 time= 14.2s
 [CV 2/5] END clf__max_features=sqrt, clf__n_estimators=187;; score=-0.401 total
 time= 7.5s
 [CV 1/5] END clf__max_features=sqrt, clf__n_estimators=187;; score=-0.380 total
 time= 8.0s
 [CV 3/5] END clf__max_features=sqrt, clf__n_estimators=187;; score=-0.375 total
 time= 7.5s
 [CV 4/5] END clf__max_features=sqrt, clf__n_estimators=187;; score=-0.390 total
 time= 7.6s
 [CV 5/5] END clf__max_features=sqrt, clf__n_estimators=187;; score=-0.381 total
 time= 8.1s
 [CV 1/5] END clf__max_features=1.0, clf__n_estimators=150;; score=-0.383 total
 time= 11.9s
 [CV 1/5] END clf__max_features=1.0, clf__n_estimators=118;; score=-0.384 total
 time= 9.7s
 [CV 3/5] END clf__max_features=1.0, clf__n_estimators=150;; score=-0.376 total
 time= 12.5s
 [CV 5/5] END clf__max_features=1.0, clf__n_estimators=150;; score=-0.380 total
 time= 11.6s
 [CV 2/5] END clf__max_features=1.0, clf__n_estimators=150;; score=-0.402 total
 time= 12.8s
 [CV 4/5] END clf__max_features=1.0, clf__n_estimators=150;; score=-0.400 total
 time= 12.2s
 [CV 2/5] END clf__max_features=1.0, clf__n_estimators=118;; score=-0.402 total
 time= 10.1s
 [CV 3/5] END clf__max_features=1.0, clf__n_estimators=118;; score=-0.376 total
 time= 10.5s

[CV 2/5] END clf__max_features=log2, clf__n_estimators=166;; score=-0.402 total
 time= 6.0s
 [CV 3/5] END clf__max_features=log2, clf__n_estimators=166;; score=-0.375 total
 time= 6.7s
 [CV 4/5] END clf__max_features=log2, clf__n_estimators=166;; score=-0.391 total
 time= 6.5s
 [CV 1/5] END clf__max_features=log2, clf__n_estimators=166;; score=-0.381 total
 time= 7.2s
 [CV 5/5] END clf__max_features=log2, clf__n_estimators=166;; score=-0.380 total
 time= 6.3s
 [CV 4/5] END clf__max_features=1.0, clf__n_estimators=118;; score=-0.401 total
 time= 10.1s
 [CV 5/5] END clf__max_features=1.0, clf__n_estimators=118;; score=-0.380 total
 time= 9.0s
 [CV 1/5] END clf__max_features=sqrt, clf__n_estimators=209;; score=-0.380 total
 time= 8.2s
 [CV 1/5] END clf__max_features=log2, clf__n_estimators=145;; score=-0.380 total
 time= 6.3s
 [CV 2/5] END clf__max_features=log2, clf__n_estimators=145;; score=-0.402 total
 time= 6.0s
 [CV 4/5] END clf__max_features=sqrt, clf__n_estimators=209;; score=-0.390 total
 time= 8.1s
 [CV 2/5] END clf__max_features=sqrt, clf__n_estimators=209;; score=-0.401 total
 time= 9.0s
 [CV 3/5] END clf__max_features=log2, clf__n_estimators=145;; score=-0.375 total
 time= 6.3s
 [CV 3/5] END clf__max_features=sqrt, clf__n_estimators=209;; score=-0.375 total
 time= 8.3s
 [CV 5/5] END clf__max_features=sqrt, clf__n_estimators=209;; score=-0.381 total
 time= 9.0s
 [CV 4/5] END clf__max_features=log2, clf__n_estimators=145;; score=-0.392 total
 time= 6.0s
 [CV 1/5] END clf__max_features=sqrt, clf__n_estimators=68;; score=-0.383 total
 time= 2.7s
 [CV 5/5] END clf__max_features=log2, clf__n_estimators=145;; score=-0.381 total
 time= 5.9s
 [CV 2/5] END clf__max_features=sqrt, clf__n_estimators=68;; score=-0.403 total
 time= 2.7s
 [CV 3/5] END clf__max_features=sqrt, clf__n_estimators=68;; score=-0.376 total
 time= 3.0s
 [CV 4/5] END clf__max_features=sqrt, clf__n_estimators=68;; score=-0.394 total
 time= 2.6s
 [CV 1/5] END clf__max_features=sqrt, clf__n_estimators=52;; score=-0.383 total
 time= 1.9s
 [CV 5/5] END clf__max_features=sqrt, clf__n_estimators=68;; score=-0.385 total
 time= 2.4s
 [CV 2/5] END clf__max_features=sqrt, clf__n_estimators=52;; score=-0.403 total
 time= 2.3s

[CV 1/5] END clf__max_features=log2, clf__n_estimators=290;; score=-0.379 total
 time= 10.8s
 [CV 3/5] END clf__max_features=sqrt, clf__n_estimators=52;; score=-0.378 total
 time= 2.0s
 [CV 4/5] END clf__max_features=sqrt, clf__n_estimators=52;; score=-0.394 total
 time= 2.0s
 [CV 3/5] END clf__max_features=log2, clf__n_estimators=290;; score=-0.374 total
 time= 11.7s
 [CV 5/5] END clf__max_features=sqrt, clf__n_estimators=52;; score=-0.386 total
 time= 2.3s
 [CV 2/5] END clf__max_features=log2, clf__n_estimators=290;; score=-0.401 total
 time= 12.0s
 [CV 5/5] END clf__max_features=log2, clf__n_estimators=290;; score=-0.380 total
 time= 12.4s
 [CV 4/5] END clf__max_features=log2, clf__n_estimators=290;; score=-0.390 total
 time= 12.6s
 [CV 1/5] END clf__max_features=sqrt, clf__n_estimators=97;; score=-0.381 total
 time= 3.6s
 [CV 3/5] END clf__max_features=sqrt, clf__n_estimators=97;; score=-0.377 total
 time= 3.6s
 [CV 2/5] END clf__max_features=sqrt, clf__n_estimators=97;; score=-0.403 total
 time= 3.6s
 [CV 4/5] END clf__max_features=sqrt, clf__n_estimators=97;; score=-0.392 total
 time= 4.1s
 [CV 5/5] END clf__max_features=sqrt, clf__n_estimators=97;; score=-0.383 total
 time= 4.2s
 [CV 1/5] END clf__max_features=log2, clf__n_estimators=88;; score=-0.380 total
 time= 3.4s
 [CV 3/5] END clf__max_features=log2, clf__n_estimators=88;; score=-0.378 total
 time= 3.4s
 [CV 2/5] END clf__max_features=log2, clf__n_estimators=88;; score=-0.403 total
 time= 3.7s
 [CV 4/5] END clf__max_features=log2, clf__n_estimators=88;; score=-0.394 total
 time= 3.4s
 [CV 5/5] END clf__max_features=log2, clf__n_estimators=88;; score=-0.384 total
 time= 3.3s
 [CV 1/5] END clf__max_features=sqrt, clf__n_estimators=146;; score=-0.380 total
 time= 6.2s
 [CV 1/5] END clf__max_features=1.0, clf__n_estimators=177;; score=-0.384 total
 time= 14.4s
 [CV 2/5] END clf__max_features=sqrt, clf__n_estimators=146;; score=-0.402 total
 time= 5.3s
 [CV 4/5] END clf__max_features=1.0, clf__n_estimators=177;; score=-0.399 total
 time= 13.9s
 [CV 3/5] END clf__max_features=1.0, clf__n_estimators=177;; score=-0.375 total
 time= 15.5s
 [CV 2/5] END clf__max_features=1.0, clf__n_estimators=177;; score=-0.402 total
 time= 15.8s

```

[CV 3/5] END clf__max_features=sqrt, clf__n_estimators=146;; score=-0.375 total
time= 5.6s
[CV 5/5] END clf__max_features=1.0, clf__n_estimators=177;; score=-0.379 total
time= 15.4s
[CV 4/5] END clf__max_features=sqrt, clf__n_estimators=146;; score=-0.392 total
time= 6.0s
[CV 5/5] END clf__max_features=sqrt, clf__n_estimators=146;; score=-0.381 total
time= 5.9s
[CV 1/5] END clf__max_features=sqrt, clf__n_estimators=62;; score=-0.383 total
time= 2.5s
[CV 2/5] END clf__max_features=sqrt, clf__n_estimators=62;; score=-0.403 total
time= 2.2s
[CV 4/5] END clf__max_features=sqrt, clf__n_estimators=62;; score=-0.394 total
time= 2.2s
[CV 3/5] END clf__max_features=sqrt, clf__n_estimators=62;; score=-0.377 total
time= 2.6s
[CV 5/5] END clf__max_features=sqrt, clf__n_estimators=62;; score=-0.385 total
time= 2.2s
[CV 1/5] END clf__max_features=1.0, clf__n_estimators=163;; score=-0.383 total
time= 13.0s
[CV 4/5] END clf__max_features=1.0, clf__n_estimators=163;; score=-0.399 total
time= 12.8s
[CV 3/5] END clf__max_features=1.0, clf__n_estimators=163;; score=-0.375 total
time= 13.3s
[CV 2/5] END clf__max_features=1.0, clf__n_estimators=163;; score=-0.402 total
time= 13.9s
[CV 5/5] END clf__max_features=1.0, clf__n_estimators=163;; score=-0.380 total
time= 13.9s

```

```

[ ]: RandomizedSearchCV(estimator=Pipeline(steps=[('scaler', StandardScaler()),
                                                    ('clf',
                                                     RandomForestRegressor(random_state=0))]),
                        n_iter=20, n_jobs=8,
                        param_distributions={'clf__max_features': ['sqrt', 1.0,
                                                                    'log2'],
                                             'clf__n_estimators': array([ 50, 51,
52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62,
63, 64, 65, 66, 67, 68, 69, 70, 71, 72, 73, 74, 75,
76, 77, 78, 79, 80, 81, 82, 83, 84, 85, 86, 87, 88,
89, 90...
219, 220, 221, 222, 223, 224, 225, 226, 227, 228, 229, 230, 231,
232, 233, 234, 235, 236, 237, 238, 239, 240, 241, 242, 243, 244,
245, 246, 247, 248, 249, 250, 251, 252, 253, 254, 255, 256, 257,
258, 259, 260, 261, 262, 263, 264, 265, 266, 267, 268, 269, 270,
271, 272, 273, 274, 275, 276, 277, 278, 279, 280, 281, 282, 283,
284, 285, 286, 287, 288, 289, 290, 291, 292, 293, 294, 295, 296,
297, 298, 299])}),

```



```
scoring='neg_mean_squared_error', verbose=3)
```

1.5 Test model

```
[ ]: clf.best_score_
```

```
[ ]: -0.38479923491844636
```

```
[ ]: clf.best_params_
```

```
[ ]: {'clf__n_estimators': 290, 'clf__max_features': 'log2'}
```

```
[ ]: clf.score(X_test, y_test)
```

```
[ ]: -0.4719475693452559
```

```
[ ]: y_pred = clf.predict(X_test)
```

```
[ ]: print("R^2: ", r2_score(y_test, y_pred))
print("mse: ", mean_squared_error(y_test, y_pred))
print("rmse: ", mean_squared_error(y_test, y_pred, squared=False))
print("mae: ", mean_absolute_error(y_test, y_pred))
print("mape: ", mean_absolute_percentage_error(y_test, y_pred))
print("Error estandar: ", stde(y_test.squeeze(),
    y_pred.squeeze(), ddof=2))
```

```
R^2: -0.20845196214129502
```

```
mse: 0.4719475693452559
```

```
rmse: 0.6869844025487448
```

```
mae: 0.44318905774581835
```

```
mape: 0.1755325182997975
```

```
Error estandar: 0.6118019412804875
```

```
[ ]: residuals = y_test - y_pred
residuals_std = residuals / residuals.std()

y_real_stage = y_test
residual_stage = residuals

#y_real_discharge = np.array([i[-1] for i in y_test])
#residual_discharge = np.array([i[-1] for i in residuals])

figure, ax = plt.subplots(ncols=2, figsize=(20, 8), dpi=80)

ax[1].scatter(y_real_stage, residual_stage / residual_stage.std(), label="stage_1
    ↪residuals")
```

```


#ax[0].scatter(y_real_discharge, residual_discharge / residual_discharge.std(),
↳ label="discharge residuals")
ax[1].axhline(y=0.0, color='r', linestyle='-')
ax[0].axhline(y=0.0, color='r', linestyle='-')

ax[1].set_title("Stage residuals")
ax[0].set_title("Discharge residuals")

ax[1].set_xlabel("Fitted values")
ax[0].set_xlabel("Fitted values")
ax[1].set_ylabel("Standardized residuals")
ax[0].set_ylabel("Standardized residuals")

plt.legend()
plt.show()

```




```

[ ]: import statsmodels.api as sm
from statsmodels.stats.diagnostic import normal_ad


#figure = sm.qqplot(residual_stage / residual_stage.std(), line='45',
↳ label='stage')
plt.show()

[ ]: figure = sm.qqplot(residual_stage / residual_stage.std(), line='45',
↳ label='discharge')
plt.show()

```



```
[ ]: plt.hist(residual_stage / residual_stage.std(), density=True, bins = 60)
plt.ylabel('Count')
plt.xlabel('Residual stage');
plt.show()
```



```
[ ]: """plt.hist(residual_discharge / residual_discharge.std(), density=True, bins =
    ↳60)
plt.ylabel('Count')
plt.xlabel('Residual discharge');
plt.show()"""
```

```
[ ]: "plt.hist(residual_discharge / residual_discharge.std(), density=True, bins =
60)\nplt.ylabel('Count')\nplt.xlabel('Residual discharge');\nplt.show()"
```

```
[ ]: stat, pval = normal_ad(residual_stage / residual_stage.std())
print("p-value:", pval)


if pval < 0.05:
    print("Hay evidencia de que los residuos no provienen de una distribución_
    ↳normal.")
else:
    print("No hay evidencia para rechazar la hipótesis de que los residuos_
    ↳vienen de una distribución normal.")
```

p-value: 0.0

Hay evidencia de que los residuos no provienen de una distribución normal.

```
[ ]: plt.figure(figsize=(8, 6), dpi=80)
plt.plot(np.arange(len(y_test)), y_test, label="Stage real")
```

```
[ ]: [ <matplotlib.lines.Line2D at 0x7f7bc35c40d0>]
```



```
[ ]: figure, ax = plt.subplots(ncols=2, figsize=(20, 8), dpi=80)

ax[0].plot(np.arange(len(y_test)), y_test, label="Stage real")
ax[0].plot(np.arange(len(y_test)), y_pred, label="Stage pred")

ax[0].set_title("Stage predictions")
ax[1].set_title("Discharge predictions")

ax[1].set_ylabel("Values")
ax[0].set_ylabel("Values")
ax[1].set_xlabel("Time")
ax[0].set_xlabel("Time")

plt.legend()
plt.show()
```

No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.

